

Strategic Plan for a Scientific Software Innovation Institute (S^2I^2) for High Energy Physics DRAFT

Peter Elmer (Princeton University)
Mike Sokoloff (University of Cincinnati)
Mark Neubauer (University of Illinois at Urbana-Champaign)

December 12, 2017



This report has been produced by the S2I2-HEP project (<http://s2i2-hep.org>) and supported by National Science Foundation grants ACI-1558216, ACI-1558219, and ACI-1558233. Any opinions, findings, conclusions or recommendations expressed in this material are those of the project participants and do not necessarily reflect the views of the National Science Foundation.

Executive Summary

The quest to understand the fundamental building blocks of nature and their interactions is one of the oldest and most ambitious of human scientific endeavors. Facilities such as CERN’s Large Hadron Collider (LHC) represent a huge step forward in this quest. The discovery of the Higgs boson, the observation of exceedingly rare decays of B mesons, and stringent constraints on many viable theories of physics beyond the Standard Model (SM) demonstrate the great scientific value of the LHC physics program. The next phase of this global scientific project will be the High-Luminosity LHC (HL-LHC) which will collect data starting circa 2026 and continue into the 2030’s. The primary science goal is to search for physics beyond the SM and, should it be discovered, to study its details and implications. During the HL-LHC era, the ATLAS and CMS experiments will record ~ 10 times as much data from ~ 100 times as many collisions as in Run 1. The NSF and the DOE are planning large investments in detector upgrades so the HL-LHC can operate in this high-rate environment. A commensurate investment in R&D for the software for acquiring, managing, processing and analyzing HL-LHC data will be critical to maximize the return-on-investment in the upgraded accelerator and detectors.

The strategic plan presented in this report is the result of a conceptualization process carried out to explore how a potential Scientific Software Innovation Institute (S^2I^2) for High Energy Physics (HEP) can play a key role in meeting HL-LHC challenges. In parallel, a Community White Paper (CWP) describing the bigger picture was prepared under the auspices of the HEP Software Foundation (HSF). Approximately 260 scientists and engineers participated in more than a dozen workshops during 2016–2017, most jointly sponsored by both HSF and the S^2I^2 -HEP project.

The conceptualization process concluded that the mission of an Institute should be two-fold: it should serve as an active center for software R&D *and* as an intellectual hub for the larger software R&D effort required to ensure the success of the HL-LHC scientific program. Four high-impact R&D areas were identified as highest priority for the U.S. university community: (1) development of advanced algorithms for data reconstruction and triggering; (2) development of highly performant analysis systems that reduce ‘time-to-insight’ and maximize the HL-LHC physics potential; (3) development of data organization, management and access systems for the Exabyte era; (4) leveraging the recent advances in Machine Learning and Data Science. In addition, sustaining the investments in the fabric for distributed high-throughput computing was identified as essential to current and future operations activities. A plan for managing and evolving an S^2I^2 -HEP identifies a set of activities and services that will enable and sustain the Institute’s mission.

As an intellectual hub, the Institute should lead efforts in (1) developing partnerships between HEP and the cyberinfrastructure communities (including Computer Science, Software Engineering, Network Engineering, and Data Science) for novel approaches to meeting HL-LHC challenges, (2) bringing in new effort from U.S. Universities emphasizing professional development and training, and (3) sustaining HEP software and underlying knowledge related to the algorithms and their implementations over the two decades required. HEP is a global, complex, scientific endeavor. These activities will help ensure that the software developed and deployed by a globally distributed community will extend the science reach of the HL-LHC and will be sustained over its lifetime.

The strategic plan for an S^2I^2 targeting HL-LHC physics presented in this report reflects a community vision. Developing, deploying, and maintaining sustainable software for the HL-LHC experiments has tremendous technical and social challenges. The campaign of R&D, testing, and deployment should start as soon as possible to ensure readiness for doing physics when the upgraded accelerator and detectors turn on. An NSF-funded, U.S. university-based S^2I^2 to lead a “software upgrade” will complement the hardware investments being made. In addition to enabling the best possible HL-LHC science, an S^2I^2 -HEP will bring together the larger cyberinfrastructure and HEP communities to study problems and build algorithms and software implementations to address issues of general import for Exabyte scale problems in big science.

50 Contributors

51 *To add: names of individual contributors to both the text of this document and to the formulation*
52 *of the ideas therein, through the workshops, meetings and discussions that took place during the*
53 *conceptualization process.*
54 Title page images are courtesy of CERN.

Contents

55	1 Introduction	1
56	2 Science Drivers	2
57	3 Computing Challenges	4
58	4 Summary of S^2I^2-HEP Conceptualization Process	7
59	5 The HEP Community	9
60	5.1 The HEP Software Ecosystem and Computing Environment	9
61	5.2 Software Development and Processes in the HEP Community	11
62	6 The Institute Role	13
63	6.1 Institute Role within the HEP Community	13
64	6.2 Institute Role in the Software Lifecycle	14
65	6.3 Institute Elements	15
66	7 Strategic Areas for Initial Investment	17
67	7.1 Rationale for choices and prioritization of a university-based S^2I^2	18
68	7.2 Data Analysis Systems	19
69	7.2.1 Challenges and Opportunities	19
70	7.2.2 Current Approaches	20
71	7.2.3 Research and Development Roadmap and Goals	21
72	7.2.4 Impact and Relevance for S^2I^2	24
73	7.3 Reconstruction and Trigger Algorithms	25
74	7.3.1 Challenges	26
75	7.3.2 Current Approaches	27
76	7.3.3 Research and Development Roadmap and Goals	27
77	7.3.4 Impact and Relevance for S^2I^2	28
78	7.4 Applications of Machine Learning	29
79	7.4.1 Opportunities	29
80	7.4.2 Current Approaches	30
81	7.4.3 Research and Development Roadmap and Goals	30
82	7.4.4 Impact and Relevance for S^2I^2	31
83	7.5 Data Organization, Management and Access (DOMA)	32
84	7.5.1 Challenges and Opportunities	32
85	7.5.2 Current Approaches	34
86	7.5.3 Research and Development Roadmap and Goals	35
87	7.5.4 Impact and Relevance for S^2I^2	36
88	7.6 Fabric of distributed high-throughput computing services (OSG)	37
89	7.7 Backbone for Sustainable Software	39
90	8 Institute Organizational Structure and Evolutionary Process	40
91	9 Building Partnerships	42
92	9.1 Partners	42
93	9.2 The Blueprint Process	46
94	10 Metrics for Success (Physics, Software, Community Engagement)	47

96	11 Training, Education and Outreach	47
97	11.1 The HEP Workforce	47
98	11.2 Current Practices	48
99	11.3 Workforce Development in HEP	49
100	11.4 Knowledge that needs to be transferred	50
101	11.5 Roadmap	51
102	11.6 Outreach	52
103	12 Broadening Participation	53
104	13 Sustainability	53
105	14 Risks and Mitigation	53
106	15 Funding Scenarios	56
107	A Appendix - S^2I^2 Strategic Plan Elements	58
108	B Appendix - Workshop List	59

1 Introduction

The High-Luminosity Large Hadron Collider (HL-LHC) is scheduled to start producing data in 2027 and extend the LHC physics program through the 2030s. Its primary science goal is to search for Beyond the Standard Model (BSM) physics, or study its details if there is an intervening discovery. Although the basic constituents of ordinary matter and their interactions are extraordinarily well described by the Standard Model (SM) of particle physics, a quantum field theory built on top of simple but powerful symmetry principles, it is incomplete. For example, most of the gravitationally interacting matter in the universe does not interact via electromagnetic or strong nuclear interactions. As it produces no directly visible signals, it is called dark matter. Its existence and its quantum nature lie outside the SM. Equally as important, the SM does not address fundamental questions related to the detailed properties of its *own* constituent particles or the specific symmetries governing their interactions. To achieve this scientific program, the HL-LHC will record data from 100 times as many proton-proton collisions as did Run 1 of the LHC.

Realizing the full potential of the HL-LHC requires large investments in upgraded hardware. The R&D preparations for these hardware upgrades are underway and the full project funding for the construction phase is expected to begin to flow in the next few years. The two general purpose detectors at the LHC, ATLAS and CMS, are operated by collaborations of more than 3000 scientists each. U.S. personnel constitute about 30% of the collaborators on these experiments. Within the U.S., funding for the construction and operation of ATLAS and CMS is jointly provided by the Department of Energy (DOE) and the National Science Foundation (NSF). Funding for U.S. participation in the LHCb experiment is provided only by the NSF. The NSF is also planning a major role in the hardware upgrade of the ATLAS and CMS detectors for the HL-LHC. This would use the Major Research Equipment and Facilities Construction (MREFC) mechanism with a possible start in 2020.

Similarly, the HL-LHC will require commensurate investment in the research and development necessary to develop and deploy the software to acquire, manage, process, and analyze the data. Current estimates of HL-LHC computing needs significantly exceed what will be possible assuming Moore's Law and more or less constant operational budgets [1]. The underlying nature of computing hardware (processors, storage, networks) is also evolving, the quantity of data to be processed is increasing dramatically, its complexity is increasing, and more sophisticated analyses will be required to maximize the HL-LHC physics yield. The magnitude of the HL-LHC computing problems to be solved will require different approaches. In planning for the HL-LHC, it is critical that all parties agree on the software goals and priorities, and that the efforts tend to complement each other. In this spirit, the HEP Software Foundation (HSF) began a planning exercise in late 2016 to prepare a Community White Paper (CWP). Its goal is to provide a roadmap for software R&D in preparation for the HL-LHC era which would identify and prioritize the software research and development investments required:

1. to enable new approaches to computing and software that can radically extend the physics reach of the detectors; and
2. to achieve improvements in software efficiency, scalability, and performance, and to make use of the advances in CPU, storage, and network technologies;
3. to ensure the long term sustainability of the software through the lifetime of the HL-LHC.

In parallel to the global CWP exercise the U.S. community executed, with NSF funding, a conceptualization process to produce a Strategic Plan for how a Scientific Software Innovation Institute (S^2I^2) could help meet the challenges. Specifically, the S^2I^2 -HEP conceptualization process [2] had three additional goals:

1. to identify specific focus areas for R&D efforts that could be part of an S^2I^2 in the U.S. university community;

2. to build a consensus within the U.S. HEP software community for a common effort; and
3. to engage with experts from the related fields of scientific computing and software engineering to identify topics of mutual interest and build teams for collaborative work to advance the scientific interests of all the communities.

This document, the “*Strategic Plan for a Scientific Software Innovation Institute (S^2I^2) for High Energy Physics*”, is the result of the S^2I^2 -HEP process.

The existing computing system of the LHC experiments is the result of almost 20 years of effort and experience. In addition to addressing the significant future challenges, sustaining the fundamental aspects of what has been built to date is also critical. Fortunately, the collider nature of this physics program implies that essentially all computational challenges are pleasantly parallel. The large LHC collaborations each produce tens of billions of events per year through a mix of simulation and data triggers recorded by their experiments, and all events are mutually independent of each other. This intrinsic simplification from the science itself permits aggregation of distributed computing resources and is well-matched to the use of *high throughput computing* to meet LHC and HL-LHC computing needs. In addition, the LHC today requires more computing resources than will be provided by funding agencies in any single location (such as CERN). Thus *distributed high-throughput computing* (DHTC) will continue to be a fundamental characteristic of the HL-LHC. Continued support for DHTC is essential for the HEP community.

Developing, maintaining and deploying sustainable software for the HL-LHC experiments, given these constraints, is both a technical and a social challenge. An NSF-funded, U.S. university-based Scientific Software Innovation Institute (S^2I^2) can play a primary leadership role in the international HEP community to prepare the “software upgrade” needed to run in parallel with the hardware upgrades planned for the HL-LHC. The Institute will exist within a larger context of international and national projects. It will help build a more cooperative, community process for developing, prototyping, and deploying software. It will drive research and development in a specific set of focus areas (see Section 7) using its own resources directly, and also leveraging them through collaborative efforts. In addition, the Institute will serve as an intellectual hub for the larger community effort in HEP software and computing – it will serve as a center for disseminating knowledge related to the current software and computing landscape, emerging technologies, and tools (see Section 6). It will work closely with its partners to evolve a common vision for future work (see Section 9). To achieve its specific goals, the Executive Director and core personnel will support backbone activities; Area Managers will organize the day to day activities of distributed efforts within each focus area. Goals and resources allocated to all projects will be reviewed on an annual basis, and updated with advice from stakeholders via the Institute’s Steering Board (see Section 8). Altogether, the Institute should serve as both an active software research and development center and as an intellectual hub for the larger software R&D effort required to ensure that the HL-LHC is able to address its Science Driver questions (see Section 2).

2 Science Drivers

An S^2I^2 focused on software required for an upgraded HL-LHC is primarily intended to enable the discovery of Beyond the Standard Model (BSM) physics, or study its details, if there is a discovery before the upgraded accelerator and detectors turn on. To understand why discovering and elucidating BSM physics will be transformative, we need to start with the key concepts of the Standard Model (SM) of particle physics, what they explain, what they do not, and how the HL-LHC will address the latter.

In the past 200 years, physicists have discovered the basic constituents of ordinary matter and they have developed a very successful theory to describe the interactions (forces) among them. All atoms, and the molecules from which they are built, can be described in terms of these constituents.

The nuclei of atoms are bound together by strong nuclear interactions. Their decays result from strong and weak nuclear interactions. Electromagnetic forces bind atoms together, and bind atoms into molecules. The electromagnetic, weak nuclear, and strong nuclear forces are described in terms of quantum field theories. The predictions of these theories are very, very precise, and they have been validated with equally precise experimental measurements. The electromagnetic and weak nuclear interactions are intimately related to each other, but with a fundamental difference: the particle responsible for the exchange of energy and momentum in electromagnetic interactions (the photon) is massless while the corresponding particles responsible for the exchange of energy and momentum in weak interactions (the W and Z bosons) are about 100 times more massive than the proton. A critical element of the SM is the prediction (made more than 50 years ago) that a qualitatively new type of particle, called the Higgs boson, would give mass to the W and Z bosons. Its discovery [3, 4] at CERN’s Large Hadron Collider (LHC) in 2012 confirmed experimentally the last critical element of the SM.

The SM describes essentially all known physics very well, but its mathematical structure and some important empirical evidence tell us that it is incomplete. These observations motivate a large number of SM extensions, generally using the formalism of quantum field theory, to describe BSM physics. For example, “ordinary” matter accounts for only 5% of the mass-energy budget of the universe, while dark matter, which interacts with ordinary matter gravitationally, accounts for 27%. While we know something about dark matter at macroscopic scales, we know nothing about its microscopic, quantum nature, *except* that its particles are not found in the SM and they lack electromagnetic and SM nuclear interactions. BSM physics also addresses a key feature of the observed universe: the apparent dominance of matter over anti-matter. The fundamental processes of leptogenesis and baryogenesis (how electrons and protons, and their heavier cousins, were created in the early universe) are not explained by the SM, nor is the required level of CP violation (the asymmetry between matter and anti-matter under charge and parity conjugation). Constraints on BSM physics come from “conventional” HEP experiments plus others searching for dark matter particles either directly or indirectly.

The LHC was designed to search for the Higgs boson and for BSM physics – goals in the realm of discovery science. The ATLAS and CMS detectors are optimized to observe and measure the direct production and decay of massive particles. They have now begun to measure the properties of the Higgs boson more precisely to test how well they accord with SM predictions.

Where ATLAS and CMS were designed to study high mass particles directly, LHCb was designed to study heavy flavor physics where quantum influences of very high mass particles, too massive to be directly detected at LHC, are manifest in lower energy phenomena. Its primary goal is to look for BSM physics in CP violation (CPV, defined as asymmetries in the decays of particles and their corresponding antiparticles) and rare decays of beauty and charm hadrons. As an example of how one can relate flavor physics to extensions of the SM, Isidori, Nir, and Perez [5] have considered model-independent BSM constraints from measurements of mixing and CP violation. They assume the new fields are heavier than SM fields and construct an effective theory. Then, they “analyze all realistic extensions of the SM in terms of a limited number of parameters (the coefficients of higher dimensional operators).” They determine bounds on an effective coupling strength couplings of their results is that kaon, B_d , B_s , and D^0 mixing and CPV measurements provide powerful constraints that are complementary to each other and often constrain BSM physics more powerfully than direct searches for high mass particles.

The Particle Physics Project Prioritization Panel (P5) issued their *Strategic Plan for U.S. Particle Physics* [6] in May 2014. It was very quickly endorsed by the High Energy Physics Advisory Panel and submitted to the DOE and the NSF. The report says, *we have identified five compelling lines of inquiry that show great promise for discovery over the next 10 to 20 years. These are the Science Drivers:*

- *Use the Higgs boson as a new tool for discovery*

- Pursue the physics associated with neutrino mass
- Identify the new physics of dark matter
- Understand cosmic acceleration: dark matter and inflation
- Explore the unknown: new particles, interactions, and physical principles.

The HL-LHC will address the first, third, and fifth of these using data acquired at twice the energy of Run 1 and with 100 times the luminosity. As the P5 report says,

The recently discovered Higgs boson is a form of matter never before observed, and it is mysterious. What principles determine its effects on other particles? How does it interact with neutrinos or with dark matter? Is there one Higgs particle or many? Is the new particle really fundamental, or is it composed of others? The Higgs boson offers a unique portal into the laws of nature, and it connects several areas of particle physics. Any small deviation in its expected properties would be a major breakthrough.

The full discovery potential of the Higgs will be unleashed by percent-level precision studies of the Higgs properties. The measurement of these properties is a top priority in the physics program of high-energy colliders. The Large Hadron Collider (LHC) will be the first laboratory to use the Higgs boson as a tool for discovery, initially with substantial higher energy running at 14 TeV, and then with ten times more data at the High-Luminosity LHC (HL-LHC). The HL-LHC has a compelling and comprehensive program that includes essential measurements of the Higgs properties.

In addition to HEP experiments, the LHC hosts the one of world’s foremost nuclear physics experiments. “The ALICE Collaboration has built a dedicated heavy-ion detector to exploit the unique physics potential of nucleus-nucleus interactions at LHC energies. [Their] aim is to study the physics of strongly interacting matter at extreme energy densities, where the formation of a new phase of matter, the quark-gluon plasma, is expected. The existence of such a phase and its properties are key issues in QCD for the understanding of confinement and of chiral-symmetry restoration.” [7] In particular, these collisions reproduce the temperatures and pressures of hadronic matter in the very early universe, and so provide a unique window into the physics of that era.

Summary of Physics Motivation: The ATLAS and CMS collaborations published letters of intent to do experiments at the LHC in October 1992, about 25 years ago. At the time, the top quark had not yet be discovered; no one knew if the experiments would discover the Higgs boson, supersymmetry, technicolor, or something completely different. Looking forward, no one can say what will be discovered in the HL-LHC era. However, with data from 100 times the number of collisions recorded in Run 1 the next 20 years are likely to bring even more exciting discoveries.

3 Computing Challenges

During the HL-LHC era (Run 4, starting circa 2026/2027), the ATLAS and CMS experiments intend to record about $10\times$ as much data from $100\times$ as many collisions as they did in in Run 1, and at twice the energy: the Run 1 integrated luminosity for each of these experiments was $\mathcal{L}_{\text{int}} \sim 30 \text{ fb}^{-1}$ at 7 and 8 TeV; for Run 4 it is designed to be $\mathcal{L}_{\text{int}} \sim 3000 \text{ fb}^{-1}$ at 14 TeV by 2035. Mass storage costs will not improve sufficiently to record so much more data, and the projection is that budgets will allow the experiments to collect only a factor of 10 more. For the LHCb experiment, this $100\times$ increase in data and processing over that of Run 1 will start in Run 3 (beginning circa 2021). The software and computing budgets for these experiments are projected to remain flat. Moore’s Law, even if it continues to hold, will not provide the required increase in computing power to enable fully processing all the data. Even assuming the experiments significantly reduce the amount of

data stored per event, the total size of the datasets will be well into the exabyte scale; they will be constrained primarily by costs and funding levels, not by scientific interest. *The overarching goal of an S^2I^2 for HEP will be to maximize the return-on-investment in the upgraded accelerator and detectors to enable break-through scientific discoveries.*

Projections for the HL-LHC start with the operating experience of the LHC to date, and account for the increased luminosity to be provided by the accelerator and the increased sophistication of the detectors. Run 2 started in the summer of 2015, with the bulk of the luminosity being delivered in 2016–2018. The April 2016 Computing Resources Scrutiny Group (CRSG) report to CERN’s Resource Review Board (RRB) report [8] estimated the ALICE, ATLAS, and CMS usage for the full period 2016–2018. A summary is shown in Table 1, along with corresponding numbers for LHCb taken from their 2017 estimate [9]. Altogether, the LHC experiments will be saving more than an exabyte of data in mass storage by the end of Run 2. In their April 2017 report [10], the CSRG says that “growth equivalent to 20%/year [...] towards HL-LHC [...] should be assumed”.

Table 1: Estimated mass storage to be used by the LHC experiments in 2018, at the end of Run 2 data-taking. Numbers extracted from the CSRG report to CERN’s RRB in April 2016 [8] for ALICE, ATLAS, & CMS and taken from LHCb-PUB-2017-019 [9] for LHCb.

Experiment	Disk Usage (PB)	Tape Usage (PB)	Total (PB)
ALICE	98	86	184
ATLAS	164	324	488
CMS	141	247	388
LHCb	41	79	120
Total	444	736	1180




Figure 1: CMS CPU and disk requirement evolution into the first two years of HL-LHC [Sexton-Kennedy2017]

While no one expects such projections to be accurate over 10 years, simple exponentiation predicts a factor of 6 growth. Naively extrapolating resource requirements using today’s software and computing models, the experiments project significantly greater needs. The magnitude of the discrepancy is illustrated in Figs. 1 and 2 for CMS and ATLAS, respectively. The CPU usages are specified in kHS06 years where a “standard” modern core corresponds to about 10 HS06 units. The disk usages are specified in PB. Very crudely, the experiments need 5 times greater resources than will be available to achieve their full science reach. An aggressive and coordinated software R&D program, such as would be possible with an S^2I^2 , can help mitigate this problem.

The challenges for processor technologies are well known [11]. While the number of transistors on integrated circuits doubles every two years (Moore’s Law), power density limitations and ag-




Figure 2: ATLAS CPU and disk requirement evolution into the first three years of HL-LHC, compared to growth rate assuming flat funding. [Campana2017]

gregate power limitations lead to a situation where “conventional” sequential processors are being replaced by vectorized and even more highly parallel architectures. To take of advantage of this increasing computing power demands major changes to the algorithms implemented in our software. Understanding how emerging architectures (from low power processors to parallel architectures like GPUs to more specialized technologies like FPGAs) will allow HEP computing to realize the dramatic growth in computing power required to achieve our science goals will be a central element of an S^2I^2 R&D effort.

Similar challenges exist with storage and network at the scale of HL-LHC [12], with implications for the persistency of data and the computing models and the software supporting them. Limitations in affordable storage pose a major challenge, as does the I/O capacity of ever larger hard disks. While wide area network capacity will probably continue to increase at the required rate, the ability to use it efficiently will need a closer integration with applications. This will require developments in software to support distributed computing (data and workload management, software distribution and data access) and an increasing awareness of the extremely hierarchical view of data, from long latency tape access and medium-latency network access through to the CPU memory hierarchy.

The human and social challenges run in parallel with the technical challenges. All algorithms and software implementations are developed and maintained by flesh and blood individuals, many with unique expertise. What can the community do to help these people contribute most effectively to the larger scientific enterprise?

- How do we train large numbers of novice developers, and smaller numbers of more expert developers and architects, in appropriate software engineering and software design principles and best practices.
- How do we foster effective collaboration within software development teams and across ex-

periments?

- How do we create a culture for designing, developing, and deploying sustainable software?

Learning how to work together as a coherent community, and engage productively with the larger scientific software community, will be critical to the success of the R&D enterprise preparing for the HL-LHC. An S^2I^2 can play a central role in guaranteeing this success.

4 Summary of S^2I^2 -HEP Conceptualization Process

The proposal “Conceptualization of an S^2I^2 Institute for High Energy Physics (S^2I^2 -HEP)” was submitted to the NSF in August 2015. Awards ACI-1558216, ACI-1558219, and ACI-1558233 were made in July 2016, and the S^2I^2 conceptualization project began in Fall 2016. Two major deliverables were foreseen from the conceptualization process in the original S^2I^2 -HEP proposal:

(1) A **Community White Paper (CWP)** [13] describing a global vision for software and computing for the HL-LHC era; this includes discussions of elements that are common to the LHC community as a whole and those that are specific to the individual experiments. It also discusses the relationship of the common elements to the broader HEP and scientific computing communities. Many of the topics discussed are relevant for a HEP S^2I^2 . The CWP document has been prepared and written as an initiative of the HEP Software Foundation. As its purview is greater than an S^2I^2 Strategic Plan, it fully engaged the international HL-LHC community, including U.S. university and national labs personnel. In addition, international and U.S. personnel associated with other HEP experiments participated at all stages. The CWP provides a roadmap for software R&D in preparation for the HL-LHC and for other HL-LHC era HEP experiments. The charge from the Worldwide LHC Computing Grid (WLCG) to the HSF and the LHC experiments [14] says it should identify and prioritize the software research and development investments required:

- to achieve improvements in software efficiency, scalability and performance and to make use of the advances in CPU, storage and network technologies,
- to enable new approaches to computing and software that can radically extend the physics reach of the detectors,
- to ensure the long term sustainability of the software through the lifetime of the HL- LHC.

(2) A separate **Strategic Plan** identifying areas where the U.S. university community can provide leadership and discussing those issues required for an S^2I^2 which are not (necessarily) relevant to the larger community. This is the document you are currently reading. In large measure, it builds on the findings of the CWP. In addition, it addresses the following questions:

- where does the U.S. university community already have expertise and important leadership roles;
- which software elements and frameworks would provide the best educational and training opportunities for students and postdoctoral fellows;
- what types of programs (short courses, short-term fellowships, long-term fellowships, etc.) might enhance the educational reach of an S^2I^2 ;
- possible organizational, personnel and management structures and operational processes; and
- how the investment in an S^2I^2 can be judged and how the investment can be sustained to assure the scientific goals of the HL-LHC.

The Strategic Plan has been prepared in collaboration with members of the U.S. DOE Laboratory community as well as the U.S. university community. Although it is not a project deliverable, an

additional goal of the conceptualization process has been to engage broadly with computer scientists and software engineers, as well as high energy physicists, to build community interest in submitting an S^2I^2 implementation proposal, should there be an appropriate solicitation.

The process to produce these two documents has been built around a series of dedicated workshops, meetings, and special outreach sessions in preexisting workshops. Many of these were organized under the umbrella of the HSF and involved the full international community. A smaller, dedicated set of workshops focused on S^2I^2 - or U.S.- specific topics, including interaction with the Computer Science community. Engagement with the computer science community has been an integral part of the S^2I^2 process from the beginning, exemplified by a workshop dedicated to fostering collaboration between HEP and computer scientists that was held at the University of Illinois in December 2016 (see the workshop report at [15]). S^2I^2 -HEP project Participant Costs funds were used to support the participation of relevant individuals in all types of workshops. A complete list of the workshops held as part of the CWP or to support the S^2I^2 -specific efforts is included in Appendix B.

The community at large was engaged in the CWP and S^2I^2 processes by building on existing communication mechanisms. The involvement of the LHC experiments (including in particular the software and computing coordinators) in the CWP process allowed for communication using the pre-existing experiment channels. To reach out more widely than just to the LHC experiments, specific contacts were made with individuals with software and computing responsibilities in the FNAL muon and neutrino experiments, Belle-II, the Linear Collider community, as well as various national computing organizations. The HSF had, in fact, been building up mailing lists and contact people beyond LHC for about 2 years before the CWP process began, and the CWP process was able to build on that.

Early in the CWP process, a number of working groups were established on topics that were expected to be important parts of the HL-LHC roadmap: *Careers, Staffing and Training; Computing Models, Facilities, and Distributed Computing; Conditions Database; Data Organization, Management and Access; Data Analysis and Interpretation; Data and Software Preservation; Detector Simulation; Event Processing Frameworks; Machine Learning; Physics Generators; Software Development, Deployment and Validation/Verification; Software Trigger and Event Reconstruction; and Visualization.*

In addition, a small set of working groups envisioned at the beginning of the CWP process failed to gather significant community interest or were integrated into the active working groups listed above. These below-threshold working groups were: *Math Libraries; Data Acquisition Software; Various Aspects of Technical Evolution (Software Tools, Hardware, Networking); Monitoring; Security and Access Control; and Workflow and Resource Management.*

The CWP process began with a kick-off workshop at UCSD/SDSC in January 2017 and concluded with a final workshop in June 2017 in Annecy, France. A large number of intermediate topical workshops and meetings were held between these. The CWP process involved a total of ~ 250 participants, listed in Appendix B. The working groups continued to meet virtually to produce their own white papers with completion targeted for early fall 2017. A synthesis full Community White Paper was planned to be ready shortly afterwards. As of early December 2017, most of the working groups have advanced drafts of their documents and the first draft of the synthesis CWP has been distributed for community review and comment; the editorial team has released the second draft, with a final version expected by mid-December 2017.

At the CWP kick-off workshop (in January 2017), each of the (active) working groups defined a charge for itself, as well as a plan for meetings, a Google Group for communication, etc. The precise path for each working group in terms of teleconference meetings and actual in-person sessions or workshops varied from group to group. Each of the active working groups has produced a working group report, which is available from the HSF CWP webpage [13].

The CWP process was intended to assemble the global roadmap for software and computing

for the HL-LHC. In addition, S^2I^2 -specific activities were organized to explore which subset of the global roadmap would be appropriate for a U.S. university-based Software Institute and what role it would play together with other U.S. efforts (including both DOE efforts, the US-ATLAS and US-CMS Operations programs and the Open Science Grid) and with international efforts. In addition the S^2I^2 -HEP conceptualization project investigated how the U.S. HEP community could better collaborate with and leverage the intellectual capacity of the U.S. Computer Science and NSF Sustainable Software (SI2) [16] communities. Two dedicated S^2I^2 HEP/CS workshops were held as well as a dedicated S^2I^2 workshop, co-located with the ACAT conference. In addition numerous outreach activities and discussions took place with the U.S. HEP community and specifically with PIs interested in software and computing R&D.

5 The HEP Community

HEP is a global science. The global nature of the community is both the context and the source of challenges for an S^2I^2 . A fundamental characteristic of this community is its globally distributed knowledge and workforce. The LHC collaborations each comprise thousands of scientists from close to 200 institutions across more than 40 countries. The large size is a response to the complexity of the endeavor. No one person or small team understands all aspects of the experimental program. Knowledge is thus collectively obtained, held, and sustained over the decades long LHC program. Much of that knowledge is curated in software. Tens of millions of lines of code are maintained by many hundreds of physicists and engineers. Software sustainability is fundamental to the knowledge sustainability required for a research program that is expected to last a couple of decades, well into the early 2040s.

5.1 The HEP Software Ecosystem and Computing Environment

The HEP software landscape itself is quite varied. Each HEP experiment requires, at a minimum, “application” software for data acquisition, data handling, data processing, simulation and analysis, as well as related application frameworks, data persistence and libraries. In addition significant “infrastructure” software is required. The scale of the computing environment itself drives some of the complexity and requirements for infrastructure tools. Over the past 20 years, HEP experiments have become large enough to require significantly greater resources than the host laboratory can provide by itself. Collaborating funding agencies typically provide in-kind contributions of computing resources rather than send funding to the host laboratory. Distributed computing is thus essential, and HEP research needs have driven the development of sophisticated software for data management, data access, and workload/workflow management.

These software elements are used 24 hours a day, 7 days a week, over the entire year. They are used by the LHC experiments in the ~ 170 computing centers and national grid infrastructures that are federated via the Worldwide LHC Computing Grid (shown in Figure 3). The U.S. contribution is organized and run by the Open Science Grid [17, 18]. The intrinsic nature of data-intensive collider physics maps very well to the use of high-throughput computing. The computing use ranges from “production” activities that are organized centrally by the experiment (e.g., basic processing of RAW data and high statistics Monte Carlo simulations) to “analysis” activities initiated by individuals or small groups of researchers for their specific research investigations.

Software Stacks: In practice much of the actual software and infrastructure is implemented *independently* by each experiment. This includes managing the software development and deployment process and the resulting software stack. Some of this is a natural result of the intrinsic differences in the actual detectors (scientific instruments) used by each experiment. Independent software stacks are also the healthy result of different experiments and groups making different algorithmic and implementation choices. And last, but not least, each experiment must have control over its

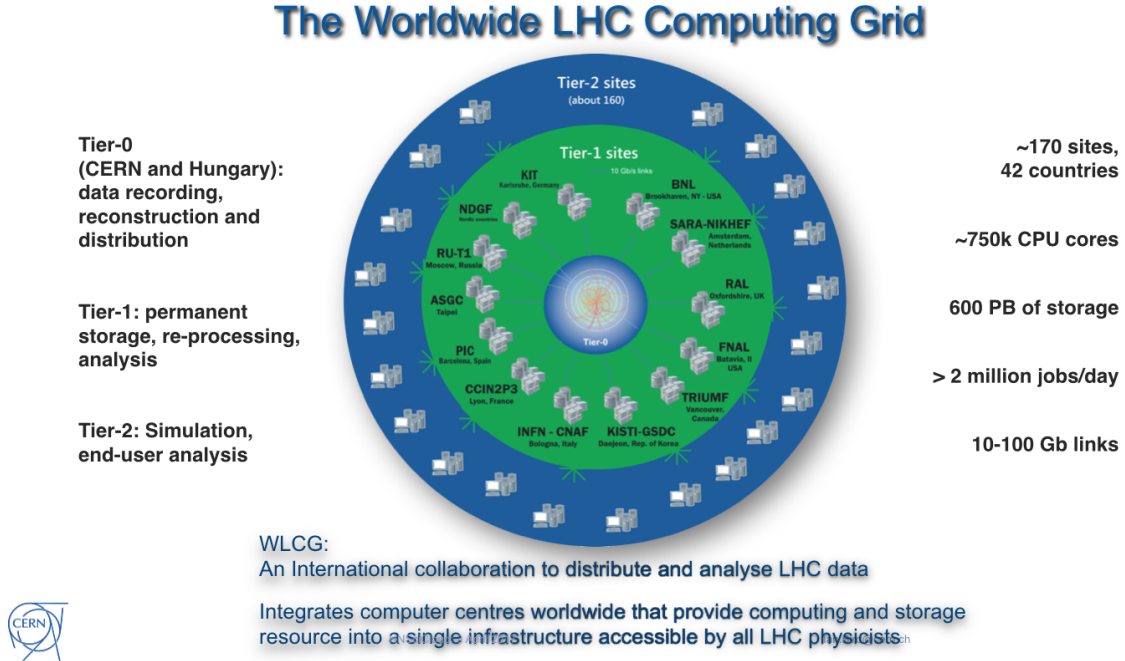


Figure 3: The Worldwide LHC Computing Grid (WLCG), which federates national grid infrastructures to provide the computing resources needed by the four LHC experiments (ALICE, ATLAS, CMS, LHCb). The numbers shown represent the WLCG resources from 2016.

own schedule to insure that it can deliver physics results in a competitive environment. This implies sufficient control over the software development process and the software itself that the experiment uses. The independence of the software processes in each experiment of course has some downsides. At times, similar functionalities are implemented redundantly in multiple experiments. Issues of long term software sustainability can arise in these cases when the particular functionality is not actually mission-critical or specific to the experiment. Obtaining human resources (both in terms of effort and in terms of intellectual input) can be difficult if the result only impacts one particular HEP experiment. Trivial technical and/or communication issues can prevent even high quality tools developed in one experiment from being adopted by another.

The HEP community has nonetheless developed an ecosystem of common software tools that are widely shared in the community. Ideas and experience with software and computing in the HEP community are shared at general dedicated HEP software/computing conferences such as CHEP [19] and ACAT [20]. In addition there are many specialized workshops on software and techniques for pattern recognition, simulation, data acquisition, use of machine learning, etc.

An important exception to the organization of software stacks by the experiments is the national grid infrastructures, such as the Open Science Grid in the U.S. The federation of computing resources from separate computing centers which at times support more than one HEP experiment or that support HEP and other scientific domains requires and creates incentives that drive the development and deployment of “common” solutions.

Application Software Examples: More than 10M lines of code have been developed within individual experiments to implement the relevant data acquisition, data handling, pattern recognition and processing, calibration, simulation and analysis algorithms. This code base includes in addition application frameworks, data persistence and related support libraries needed to structure than myriad algorithms into single data processing applications. Much of the code is experiment-specific

due to real differences in the detectors used by each experiment and the techniques appropriate to the different instruments. Some code is however simply redundant development of different implementations of the same functionalities. This code base contains significant portions which are a by-product of the physics research program (i.e. the result of R&D by postdocs and graduate students) and typically without with the explicit aim of producing sustainable software. Long term sustainability issues exist in many places in such code. One obvious example is the need to develop parallel algorithms and implementations for the increasingly computationally intensive charged particle track reconstruction.

The preparations for the LHC have nonetheless yielded important community software tools for data analysis like ROOT [21] and detector simulation GEANT4 [22–24], both of which have been critical not only for LHC but in most other areas of HEP and beyond. Other tools have been shared between some, but not all, experiments. Examples include the GAUDI [25] event processing framework, IgProf [26] for profiling very large C++ applications like those used in HEP, RooFit [27] for data modeling and fitting and the TMVA [28] toolkit for multivariate data analysis.

In addition software is a critical tool for the interaction and knowledge transfer between experimentalists and theorists. Software provide an important physics input by the theory community to the LHC experimental program, for example through event generators such as SHERPA [29] and ALPGEN [30] and through jet finding tools like FastJet [31, 32].

Infrastructure Software Examples: As noted above, the need for “infrastructure” tools which can be deployed as services in multiple computer centers creates incentives for the development of common tools which can be used by multiple HEP experiments and perhaps with other sciences. Examples include FRONTIER [33] for cached access to databases, XROOTD [34] and dCache [35] for distributed access to bulk file data, EOS [36, 37] for distributed disk storage cluster management, FTS [38] for data movement across the distributed computing system, CERNVM-FS [39] for distributed and cached access to software, GlideinWMS [40] and PanDA [41, 42] for workload management. Although not developed specifically for HEP, HEP has been an important domain-side partner in the development of tools such as HTCondor [43] for distributed high throughput computing and the Parrot [44] virtual file system.

Global scientific collaborations need to meet and discuss, and this has driven the development of the scalable event organization software Indico [45, 46]. Various tools have [XXX \(data and software preservation, Inspire-hep\)...](#)

5.2 Software Development and Processes in the HEP Community

The HEP community has by necessity developed significant experience in creating software infrastructure and processes that integrate contributions from large, distributed communities of physics researchers. To build its software ecosystem, each of the major HEP experiments provides a set of “software architectures and lifecycle processes, development, testing and deployment methodologies, validation and verification processes, end usability and interface considerations, and required infrastructure and technologies” (to quote the NSF S^2I^2 solicitation [47]). Computing hardware to support the development process for the application software (such as continuous integration and test machines) is typically provided by the host laboratory for the experiments, e.g., CERN for the LHC experiments. Each experiment manages software release cycles for its own unique application software code base, as well as external software elements it integrates into its software stack, in order to meet goals ranging from physics needs to bug and performance fixes. The software development infrastructure is also designed to allow individuals to write, test and contribute software from any computing center or laptop/desktop. The software development and testing support for the “infrastructure” part of the software ecosystem, supporting the distributed computing environment, is more diverse and not centralized at CERN. It relies much more heavily on resources such as the Tier-2 centers and the Open Science Grid in the U.S. The integration and testing is more

564 complex for the computing infrastructure software elements, however the full set of processes has
 565 also been put in place by each experiment.




Figure 4: Evolution of the number of individuals making contributions to the CMS application software release each month over the period from 2007 to 2016. Also shown is how the developer community was maintained through large changes to the technical infrastructure, in this case the evolution of the version control system from CVS hosted at CERN to git hosted in GitHub. This plot shows only the application software managed in the experiment-wide software release (CMSSW) and not “infrastructure” software (e.g., for data and workflow management) or “analysis” software developed by individuals or small groups.

566 For the most part, the HEP community has not formally adopted any explicit development
 567 methodology or model, however the de-facto method adopted is very similar to agile software de-
 568 velopment [48]. On slightly longer time scales, the software development efforts within the experi-
 569 ments must respond to various challenges including evolving physics goals and discoveries, general
 570 infrastructure and technology evolution, as well as the evolution of the experiments themselves
 571 (detector upgrades, accelerator energy, and luminosity increases, etc.). HEP experiments have also
 572 maintained these software infrastructures over time scales ranging from years to decades and in
 573 projects involving hundreds to thousands of developers. Figure 4 shows the example of the ap-
 574 plication software release (CMSSW) of CMS experiment at the LHC. Over a ten year period, up
 575 to 300 people were involved in making changes to the software each month. The software process
 576 shown in the figure results in the integration, testing and deployment of tens of releases per year

on the global computing infrastructure. The figure also shows an example of the evolution in the technical infrastructure, in which the code version control system was changed from CVS (hosted at CERN) to git (hosted on GitHub [49]). Similar software processes are also in routine use to develop, integrate, test and deploy the computing infrastructure elements in the software ecosystem which support distributed data management and high throughput computing.

In this section, we described ways in which HEP community develops its software and manages its computing environment to produce physics results. In the next section (Section 6), we present the role of the Institute to facilitate a successful HL-LHC physics program through targeted software development and leadership, more generally, within the HEP software ecosystem.

6 The Institute Role

6.1 Institute Role within the HEP Community

The mission of a Scientific Software Innovation Institute (S^2I^2) for HL-LHC physics should be to serve as both an active software research and development center and as an intellectual hub for the larger R&D effort required to ensure the success of the HL-LHC scientific program. The timeline for the LHC and HL-LHC is shown in Figure 5. A Software Institute operating roughly in the 5 year period from 2019 to 2023 (inclusive) will coincide with two important steps in the ramp up to the HL-LHC: the delivery of the Computing Technical Design Reports (CTDRs) of ATLAS and CMS in ~ 2020 and LHC Run 3 in 2021-2023. The CTDRs will describe the experiments' technical blueprints for building software and computing to maximize the HL-LHC physics reach, given the financial constraints defined by the funding agencies. For ATLAS and CMS, the increased size of the Run 3 data sets relative to Run 2 will not be a major challenge, and changes to the detectors will be modest compared to the upgrades anticipated for Run 4. As a result, ATLAS and CMS will have an opportunity to deploy prototype elements of the HL-LHC computing model during Run 3 as real road tests, even if not at full scale. In contrast, LHCb is making its major transition in terms of how much data will be processed at the onset of Run 3. Some Institute deliverables will be deployed at full scale to directly maximize LHCb physics and provide valuable experience the larger experiments can use to prepare for the HL-LHC.

The Institute will exist within a larger context of international and national projects that are required for software and computing to successfully enable science at the LHC, both today, and in the future. Most importantly at the national level, this includes the U.S. LHC "Operations Programs" jointly funded by DOE and NSF, as well as the Open Science Grid project. In the present section we focus on the role of the Institute while its relationships to these national and international partners are elaborated on in Section 9.

The Institute's mission will be realized by building a more cooperative, community process for developing, prototyping, and deploying software. The Institute itself should be greater than the sum of its parts, and the larger community efforts it engenders should produce better and more sustainable software than would be possible otherwise. Consistent with this mission, the role of the Institute within the HEP community will be to:

1. drive the software R&D process in specific focus areas using its own resources directly, and also leveraging them through collaborative efforts (see Section 7).
2. work closely with the LHC experiments, their U.S. Operations Programs, the relevant national laboratories, and the greater HEP community to identify the highest priority software and computing issues and then create collaborative mechanisms to address them.
3. serve as an intellectual hub for the larger community effort in HEP software and computing. For example, it will bring together a critical mass of experts from HEP, other domain



Figure 5: Timeline for the LHC and HL-LHC [50], indicating both data-taking periods and “shut-down” periods which are used for upgrades of the accelerator and detectors. Data-taking periods are indicated by green lines showing the relative luminosity and red lines showing the center of mass energy. Shutdowns with no data-taking are indicated by blue boxes (LS = Long Shutdown, EYETS = Extended Year End Technical Stop). The approximate periods of execution for an S^2I^2 for HEP and the writing and delivery of the CTRDs are shown in green.

sciences, academic computer science, and the private sector to advise the HEP community on sustainable software development. Similarly, the Institute will serve as a center for disseminating knowledge related to the current software and computing landscape, emerging technologies, and tools. It will provide critical evaluation of new proposed software elements for algorithm essence (e.g. to avoid redundant efforts), feasibility and sustainability, and provide recommendations to collaborations (both experiment and theory) on training, workforce, and software development.

4. deliver value through its (a) contributions to the development of the CTRDs for ATLAS and CMS and (b) research, development and deployment of software that is used for physics during Run 3.

6.2 Institute Role in the Software Lifecycle

Figure 6 shows the elements of the software life cycle, from development of *core concepts and algorithms*, through *prototypes* to deployment of *software products* and *long term support*. The community vision for the Institute is that it will focus its resources on developing innovative ideas and concepts through the prototype stage and along the path to become software products used by the wider community. It will partner with the experiments, the U.S. LHC Operations Programs and others to transition software from the prototype stage to the software product stage. As described in Section 5.2 the experiments already provide full integration, testing deployment and lifecycle processes. The Institute will not duplicate these, but instead collaborate with the experiments and Operations Programs on the efforts required for software integration activities and activities associated to initial deployments of new software products. This may also include the phasing out of older software elements, the transition of existing systems to new modes of working and the consolidation of existing redundant software elements.

The Institute will have a finite lifetime of 5 years (perhaps extensible in a 2nd phase to 10 years), but this is still much shorter than the planned lifetime of HL-LHC activities. The Institute will thus also provide technical support to the experiments and others to identify sustainability and support models for the software products developed. It may at times provide technical support for driving transitions in the HEP software ecosystem which enhance sustainability. In its role as an intellectual hub for HEP software innovation, it will provide advice and guidance broadly on software development within the HEP ecosystem. For example, a new idea or direction under consideration by an experiment could be critically evaluated by the Institute in terms of its essence, novelty, sustainability and impact which would then provide written recommendations for the proposed activity. This will be achieved through having a critical mass of experts in scientific software development inside and outside of HEP and the computer science community who partner with the Institute.

Software Life Cycle

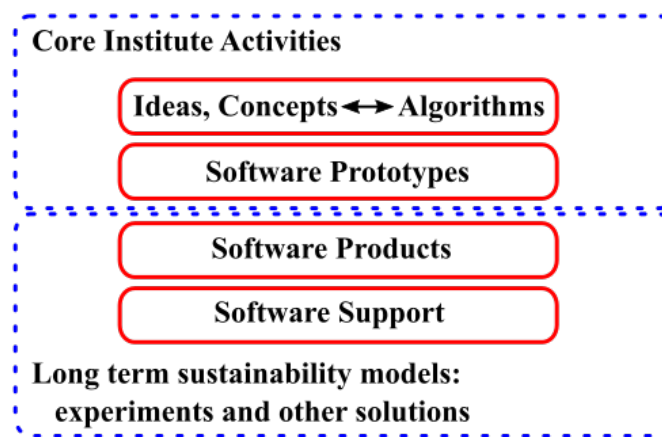


Figure 6: Roles of the Institute in the Software Life Cycle

6.3 Institute Elements

The Institute will have a number of internal functional elements, as shown in Figure 7. (External interactions of the institute will be described in Section 9.)

Institute Management: In order to accomplish its mission, the institute will have a well-defined internal management structure, as well as external governance and advisory structures. Further information on this aspect is provided in Section 8.

Focus Areas: The Institute will have N focus areas, which will pursue the main R&D goals being pursued by the Institute. High priority candidates for these focus areas are described in Section 7. How many of these will be implemented in an Institute implementation will depend on available funding, as described in Section 15. Each focus area will have its own specific plan of work and metrics for evaluation.

Institute Blueprint: The Institute Blueprint activity will maintain the software vision for the Institute and, 3-4 times per year, will bring together expertise to answer specific key questions within the scope of the Institute's activities or, as needed, within the wider scope of HEP software/computing. Blueprint activities will be an essential element to build a common vision with other HEP and HL-LHC R&D efforts, as described in Section 9. The blueprints will then inform

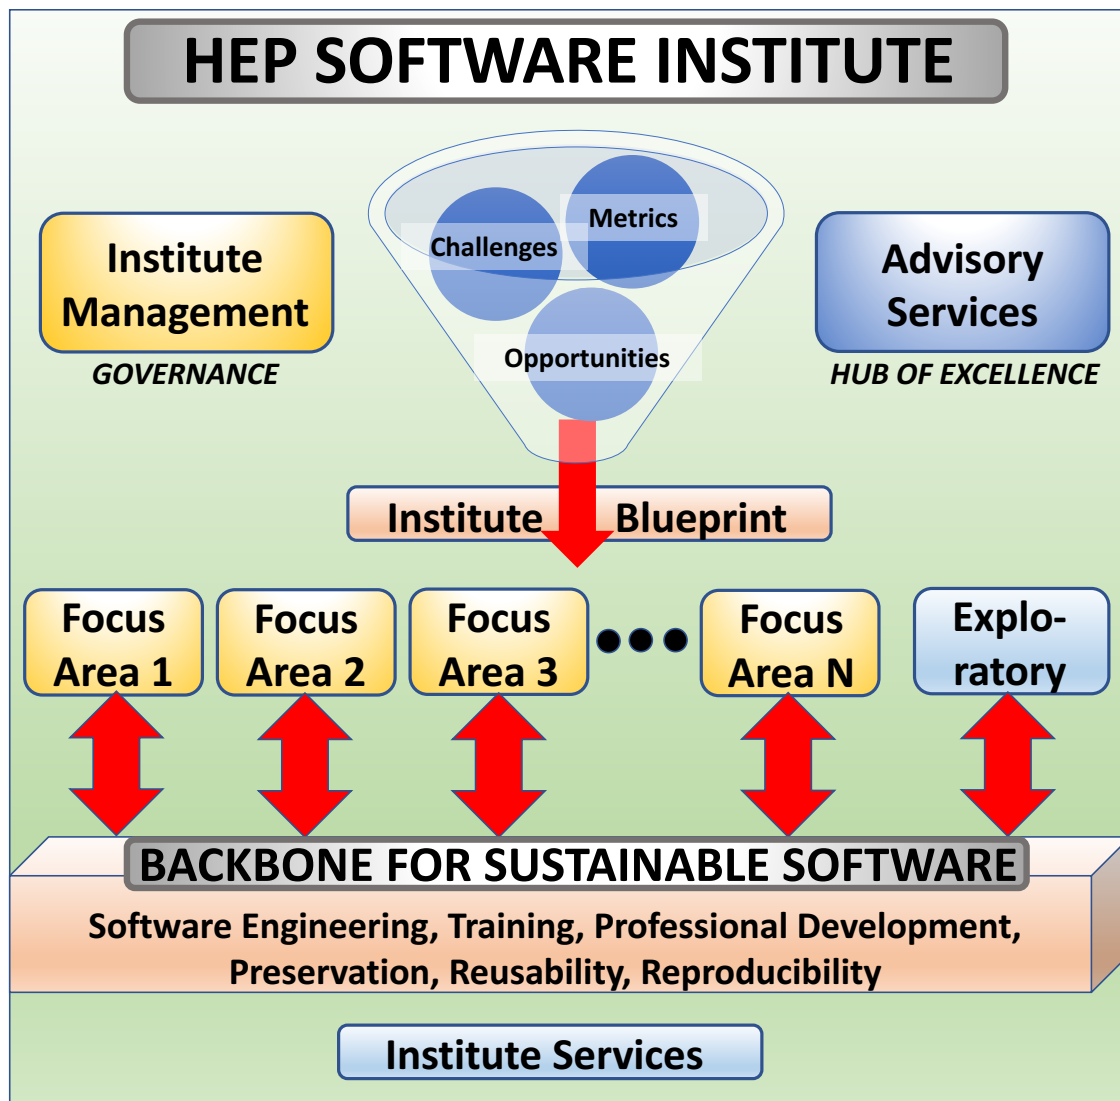


Figure 7: Internal elements of the Institute.

the evolution of both the Institute activities and the overall community HL-LHC R&D objectives in the medium and long term.

Exploratory: From time to time the Institute may deploy modest resources for short term exploratory R&D projects of relevance to inform the planning and overall mission of the Institute.

Backbone for Sustainable Software: In addition to the specific technical advances which will be enabled by the Institute, a dedicated “backbone” activity will focus on how these activities are communicated to students and researchers, identifying best practices and possible incentives, developing and providing training and making data and tools available to the public. Further information on this activity is included in Section 7.7.

Advisory Services: The Institute will play a role in the larger research software community (in HEP and beyond) by being available to provide technical and planning advice to other projects and by participating in reviews. The Institute will execute this functionality both with individuals

directly employed by the Institute and by involving others through its network of partnerships.

Institute Services: The Institute may provide other services in support of its software R&D activities. Possible examples include access to build platforms and continuous integration systems; software stack build and packaging services; technology evaluation services; performance benchmarking services; access to computing resources and related services required for testing of prototypes at scale in the distributed computing environment. In most cases, the actual services will not be owned by the Institute, but instead by one its many partners. The role of the Institute in this case will be to guarantee and coordinate access to the services in support of its mission.

7 Strategic Areas for Initial Investment

A university-based S^2I^2 focused on software needed to ensure the scientific success of the HL-LHC will be part of a larger research, development, and deployment community. It will directly fund and lead some of the R&D efforts; it will support related deployment efforts by the experiments; and it will serve as an intellectual hub for more diverse efforts. The process leading to the CWP, discussed in Section 4, identified three *impact criteria* for judging the value of additional investments, regardless of who makes the investments:

- **Impact - Physics:** Will efforts in this area enable new approaches to computing and software that maximize, and potentially radically extend, the physics reach of the detectors?
- **Impact - Resources/Cost:** Will efforts in this area lead to improvements in software efficiency, scalability and performance and make use of the advances in CPU, storage and network technologies, that allow the experiments to maximize their physics reach within their computing budgets?
- **Impact - Sustainability:** Will efforts in this area significantly improve the long term sustainability of the software through the lifetime of the HL-LHC?

These are key questions for HL-LHC software R&D projects funded by any mechanism, especially an S^2I^2 . During the CWP process, Working Groups (WGs) formed to consider potential activities in areas spanning the HL-LHC software community:

- Careers, Staffing and Training
- Conditions Database
- Computing Models, Facilities and Distributed Computing
- Data Access, Organization and Management
- Data Analysis and Interpretation
- Data and Software Preservation
- Detector Simulation
- Event Processing Frameworks
- Machine Learning
- Physics Generators
- Software Development, Deployment and Validation/Verification
- Software Trigger and Event Reconstruction
- Visualization
- Workflow and Resource Management

Each WG was asked to prepare a section of the CWP including the research and development topics identified in a roadmap for software and computing R&D in HEP for the 2020s, and to evaluate these activities in terms of the impact criteria.

7.1 Rationale for choices and prioritization of a university-based S^2I^2

The S^2I^2 will not be able to solve all of the challenging software problems for the HL-LHC, and it should not take responsibility for deploying and sustaining experiment-specific software. It should instead focus its efforts in targetted areas where R&D will have a high impact on the HL-LHC program. The S^2I^2 needs to align its activities with the expertise of the U.S. university program and with the rest of the community. In addition to identifying areas in which it will lead efforts, the Institute should clearly identify areas in which it will not. These will include some where it will have no significant role at all, and others where it might participate with lower priority.

The S^2I^2 process was largely community-driven. During this process, *additional* S^2I^2 -specific criteria were developed for identifying Focus Areas for the Institute and specific initial R&D topics within each:

- **Interest/Expertise:** Does the U.S. university community have strong interest and expertise in the area?
- **Leadership:** Are the proposed focus areas complementary to efforts funded by the US-LHC Operations programs, the DOE, and international partners?
- **Value:** Is there potential to provide value to more than one HL-LHC experiment and to the wider HEP community?
- **Research/Innovation:** Are there opportunities for combining research and innovation as part of partnerships between the HEP and Computer Science/Software Engineering/Data Science communities?

At the end of the S^2I^2 process, there was a general consensus that highest priority Focus Areas where an S^2I^2 can play a leading role are:

- **Data Analysis Systems:** Modernize and evolve tools and techniques for analysis of high-energy physics data sets. Potential focus areas include adoption of data science tools and approaches, development of analysis systems, analysis resource management, analysis preservation, and visualization for data analytics.
- **Machine Learning Applications:** Exploit Machine Learning approaches to improve the physics reach of HEP data sets. Potential focus areas include track and vertex reconstruction, raw data compression, parameterized simulation methods, and data visualization.
- **Data Organization, Management and Access (DOMA):** Modernize the way HEP organizes, manages, and accesses its data. Potential focus areas include approaches to data persistence, caching, federated data centers, and interactions with networking resources.
- **Reconstruction Algorithms and Software Triggering:** Develop algorithms able to exploit next-generation detector technologies and next-generation computing platforms and programming techniques. Potential focus areas include algorithms for new computing architectures, modernized programming techniques, real-time analysis techniques, and anomaly detection techniques and other approaches that target precision reconstruction and identification techniques enabled by new experimental apparatus and larger data rates.

Two additional potential Focus Areas were identified as medium priority for an S^2I^2 :

- Production Workflow, Workload and Resource Management
- Event Visualization techniques, primarily focusing on collaborative and immersive event displays

Production workflow as well as workload and resource management are absolutely critical software elements for the success of the HL-LHC that will require sustained investment to keep up with the increasing demands. However, the existing operations programs plus other DOE-funded projects are

leading the efforts in these areas. One topic in this area where an S^2I^2 may collaborate extensively is workflows for compute-intensive analysis. Within the S^2I^2 , this can be addressed as part of the Data Analysis Systems focus area. Similarly, there are likely places where the S^2I^2 will collaborate with the visualization community. Specifically, visualization techniques for data analytics and ML analytics can be addressed as part of Data Analysis Systems and ML Applications, respectively.

Although software R&D efforts in each of the following areas will be critical for the success of the HL-LHC, there was a general consensus that other entities are leading the efforts, and these areas should be low priority for S^2I^2 efforts and resources:

- Conditions Database
- Event Processing Frameworks
- Data Acquisition Software
- General Detector Simulation
- Physics Generators
- Network Technology

As is evident from our decision to include elements of production workflow and visualization into higher priority focus areas, the definitions of focus areas are intentionally fluid. In addition, some of the proposed activities intentionally cross nominal boundaries.

7.2 Data Analysis Systems

At the heart of experimental HEP is the development of facilities (e.g. particle colliders, underground laboratories) and instrumentation (e.g. detectors) that provide sensitivity to new phenomena. The analysis and interpretation of data from sophisticated detectors enables HEP to understand the universe at its most fundamental level, including the constituents of matter and their interactions, and the nature of space and time itself. The final stages of data analysis are undertaken by small groups, or individual researchers. The baseline analysis model utilizes successive stages of data reduction, finally reaching a compact dataset for quick real-time iterations. This approach aims at exploiting the maximum possible scientific potential of the data, whilst minimising the “time-to- insight” for a large number of different analyses performed in parallel. Optimizing analysis systems is a complicated combination of diverse constraints, ranging from the need to make efficient use of computing resources to navigating through specific policies of experimental collaborations. Any analysis system has to be flexible enough to handle bursts of activity driven primarily by conference schedules. Future analysis models must also be nimble enough to adapt to new opportunities for discovery (intriguing hints in the data or new experimental signatures), massive increases in data volume by the experiments, and potentially significantly more complex analyses, while still retaining this essential “time-to-insight” optimization.

7.2.1 Challenges and Opportunities

Over the past 20 years the HEP community has developed and primarily utilized an analysis ecosystem centered on ROOT [51]. This software ecosystem currently both dominates HEP analysis and impacts the full event processing chain, providing the core libraries, I/O services, and analysis tools. This approach has certain advantages for the HEP community as compared with other scientific disciplines. It provides an integrated and validated toolkit, lowering the barrier for analysis productivity and enabling the community to speak in a common analysis language. It also facilitates improvements and additions to the toolkit being made available quickly to the community and therefore benefiting a large number of analyses. More recently, open source tools for analysis have become widely available from industry and data science. This newer ecosystem includes data analysis platforms, machine learning tools, and efficient data storage protocols. In many cases, these

tools are evolving very quickly and surpass the HEP efforts both in total investment in analysis software development and the size of communities that use and maintain these tools.

The maintenance and sustainability of the current analysis ecosystem is a challenge. The ecosystem supports a number of use cases and integrates and maintains a wide variety of components. Support for these components has to be prioritized to fit within the available effort, which is provided by a few institutions and not very distributed across the community. Legacy and less used parts of the ecosystem are hard to retire and their continued support strains the available effort. The emergence and abundance of alternative and new analysis components and techniques coming from industry open source projects is also a challenge for the HEP analysis software ecosystem. The community is very interested in using these new techniques and technologies. This leads to additional support needs in order to use the new technologies together with established components of the ecosystem and also be able to interchange old components with new open source components.

Reproducibility is the cornerstone of scientific results. It is currently difficult to repeat most HEP analyses in the same the manner they were originally performed. This difficulty mainly arises due to the number of scientists involved, the large number of steps in a typical HEP analysis workflow, and the complexity of the analyses themselves. A challenge specific to data analysis and interpretation is tracking the evolution of relationships between all the different components of an analysis. Better methods for the preservation of analysis workflows and reuse of analysis software and data products would improve the quality of HEP physics results and reduce the time-to-insight because it would be easier for analyses to progress through increases in data volume and changes in analyses personnel.

Robust methods for data reinterpretation are also critical. Collaborations typically interpret results in the context of specific models for new physics searches and sometimes reinterpret those same searches in the context of alternative theories. However, understanding the full implications of these searches requires the interpretation of the experimental results in the context of many more theoretical models than are currently explored at the time of publication. Analysis reproducibility and reinterpretation strategies need to be considered in all new approaches under investigation, so that they become a fundamental component of the system as a whole.

7.2.2 Current Approaches

Methods for analyzing the data at the LHC experiments have been developed over the years and successfully applied to LHC data to produce physics results during Run 1 and Run 2. The amount of data typically used by a LHC Run 2 data analysis at the LHC (hundreds of TB or PBs) is far too large to be delivered locally to the user. The baseline analysis model utilizes successive stages of data reduction, finally analyzing a compact dataset with quick real time iteration. Experiments and their analysts use a series of processing steps to reduce large input datasets down to sizes suitable for laptop-scale analysis. The line between managed production-like analysis processing and individual analysis, as well as the balance between harmonized vs. individualized analysis data formats differs by experiment, based on their needs and optimization level and the maturity of an experiment in its life cycle.

An evolution of this baseline approach is to produce physics-ready data right from the output of the high-level trigger of the experiment, avoiding the need for any further processing of the data with updated or new software algorithms or detector conditions. The online calibrations are not of sufficient quality to yet enable this approach for all types of analysis, however this approach is now in use across all of the LHC experiments, and will be the primary method used by LHCb in Run 3. Referred to as “real-time analysis”, this technique could be a key enabler of a simplified analysis model that allows simple stripping of data and very efficient data reduction.

The technologies to enable both analysis reproducibility and analysis reinterpretation are evolving quickly. Both require preserving the data and software used for an analysis in some form. This

“analysis capture” is best performed while the analysis is being developed, or at least before it has been published. Recent progress using workflow systems and containerization technology have rapidly transformed this area to provide robust solutions to help analysts adopt techniques that enable reproducibility and reinterpretation of their work.

The LHC collaborations are pursuing a vast number of searches for new physics. Interpretation of these analyses sits at the heart of the LHC physics priorities, and aligns with using the Higgs as a tool for discovery, identify the new physics of dark matter, and explore the unknown of new particles, interactions, and physical principles. The collaborations typically interpret these results in the context of specific models for new physics searches and sometimes reinterpret those same searches in the context of alternative theories. However, understanding the full implications of these searches requires the interpretation of the experimental results in the context of many more theoretical models than are currently explored by the experiments. This is a very active field, with close theory-experiment interaction and with several public tools in development.

For example, a forum [52] on the interpretation of the LHC results for Beyond Standard Model (BSM) studies was initiated to discuss topics related to the BSM (re)interpretation of LHC data, including the development of the necessary public recasting tools [53] and related infrastructure, and to provide a platform for a continued interaction between the theorists and the experiments.

The infrastructure needed for analysis reinterpretation is a focal point of other cyber infrastructure components including the INSPIRE literature database [54], the HEPData data repository [55,56], the CERN Analysis Preservation framework [57,58], and the REANA cloud-based workflow execution system [59]. Critically, this cyber infrastructure sits at the interface between the theoretical community and various experimental collaborations. As a result, this type of infrastructure is not funded through the experiments and tends to fall through the cracks. Thus, it is the perfect topic for a community-wide, cross-collaboration effort.

7.2.3 Research and Development Roadmap and Goals

The goal for future analysis models is to reduce the “time-to-insight” while exploiting the maximum possible scientific potential of the data within the constraints of computing and human resources. Analysis models aim towards giving scientists access to the data in the most interactive and reproducible way possible, to enable quick turn-around in iteratively learning new insights from the data.

Many analyses have common deadlines defined by conference schedules and the availability of physics-quality data samples. The increased analysis activity before these deadlines require the analysis system to be sufficiently elastic to guarantee a rich physics harvest. Models must evolve to take advantage of new computing hardware such as GPUs and new memory as they emerge to reduce the ‘time-to-insight’ further.

Diversification of the Analysis Ecosystem. ROOT and its ecosystem currently dominate HEP analysis and impact the full event processing chain in HEP, providing foundation libraries, I/O services, etc. The analysis tools landscape is now evolving in ways that will influence on the analysis and core software landscape for HL-LHC.

- Data-intensive analysis is growing in importance in other science domains as well as the wider world. Powerful tools from data science and new development initiatives, both within our field and in the wider open source community, have emerged. These tools include software and platforms for visualizing large volumes of complex data and machine learning applications such as TensorFlow, Dask, Pachyderm, Blaze, Parsl, and Thrill. R&D into these tools is needed to enable widespread adoption of these tools in HEP in cases where they can have a big impact on the efficiency of HEP analysts.

- One increasingly important aspect is automation of workflows and the use of automated analysis pipelines. Technologies behind these often leverage open source software such as continuous integration tools. With a lower bar to adoption, these pipeline toolkits could become much more widespread in HEP, with benefits including reduced mechanical work by analysts and enabling analysis reproducibility at a very early stage.
- Notebook interfaces have already demonstrated their value for tutorials and exercises in training sessions and facilitating reproducibility. Remote services like notebook-based analysis-as-a-service should be explored and HEP research tool in addition to education and outreach.
- The HEP community should leverage data formats which are standard within data science, which is critical for gaining access to non-HEP tools, technologies and expertise from computer scientists. We should investigate optimizing some of the more promising formats for late-stage HEP analysis workflows.

Connecting to Modern Cyberinfrastructure. Facilitating easy access and efficient use of modern cyberinfrastructure for analysis workflows will be very important during the HL-LHC due to the anticipated proliferation of such platforms and an increased demand for analysis resources to achieve the physics goals. These include scalable platforms, campus clusters, clouds, and HPC systems, which employ modern and evolving architectures such as GPUs, TPUs, FPGAs, memory-intensive systems, and web services. We should develop mechanisms to instantiate resources for analysis from shared infrastructure as demand arises and share them elastically to support easy, efficient use. An approach gaining a lot of interest for deployment of analysis job payload is containers on grid, cloud, HPC and local resources. The goal is to develop approaches to data analysis which make it easy to utilize heterogeneous resources for analysis workflows. The challenges include making heterogeneous look not so to the analyzers and adapting to changes on resources not directly controlled by analysts or their experiments (both technically and financially).

Functional, Declarative Programming. In a functional approach to programming, an analyst defines what tasks she or he would like the computing system to perform, rather than telling the system how to do it. In this way, scientists express the intended data transformation as a query on data. Instead of having to define and control the ‘how’, the analyst would declare the ‘what’ of their analysis, essentially removing the need to define the event loop in an analysis and leave it to underlying services and systems to optimally iterate over events. This model allows (and gives the responsibility to) the underlying infrastructure to optimize all aspects of the application, including data access patterns and execution concurrency. HEP analysis throughput could be greatly enhanced by switching to a functional or declarative programming model. The HEP community is already investing in R&D projects to enable a functional programming approach (for example TDataFrame in ROOT). Additional R&D projects are needed to develop a full functional or declarative programming language model for HEP analysts.

Improved Non-event Data Handling. An important area that has not received sufficient development is the access to non-event data required for analysis. Example data types include cross-section values, scale factors, efficiencies and fake rate tables, and potentially larger data tables produced by methods such as BDTs or neural networks. Easy storage of non-event data of all sorts of different content, during the analysis step is needed to bring reliable and reproducible access to non-event data just as it currently exists for event data. While a number of ways of doing this have been developed, no commonly accepted and supported way has yet emerged.

High-throughput, Low-latency Analysis Systems. An interesting alternative approach to the current approach to analysis data reduction via a series of time-intensive processing steps is a very low-latency analysis system. To be of interest, an analysis facility would need to provide results, such as histograms, on time-scales short enough to allow many iterations per day by the analyzer.

Two promising, new approaches to data analysis systems in this area are:

- **Spark-like analysis systems.** A new model of data analysis, developed outside of HEP, maintains the concept of sequential ntuple reduction but mixes interactivity with batch processing. Spark is one such system, but TensorFlow, Dask, Pachyderm, and Thrill are others. Distributed processing is either launched as a part of user interaction at a command prompt or wrapped up for batch submission. The key differences from the above are:

1. parallelization is implicit through map/filter/reduce functions.
2. data are abstracted as remote, distributed datasets, rather than files.
3. computation and storage are mixed for data locality: a specialized cluster must be prepared, but can yield higher throughput.

A Spark-like analysis facility would be a shared resource for exploratory data analysis (e.g., making quick plots on data subsets through the spark-shell) and batch submission with the same interface (e.g., substantial jobs through spark-submit). The primary advantage that software products like Spark introduce is in simplifying the user's access to data, lowering the cognitive overhead to setting up and running parallel jobs.

- **Query-based analysis systems.** In one vision for a query-based analysis approach, a series of analysis cycles, each of which provides minimal input (queries of data and code to execute), generates the essential output (histograms, ntuples, etc.) that can be retrieved by the user. The analysis workflow should be accomplished without focus on persistence of data traditionally associated with data reduction, however transient data may could be generated in order to efficiently accomplish this workflow and optionally could be retained to a facilitate an analysis 'checkpoint' for subsequent execution. In this approach, the focus is on obtaining the analysis end-products in a way that does not necessitate a data reduction campaign and associated provisioning of resources. Advantages of a query-based analysis system and its key components include:

1. *Reduced resource needs for Analysis.* A critical consideration of the currently ntuple-driven analysis method is the large CPU and storage requirements for the intermediate data samples. The query-based system provides only the final outcomes of interest (histograms, etc).
2. *Sharing resources with traditional systems.* Unlike a traditional batch system, access to this query system is intermittent and extremely bursty, so it would be hard to justify allocating exclusive resources to it. The query system must share resources with a traditional batch system in such a way that it could elastically scale in response to load.
3. *Fast Columnar Data Caching.* Presenting column partitions ("Columnar cache") to an analysis system as the fundamental unit of data management as opposed to files is an essential feature of the query system. It facilitates retaining input data between queries, which are usually repeated with small modifications (intentionally as part of a systematics study or unplanned as part of normal data exploration).
4. *Provenance.* The query system should also attach enough provenance to each dataset that it could be recreated from the original source data, which is considered immutable. User datasets, while they can't be modified in-place, can be deleted, so a dataset's paper trail must extend all the way back to source data.

Data Interpretation. The LHC provides a large increase in center-of-mass energy over previous collider experiments, starting from 7-8 TeV in Run 1 to 13 TeV during Run 2. The associated large increase in gluon luminosity provided the necessary conditions for discovery of the Higgs boson by the ATLAS and CMS collaborations [60, 61]. Searches for other new particles at high mass

has been a primary focus at the LHC, with lower limits on new particle masses reaching several TeV in many new physics models. The HL-LHC will be an era of increased integrated luminosity rather than increased collision energy. It is conceivable, maybe even likely, that the focus of many analyses during the HL-LHC will shift from direct searches for new particle production to indirect searches for new states with masses beyond the direct reach of the experiments. In this scenario, many LHC analyses will be searching for virtual effects from particles at high-scale, evident only through a detailed study of the kinematics of many events and correlations among many observables. Given its general parameterization of new physics at high-scale, a central framework for this type of analysis is the Effective Field Theory (EFT) extension of the SM [62]. Constraining possible higher-dimensional operators within the context of EFT or other model parameter estimation in high-dimensional spaces using the large datasets afforded by the HL-LHC will be both a challenge and an opportunity for HEP, demanding improvements in analysis techniques, software and computing. An Institute could bring together HEP theorists and experimentalists and computer scientists to tackle the challenges associated with these kinds of generalized interpretations of LHC data. Examples include developing better high-dimensional minimization methods and machine learning approaches to approximate event probability densities [63] and provide likelihood-free inference [64].

Analysis Reproducibility and Reinterpretation. To be successful, analysis reproducibility and reinterpretation need to be considered in all new approaches under investigation and needs to be a fundamental component of the analysis ecosystem as a whole. These considerations become even more critical as we explore analysis models with more heterogeneous hardware and analysis techniques. One specific piece of infrastructure that is currently missing is an analysis database able to represent the many-to-many mapping between publications, logical labels for the event selection defining signal and control regions, data products associated to the application of those event selections to specific datasets, the theoretical models associated to simulated datasets, the multiple implementations of those analyses from the experiments and theoretical community created for the purpose of analysis interpretation, and the results of those interpretations. The protocol for analysis (re)interpretation is clear and narrowly scoped, which makes it possible to offer it as a service. This type of activity lends itself to an “Interpretation Gateway” concept, whose goal is to facility access to shared data, software, computing services, instruments, and related educational materials [65]. Much of the necessary infrastructure is in place to create it [53, 66, 67]. Such an interpretation service would greatly enhance the physics impact of the LHC and also enhance the legacy of the LHC well into the future. An Institute could potentially drive the integration analysis facilities, analysis preservation infrastructure, data repositories, and recasting tools.

7.2.4 Impact and Relevance for S^2I^2

Physics Impact: The very fast turnaround of analysis results that could be possible with new approaches to data access and organization would lead to rapid turnaround for new science.

Resources Impact: Optimized data access for analysis will lead to more efficient use of both CPU and (especially) storage resources. This is essential holding down the overall costs of computing.

Sustainability Impact: This effort would improve the reproducibility and provenance tracking for workflows (especially analysis workflows), making physics analyses more sustainable through the lifetime of the HL-LHC.

Interest/Expertise: University groups have already pioneered significant changes to the data access model for the LHC through the development of federated storage systems, and are prepared to take this further. Other groups are currently exploring the features of modern storage systems and their possible implementation in experiments.

Leadership: Universities are where data analyses for Ph.D. theses are done, together with postdocs and professors. There is also much to be gained within the US physics effort for HL-LHC by focusing on improving the last-mile of analysis computing. Therefore, it is natural for the US Universities to lead in the development of data analysis systems, especially considering the potential for computer science colleagues to collaborate on innovative approaches to such systems. This is also an area where partnership with national labs can be very productive, since the much of technical infrastructure to develop these systems at required scales are at the labs.

Value: All LHC experiments will benefit from new methods of data access and organization, although the implementations may vary due to the different data formats and computing models of each experiment.

Research/Innovation: This effort would rely on partnerships with data storage and access experts in the CS community, some of whom are already providing consultation in this area.

7.3 Reconstruction and Trigger Algorithms

The real-time processing in the trigger and reconstruction of raw detector data (real and simulated) represent major components of today's computing requirements in HEP. A recent projection [1] of the ATLAS 2016 computing model results in >85% of the HL-LHC CPU resources being spent on the reconstruction of data or simulated events. Several types of software algorithms are essential to the transformation of raw detector data into analysis-level objects. Specifically, these algorithms can be broadly grouped:

1. Online: Algorithms, or sequences of algorithms, executed on events read out from the detector in near-real-time as part of the software trigger, typically on a computing facility located close to the detector itself.
2. Offline: As distinguished from online, any algorithm or sequence of algorithms executed on the subset of events preselected by the trigger system, or generated by a Monte Carlo simulation application, typically in a distributed computing system.
3. Reconstruction: The transformation of raw detector information into higher level objects used in physics analysis. A defining characteristic of 'reconstruction' that separates it from 'analysis' is that the quality criteria used in the reconstruction to, for example, minimize the number of fake tracks, are independent of how those tracks will be used later on. This usually implies that reconstruction algorithms use the entirety of the detector information to attempt to create a full picture of each interaction in the detector. Reconstruction algorithms are also typically run as part of the processing carried out by centralized computing facilities.
4. Trigger: the online classification of events which reduces either the number of events which are kept for further 'offline' analysis, the size of such events, or both. Software triggers, whose defining characteristic is that they process data without a fixed latency, are part of the real-time processing path and must make decisions quickly enough to keep up with the incoming data, possibly using substantial disk buffers.
5. Real-time analysis: Data processing that goes beyond object reconstruction, and is performed online within the trigger system. The typical goal of real-time analysis is to combine the products of the reconstruction algorithms (tracks, clusters, jets...) into complex objects (hadrons, gauge bosons, new physics candidates...) which can then be used directly in analysis without an intermediate reconstruction step.

7.3.1 Challenges

Software trigger and event reconstruction techniques in HEP face a number of new challenges in the next decade. These are broadly categorized into 1) those from new and upgraded accelerator facilities, 2) from detector upgrades and new detector technologies, 3) increases in anticipated event rates to be processed by algorithms (both online and offline), and 4) from evolutions in software development practices.

Advances in facilities and future experiments bring a dramatic increase in physics reach, as well as increased event complexity and rates. At the HL-LHC, the central challenge for object reconstruction is thus to maintain excellent efficiency and resolution in the face of high pileup values, especially at low object p_T . Detector upgrades such as increases in channel density, high precision timing and improved detector geometric layouts are essential to mitigate these problems. For software, particularly for triggering and event reconstruction algorithms, there is a critical need not to dramatically increase the processing time per event.

A number of new detector concepts are proposed on the 5-10 year timescale to help in overcoming the challenges identified above. In many cases, these new technologies bring novel requirements to software trigger and event reconstruction algorithms or require new algorithms to be developed. Ones of particular importance at the HL-LHC include high-granularity calorimetry, high precision timing detectors, and hardware triggers based on tracking information which may seed later software trigger and reconstruction algorithms.

Trigger systems for next-generation experiments are evolving to be more capable, both in their ability to select a wider range of events of interest for the physics program of their experiment, and their ability to stream a larger rate of events for further processing. ATLAS and CMS both target systems where the output of the hardware trigger system is increased to $10\times$ the current capability, up to 1 MHz [68, 69]. In other cases, such as LHCb [70] and ALICE [71], the full collision rate (between 30 to 40 MHz for typical LHC operations) will be streamed to real-time or quasi-realtime software trigger systems starting in Run 3. The increase in event complexity also brings a ‘problem’ of overabundance of signal to the experiments, and specifically the software trigger algorithms. The evolution towards a genuine real-time analysis of data has been driven by the need to analyze more signal than can be written out for traditional processing, and technological developments which make it possible to do this without reducing the analysis sensitivity or introducing biases.

The evolution of computing technologies presents both opportunities and challenges. It is an opportunity to move beyond commodity x86 technologies, which HEP has used very effectively over the past 20 years, to performance-driven architectures and therefore software designs. It is also a significant challenge to derive sufficient event processing throughput per cost to reasonably enable our physics programs [72]. Specific items identified included 1) the increase of SIMD capabilities (processors capable of running a single instruction set simultaneously over multiple data), 2) the evolution towards multi- or many-core architectures, 3) the slow increase in memory bandwidth relative to CPU capabilities, 4) the rise of heterogeneous hardware, and 5) the possible evolution in facilities available to HEP production systems.

The move towards open source software development and continuous integration systems brings opportunities to assist developers of software trigger and event reconstruction algorithms. Continuous integration systems have already allowed automated code quality and performance checks, both for algorithm developers and code integration teams. Scaling these up to allow for sufficiently high statistics checks is still among the outstanding challenges. As the timescale for recording and analyzing data increases, maintaining and supporting legacy code will become more challenging. Code quality demands increase as traditional offline analysis components migrate into trigger systems, and, more generally, into algorithms that are run only once.

7.3.2 Current Approaches

Substantial computing facilities are in use for both online and offline event processing across all experiments surveyed. Online facilities are dedicated to the operation of the software trigger, while offline facilities are shared for operational needs including event reconstruction, simulation (often the dominant component) and analysis. CPU use by experiments is typically at the scale of tens or hundreds of thousands of x86 processing cores. The projections of CPU requirements discussed in Section 3 clearly demonstrate the need for either much larger facilities than anticipated or correspondingly more performant algorithms.

The CPU time needed for event reconstruction tends to be dominated by that used by charged particle reconstruction (tracking), especially as the need for efficiently reconstructing low p_T particles is considered. Calorimetric reconstruction, particle flow reconstruction and particle identification algorithms also make up significant parts of the CPU budget in some experiments.

Disk storage is currently 10s to 100s of PB per experiment. It is dominantly used to make the output of the event reconstruction, for both real data and simulated data, available for analysis.

Current generation experiments have moved towards smaller, but still flexible, data tiers for analysis. These tiers are typically based on the ROOT [51] file format and constructed to facilitate both skimming of interesting events and the selection of interesting pieces of events by individual analysis groups or through centralized analysis processing systems. Initial implementations of real-time analysis systems are in use within several experiments. These approaches remove the detector data that typically makes up the raw data tier kept for offline reconstruction, and to keep only final analysis objects [73–75].

Detector calibration and alignment requirements were surveyed. Generally a high level of automation is in place across experiments, both for very frequently updated measurements and more rarely updated measurements. Often, automated procedures are integrated as part of the data taking and data reconstruction processing chain. Some longer term measurements, requiring significant data samples to be analyzed together remain as critical pieces of calibration and alignment work. These techniques are often most critical for a subset of high precision measurements rather than for the entire physics program of an experiment.

7.3.3 Research and Development Roadmap and Goals

The CWP identified seven broad areas which will be critical for software trigger and event reconstruction work over the next decade. These are:

Roadmap area 1: Enhanced vectorization programming techniques - HEP=developed toolkits and algorithms typically make poor use of vector processors on commodity computing systems. Improving this will bring speedups to applications running on both current computing systems and most future architectures. The goal for work in this area is to evolve current toolkit and algorithm implementations, and best programming techniques to better use SIMD capabilities of current and future computing architectures.

Roadmap area 2: Algorithms and data structures to efficiently exploit many-core architectures - Computing platforms are generally evolving towards having more cores to increase processing capability. This evolution has resulted in multi-threaded frameworks in use, or in development, across HEP. Algorithm developers can improve throughput by being thread safe and enabling the use of fine-grained parallelism. The goal is to evolve current event models, toolkits and algorithm implementations, and best programming techniques to improve the throughput of multi-threaded software trigger and event reconstruction applications.

Roadmap area 3: Algorithms and data structures for non-x86 computing architectures (e.g. GPUs, FPGAs) - Computing architectures using technologies beyond CPUs offer an interesting alternative for increasing throughput of the most time consuming trigger or reconstruction algorithms. Such architectures (e.g. GPUs, FPGAs) could be easily integrated into dedicated

trigger or specialized reconstruction processing facilities (e.g. online computing farms). The goal is to demonstrate how the throughput of toolkits or algorithms can be improved through the use of new computing architectures in a production environment. The adoption of these technologies will particularly affect the research and development needed in other roadmap areas.

Roadmap area 4: Enhanced QA/QC for reconstruction techniques - HEP experiments have extensive continuous integration systems, including varying code regression checks that have enhanced the quality assurance (QA) and quality control (QC) procedures for software development in recent years. These are typically maintained by individual experiments and have not yet reached the scale where statistical regression, technical, and physics performance checks can be performed for each proposed software change. The goal is to enable the development, automation, and deployment of extended QA and QC tools and facilities for software trigger and event reconstruction algorithms.

Roadmap area 5: Real-time analysis - Real-time analysis techniques are being adopted to enable a wider range of physics signals to be saved by the trigger for final analysis. As rates increase, these techniques can become more important and widespread by enabling only the parts of an event associated with the signal candidates to be saved, reducing the required disk space. The goal is to evaluate and demonstrate the tools needed to facilitate real-time analysis techniques. Research topics include compression and custom data formats; toolkits for real-time detector calibration and validation which will enable full offline analysis chains to be ported into real-time; and frameworks which will enable non-expert offline analysts to design and deploy real-time analyses without compromising data taking quality.

Roadmap area 6: High precision physics-object reconstruction, identification and measurement techniques - The central challenge for object reconstruction at the HL-LHC is to maintain excellent efficiency and resolution in the face of high pileup values, especially for low object p_T . Both trigger and reconstruction algorithms must exploit new techniques and higher granularity detectors to maintain, or even improve, future physics measurements. It is also becoming clear that reconstruction in very high pileup environments at the HL-LHC will only be possible by adding timing information to our detectors. Designing appropriate detectors requires that the corresponding reconstruction algorithms be developed and demonstrated to work well in complex environments.

Roadmap area 7: Fast software trigger and reconstruction algorithms for high-density environments - Future experimental facilities will bring a large increase in event complexity. The scaling of current-generation algorithms with this complexity must be improved to avoid a large increase in resource needs. Where possible, toolkits and algorithms will be evolved or rewritten, focusing on their physics and technical performance at high event complexity (e.g. high pileup). It is likely also necessary to deploy new algorithms and new approaches, including advanced machine learning techniques developed in other fields, in order to solve these problems. One possible approach is that of anomaly detection, where events not consistent with known processes or signatures are identified and retained. The most important targets are those which limit expected throughput performance (e.g. charged-particle tracking).

7.3.4 Impact and Relevance for S^2I^2

Reconstruction algorithms are projected to be the biggest CPU consumers at the HL-LHC. Code modernization and new approaches are needed to address the large increases in pileup (4x) and trigger output rates (5-10x). Trigger/Reco algorithm enhancements (and new approaches) enable extended physics reach even in more challenging detection environments (e.g., pileup). Moreover, Trigger/Reco algorithm development is needed to take full advantage of enhanced detector capabilities (e.g., timing detectors, high-granularity calorimeters). ‘Real time analysis’ promises to effectively increase achievable trigger rates (for fixed budgets) through making reduced-size,

analysis-ready output from online trigger(-less) systems.

Physics Impact: Effectively selecting datasets to be persisted, and processing them sufficiently rapidly while maintaining the quality of the reconstructed objects, will allow analysts to use the higher collision rates in the more complex environments to address the broadest range of physics questions.

Resources Impact: Technical improvements achieved in trigger or reconstruction algorithms directly reduce the computing resources needed for HL-LHC computing. In addition, targeted optimizations of existing code will allow HL-LHC experiments to fully take advantage of the significant computing resources at HPC centers that may become available at little direct cost.

Sustainability Impact: University personnel, including graduate students and post-docs working in the research program, frequently develop and maintain trigger, reconstruction, and real-time analysis algorithms and implementations. Doing so in the context of an S^2I^2 will focus efforts on best practices related to reproducible research, including design and documentation.

Interest/Expertise: U.S. university groups are already leading many of the efforts in these areas. They are also working on designs of detector upgrades that require improved algorithms to take advantage of new features such as high precision timing. Similarly, they are already studying the use of more specialized chipsets, such as FPGAs and GPUs, for HEP-specific applications such as track pattern recognition and parameter estimation.

Leadership: As in the bullet above.

Value: All LHC experiments will benefit from these techniques, although detailed implementations will be experiment-specific given the differing detector configurations.

Research/Innovation: The CPU evolution requirements described in Section 3 are about $6\times$ greater than those promised by Moore's Law. Achieving this level of performance will require significant algorithmic innovation and software engineering research to take advantage of vector processors and other emerging technologies. Machine learning also promises the ability to replace some of the most CPU-intensive algorithms with fast inference engines trained on mixtures of simulated and real data. These efforts will require collaboration of domain experts with software engineers, computer scientists, and data scientists with complementary experience.

7.4 Applications of Machine Learning

Machine Learning (ML) is a rapidly evolving approach to characterizing and describing data with the potential to radically change how data is reduced and analyzed. Some applications will qualitatively improve the physics reach of data sets. Others will allow much more efficient use of processing and storage resources, effectively extending the physics reach of the HL-LHC experiments. Many of the activities in this focus area will explicitly overlap with those in the other focus areas. Some will be more generic. As a first approximation, the HEP community will build domain-specific applications on top of existing toolkits and ML algorithms developed by computer scientists, data scientists, and scientific software developers from outside the HEP world. HEP developers will also work with these communities to understand where some of our problems do not map onto existing paradigms well, and how these problems can be re-cast into abstract formulations of more general interest.

7.4.1 Opportunities

The world of data science has developed a variety of very powerful ML approaches for classification (using pre-defined categories), clustering (where categories are discovered), regression (to produce

continuous outputs), density estimation, dimensionality reduction, etc. Some have been used productively in HEP for more than 20 years; others have been introduced relatively recently. More are on their way. A key feature of these algorithms is that most have open software implementations that are reasonably well documented. HEP has been using ML algorithms to improve software performance in many types of software for more than 20 years, and ML has already become ubiquitous in some types of applications. For example, particle identification algorithms that require combining information from multiple detectors to provide a single figure of merit use a variety of BDTs and neural nets. With the advent of more powerful hardware and more performant ML algorithms, we want to use these tools to develop application software that could:

- replace the most computationally expensive parts of pattern recognition algorithms and algorithms that extract parameters characterizing reconstructed objects;
- compress data significantly with negligible loss of fidelity in terms of physics utility;
- extend the physics reach of experiments by qualitatively changing the types of analyses that can be done.

The abundance of ML algorithms and implementations presents both opportunities and challenges for HEP. Which are most appropriate for our use? What are the tradeoffs of one compared to another? What are the tradeoffs of using ML algorithms compared to using more traditional software? These issues are not necessarily factorizable, and a key goal of an Institute will be making sure that the lessons learned by one any research team are usefully disseminated to the greater HEP world. In general, the Institute will serve as a repository of expertise. Beyond the R&D projects it sponsors directly, the Institute will help teams develop and deploy experiment-specific ML-based algorithms in their software stacks. It will provide training to those developing new ML-based algorithms as well as those planning to use established ML tools.

7.4.2 Current Approaches

The use of ML in HEP analyses has become commonplace over the past two decades. Many analyses use the HEP-specific software package TMVA [28] included in the CERN ROOT [21] project. Recently, many HEP analysts have begun migrating to ML packages developed outside of HEP, such as SCIKIT-LEARN [76] and KERAS [77]. Data scientists at Yandex created a Python package that provides a consistent API to most ML packages used in HEP [78], and another that provides some HEP-specific ML algorithms [79]. Packages like SPEARMINT [80] perform Bayesian optimization and can improve HEP Monte Carlo [81, 82].

The keys to successfully using ML for any problem are:

- creating/identifying the optimal training, validation, and testing data samples;
- designing and selecting feature sets; and
- defining appropriate problem-specific loss functions.

While each experiment is likely to have different specific use cases, we expect that many of these will be sufficiently similar to each other that much of the research and development can be done commonly. We also expect that experience with one type of problem will provide insights into how to approach other types of problems.

7.4.3 Research and Development Roadmap and Goals

The following specific examples illustrate possible first-year activities.

- Charged track and vertex reconstruction is one of the most CPU intensive elements of the software stack. The algorithms are typically iterative, alternating between selecting hits associated with tracks and characterizing the trajectory of a track (a collection of hits). Similarly,

vertices are built from collections of tracks, and then characterized quantitatively. ML algorithms have been used extensively outside HEP to recognize, classify, and quantitatively describe objects. We will investigate how to replace components of the pattern recognition algorithms and the ‘fitting’ algorithms that extract parameters characterizing the reconstructed objects. As existing algorithms already produce high-quality physics, the primary goal of this activity will be developing replacement algorithms that execute much more quickly while maintaining sufficient fidelity.

- ML algorithms can often discover patterns and correlations more powerfully than human analysts alone. This allows qualitatively better analysis of recorded data sets. For example, ML algorithms can be used to characterize the substructure of “jets” observed in terms of underlying physics processes. ATLAS, CMS, and LHCb already use ML algorithms to separate jets into those associated with b-quark, c-quarks, or lighter quarks. ATLAS and CMS have begun to investigate whether sub-jets can be reliably associated with quarks or gluons. If this can be done with both good efficiency and accurate understanding of efficiency, the physics reach of the experiments will be radically extended .
- The ATLAS, CMS, and LHCb detectors all produce much more data than can be moved to permanent storage. The process of reducing the size of the data sets is referred to as the trigger. Electronics sparsify the data stream using zero suppression and they do some basic data compression. While this will reduce the data rate by a factor of 100 (or more, depending on the experiment) to about 1 terabyte per second, another factor of order 1500 is required before the data can be written to tape (or other long-term storage). ML algorithms have already been used very successfully to rapidly characterize which events should be selected for additional consideration and eventually persisted to long-term storage. The challenge will increase both quantitatively and qualitatively as the number of proton-proton collisions per bunch crossing increases.
- All HEP experiments rely on simulated data sets to accurately compare observed detector response data with expectations based on the hypotheses of the Standard Model or models of new physics. While the processes of subatomic particle interactions with matter are known with very good precision, computing detector response analytically is intractable. Instead, Monte Carlo simulation tools, such as GEANT4 [22–24], have been developed to simulate the propagation of particles in detectors. They accurately model trajectories of charged particles in magnetic fields, interactions and decays of particles as they traverse the fiducial volume, etc. Unfortunately, simulating the detector response of a single LHC proton-proton collision takes on the order of several minutes. *Fast simulation* replaces the slowest components of the simulation chain with computationally efficient approximations. Often, this is done using simplified parameterizations or look-up tables which don’t reproduce detector response with the required level of precision. A variety of ML tools, such as Generative Adversarial Networks and Variational Auto-encoders, promise better fidelity and comparable executions speeds (after training). For some of the experiments (ATLAS and LHCb), the CPU time necessary to generate simulated data will surpass the CPU time necessary to reconstruct the real data. The primary goal of this activity will be developing fast simulation algorithms that execute much more quickly than full simulation while maintaining sufficient fidelity.

7.4.4 Impact and Relevance for S^2I^2

Physics Impact: Software built on top of machine learning will provide the greatest gains in physics reach by providing new types of reconstructed object classification and by allowing triggers to more quickly and efficiently select events to be persisted.

Resources Impact: Replacing the most computationally expensive parts of reconstruction will

allow the experiments to use computing resources more efficiently. Optimizing data compression will allow the experiments to use data storage and networking resources more efficiently.

Sustainability Impact: Building our domain-specific software on top of ML tools from the larger scientific software community should reduce the need to maintain equivalent tools we built (or build) ourselves, but it will require that we help maintain the toolkits we use.

Interest/Expertise: U.S. university personnel are already leading significant efforts in using ML, from reconstruction and trigger software to tagging jet flavors to identifying jet substructures.

Leadership: There is a natural area for Institute leadership: in addition to the existing interest and expertise in the university HEP community, this is an area where engaging academics from other disciplines will be a critical element in making the greatest possible progress.

Value: All LHC experiments will benefit from using ML to write more performant software. Although specific software implementations of algorithms will differ, much of the R&D program can be common. Sharing insights and software elements will also be valuable.

Research/Innovation: ML is evolving very rapidly, so there are many opportunities for basic and applied research as well as innovation. As most of the work developing ML algorithms and implementing them in software (as distinct from the applications software built using them) is done by experts in the computer science and data science communities, HEP needs to learn how to effectively use toolkits provided by the open scientific software community. At the same time, some of the HL-LHC problems may be of special interest to these other communities, either because the sizes of our data sets are large (multi-exabyte) or because they have unique features.

7.5 Data Organization, Management and Access (DOMA)

Experimental HEP has long been a data intensive science and it will continue to be through the HL-LHC era. The success of HEP experiments is built on their ability to reduce the tremendous amounts of data produced by HEP detectors to physics measurements. The reach of these data-intensive experiments is limited by how quickly data can be accessed and digested by the computational resources. Both changes in technology and large increases in data volume require new computational models [12]. HL-LHC and the HEP experiments of the 2020s will be no exception.

Extending the current data handling methods and methodologies is expected to be intractable in the HL-LHC era. The development and adoption of new data analysis paradigms gives the field, as a whole, a window in which to adapt our data access and data management schemes to ones which are more suited and optimally matched to a wide range of advanced computing models and analysis applications. This type of shift has the potential for enabling new analysis methods and allowing for an increase in scientific output.

7.5.1 Challenges and Opportunities

The LHC experiments currently provision and manage about an exabyte of storage, approximately half of which is archival, and half is traditional disk storage. The storage requirements per year are expected to jump by a factor of 10 for the HL-LHC. This itself is faster than projected Moore's Law gains and will present major challenges. Storage will remain one of the visible cost drivers for HEP computing, however the projected growth and cost of the computational resources needed to analyze the data is also expected to grow even faster than the base storage costs. The combination of storage and analysis computing costs may restrict scientific output and potential physics reach of the experiments, thus new techniques and algorithms are likely to be required.

These three main challenges for data in the HL-LHC era can thus be summarized:

1. **Big Data:** the HL-LHC will bring significant increases to both the data rate and the data volume. The computing systems will need to handle this without significant cost increases and within evolving storage technology limitations.
2. **Dynamic Distributed Computing:** In addition, the significantly increased computational requirements for the HL-LHC era will also place new requirements on data. Specifically the use of new types of compute resources (cloud, HPC, and hybrid platforms) with different dynamic availability and characteristics will require more dynamic DDM systems.
3. **New Applications:** New applications such as machine learning training or high rate data query systems for analysis will likely be employed to meet the computational constraints and to extend the physics reach of the HL-LHC. These new applications will place new requirements on how and where data is accessed and produced. For example, specific applications (e.g. training for machine learning) may require use of specialized processor resources such as GPUs, placing further requirements on data formats.

The projected event complexity of data from future LHC runs will require advanced reconstruction algorithms and analysis tools. The precursors of these tools, in the form of new machine learning paradigms, pattern recognition algorithms, and fast simulations, already show promise in reducing CPU needs for HEP experiments. As these techniques continue to grow and blossom, they will place new requirements on the computational resources that need to be leveraged by all of HEP. The storage systems that are developed, and the data management techniques that are employed will need to directly support this wide range of computational facilities, and will need to be matched to the changes in the computational work, so as not to impede the improvements that they are bringing.

As with CPU, the landscape of storage protocols accessible to us is trending towards heterogeneity. Thus, the ability to leverage new storage technologies as they become available into existing data delivery models becomes a challenge that we must be prepared for.

In part, this also means HEP experiments should be prepared to separate storage abstractions from the underlying storage resource systems [83]. Like opportunistic CPU, opportunistic or storage resources available for limited duration (e.g. from a cloud provider) require data management and provisioning systems that can exploit such resources on short notice. Much of this change can be aided by active R&D into our own IO patterns, which are yet to be fully studied and understood in HEP.

On the hardware side, R&D is needed in alternative approaches to data archiving to determine the possible cost/performance tradeoffs. Currently, tape is extensively used to hold data that cannot be economically made available online. While the data is still accessible, it comes with a high latency penalty, limiting possible analysis. We suggest investigating either separate direct access-based archives (e.g. disk or optical) or new models that overlay online direct access volumes with archive space. This is especially relevant when access latency is proportional to storage density.

Closely related is the choice of what data objects should always be retrieved together (i.e. the “atomic size”) which has implications at all levels in the software, storage and network infrastructure. In storage systems, as the atomic size increases so does memory pressure, CPU cycles spent on copying/moving data, and the likelihood of hot spots on data servers. As atomic size decreases, the CPU cycles spent on requesting data, round-trip times, and metadata overhead increase while locality is reduced. Luckily, modern storage systems like Ceph have a number of effective knobs to navigate these trade-offs, including sizing of objects, partitioning and striping of data to objects, and co-location of objects. However, these currently must be manually tuned for the workflow being optimized. Research in automating and “learning” which sets of storage system parameters yield optimal access performance is needed.

In the end, the results have to be weighed against the storage deployment models that, currently, differ among the various experiments. In the near term, this offers an opportunity to evaluate the effectiveness of chosen approaches at scale. The lessons drawn will provide guidance going forward into the HL-LHC era.

While our focus is convergence within the LHC community we do not want to imply that efforts to broaden that convergence to include non-LHC experiments should not be pursued. Indeed, as the applicable community increases, costs are typically driven lower. and sustainability of the devised solutions increases. This needs to be explored as it is not clear to what extent LHC-focused solutions can be used in other communities that ostensibly have different cultures, processing needs, and even funding models. We should caution that making any system cover an ever wider range of requirements inevitably leads to more complex solutions that are difficult to maintain and while they perform well on average they rarely perform well for any specific use.

Finally, any and all changes undertaken must not make the ease of access to data any worse than it is under current computing models. We must also be prepared to accept the fact that the best possible solution may require significant changes in the way data is handled and analyzed. What is clear is that what is being done today will not scale to the needs of HL LHC [84].

7.5.2 Current Approaches

The original LHC computing models (circa 2005) were built up from the simpler models used before distributed computing was a central part of HEP computing. This allowed for a reasonably clean separation between three different aspects of interacting with data: organization, management and access.

Data Organization: This is essentially how data is structured as it is written. Most data is written in flat files, in ROOT [51] format, typically with a column-wise organization of the data. The records corresponding to these columns are compressed. The internal details of this organization are typically visible only to individual software applications.

Data Management: The key challenge here was the transition to the use of distributed computing in the form of the grid. The experiments developed dedicated data transfer and placement systems, along with catalogs, to move data between computing centers. To first order the computing models were rather static: data was placed at sites and the relevant compute jobs were sent to the right locations. Applications might interact with catalogs or, at times, the workflow management systems does this on behalf of the applications.

Data Access: Various protocols are used for direct reads (rfio, dcap, xrootd, etc.) with a given computer center and/or explicit local stagein and caching for read by jobs. Application access may use different protocols than those used by the data transfers between site.

Before the LHC turn-on and in the first years of the LHC, these three areas were to first order optimized independently. Many of the challenges were in the area of “Data Management (DM)” as the Worldwide LHC Computing Grid was commissioned. As the LHC computing matured through Run 1 and Run 2, the interest has turned to optimizations spanning these three areas. For example, the recent use of “Data Federations” [85, 86] mixes up the Data Management and Data Access aspects. As we will see below, some of the foreseen opportunities towards HL-LHC may require global optimizations.

Thus in this document we take a broader view than traditional “DM”, and consider the combination of “Data Organization, Management and Access (DOMA)” together. We believe that by treating this area as a this full picture of data needs in HEP will provide important opportunities for efficiency and scalability as we enter the many-Exabyte era.

7.5.3 Research and Development Roadmap and Goals

In the following, we identify task areas and goals to demonstrate that the increased volume and complexity of data expected over the coming decade can be stored, accessed, and analyzed at an affordable cost.

Atomic Size of Data: An important goal is to create abstractions that make questions like the atomic size of data go away because that size is determined automatically. In higher layers of abstraction, we generally mean sub-file granularity, e.g. event-based. This should be studied to see whether it can be implemented efficiently and in a scalable, cost-effective manner. Applications making use of event selection can be assessed as to whether it offers an advantage over current file-based granularity. The following tasks should be completed by 2020:

- Quantify the impact on performance and resource utilization (CPU, storage, network) for the main type of access patterns (simulation, reconstruction, analysis).
- Assess impact of different access patterns on catalogs and data distribution.
- Assess whether event-granularity makes sense in object stores that tend to require large chunks of data for efficiency.
- Test for improvement in recoverability from preemption, in particular when using cloud spot resources and/or dynamic HPC resources.

Data Organization Paradigms: We will seek to derive benefits from data organization and analysis technologies adopted by other big data users. A proof-of-concept that involves the following tasks needs to be established by 2020 to allow full implementations to be made in the years that follow.

- Study the impact of column-wise, versus row-wise, organization of data on the performance of each kind of access, including the associated storage format.
- Investigate efficient data storage and access solutions that support the use of map-reduce or Spark-like analysis services.
- Evaluation of declarative interfaces and in-situ processing.
- Evaluate just-in-time decompression schemes and mappings onto hardware architectures considering the flow of data, from spinning disk to memory and application.
- Investigate the long term replacement of gridftp as the primary data transfer protocol. Define metrics (performance, etc.) for evaluation.
- Benchmark end-end data delivery for the main use cases (reco, MC, various analysis workloads, etc.), what are the impediments to efficient data delivery to the CPU to and from (remote) storage? What are the necessary storage hierarchies, and how does that map into technologies foreseen?

Data Placement and Caching: Discover the role that data placement optimizations can play, including the use of data caches distributed over the WAN, in order to use computing resources more effectively. Investigate and or develop the technologies that can be used to achieve this. The following tasks should be completed by 2020:

- Quantify the benefit of placement optimization for the main use cases i.e. reconstruction, analysis, and simulation.
- Assess the benefit of caching for machine learning-based applications (in particular for the learning phase) and follow-up the evolution of technology outside HEP itself.
- Assess the potential benefit of content delivery networks in the HEP data context.

- Assess the feasibility and potential benefit of a named data network component in a HL-LHC data management system, in both medium and long-term as this new technology matures [87].

Federated Data Centers (a prototype “Data-Lake”)

- Understanding the needed functionalities, including policies for managing data and replications, availability, quality of service, service levels, etc.;
- Understand how to interface a data-lake federation with heterogeneous storage systems in different sites
- Investigate how to define and manage the interconnects, network performance and bandwidth, monitoring, service quality etc. Integration of networking information and testing of advanced networking infrastructure.
- Investigate policies for managing and serving derived data sets, lifetimes, re-creation (on-demand?), caching of data, etc.

In the longer term, the benefits that can be derived from using different approaches to the way HEP is currently managing its data delivery systems should be investigated. Different content delivery methods should be studied for their utility in the HL-LHC context, namely Content Delivery Networks (CDN) and Named Data Networking (NDN).

Support for Query-based analysis techniques: Data analysis is currently tied in with ROOT-based formats. In many currently-used paradigms, physicists consider all events at an equivalent level of detail and in the format offering the highest level of detail that needs to be considered in an analysis. However, not every event considered in analysis requires the same level of detail. One consideration to improve event access throughput is to design event tiers with different abstractions, and thus data sizes. All events can be considered at a lighter-weight tier while events of interest only can be accessed with a more information-rich tier.

For more scalable analysis, another opportunity to evaluate is how much work can be offloaded to a storage system, for example caching uncompressed or reordered data for fast access. The idea can be extended to virtual data and to query interfaces which would perform some of the transformation logic currently executed on CPU workers. Interactive querying of large datasets is an active field in the Big Data industry; examples include Spark-SQL, Impala, Kudu, Hawq, Apache Drill, and Google Dremel/BigQuery. A key question is about the usability of these techniques in HEP and we need to assess if our data transformations are not too complex for the SQL-based query languages used by these products. We also need to take into account that the adoption of these techniques, if they prove to be beneficial, would represent a disruptive change which directly impacts the end user and therefore promoting acceptance through intermediate solutions would be desirable.

7.5.4 Impact and Relevance for S^2I^2

Physics Impact: The very fast turnaround of analysis results that could be possible with new approaches to data access and organization would lead to rapid turnaround for new science.

Resources Impact: Optimized data access will lead to more efficient use of resources. In addition, by changing the analysis models, and by reducing the number of data replicas required, the overall costs of storage can be reduced.

Sustainability Impact: This effort would improve the reproducibility and provenance tracking for workflows (especially analysis workflows), making physics analyses more sustainable through the lifetime of the HL-LHC.

Interest/Expertise: University groups have already pioneered significant changes to the data access model for the LHC through the development of federated storage systems, and are prepared to take this further. Other groups are currently exploring the features of modern storage systems and their possible implementation in experiments.

Leadership: CS research and technology innovation in several pertinent areas are being carried out by university groups, including research on methods for large scale adaptive and elastic database systems that support intensive mixed workloads (e.g. high data ingest, online analytics, and transactional updates). Also universities are leading centers for work addressing critical emerging problems across many science domains, including analytical systems that benefit from column-oriented storage, where data is organized by attributes instead of records, thus enabling efficient disk I/O and which support directly querying over losslessly compressed data. As many teams perform data analytics in a collaborative way, where several users contribute to cleaning, modeling, analyzing, and integrating new data. To allow users to work on these tasks in isolation and selectively share the results, research at universities is actively developing systems to support lightweight dataset versioning, that is similar to software control systems like Git, but for structured data.

Value: All LHC experiments will benefit from new methods of data access and organization, although the implementations may vary due to the different data formats and computing models of each experiment.

Research/Innovation: This effort would rely on partnerships with data storage and access experts in the CS community, some of whom are already providing consultation in this area.

7.6 Fabric of distributed high-throughput computing services (OSG)

Since its inception, the Open Science Grid (OSG) has evolved into an internationally-recognized element of the U.S. national cyberinfrastructure, enabling scientific discovery across a broad range of disciplines. This has been accomplished by a unique partnership that cuts across science disciplines, technical expertise, and institutions. Building on novel software and shared hardware capabilities, the OSG has been expanding the reach of high-throughput computing (HTC) to a growing number of communities. Most importantly, in terms of the HL-LHC, it provides essential services to US-ATLAS and US-CMS.

The importance of the fabric of distributed high-throughput computing (DHTC) services was identified by the National Academies of Science (NAS) 2016 report on NSF Advanced Computing Infrastructure: *Increased advanced computing capability has historically enabled new science, and many fields today rely on high-throughput computing for discovery* [88]. HEP in general, and the HL-LHC science program in particular, already relies on DHTC for discovery; we expect this to become even more true in the future. While we will continue to use existing facilities for HTC, and similar future resources, we must be prepared to take advantage of new methods for accessing both “traditional” and newer types of resources.

The OSG provides the infrastructure for accessing all different types of resources as transparently as possible. Traditional HTC resources include dedicated facilities at national laboratories and universities. The LHC is also beginning to use allocations at national HPC facilities, (e.g., NSF- and DOE- funded leadership class computing centers) and elastic, on-demand access to commercial clouds. It is sharing facilities with collaborating institutions in the wider national and international community. Moving beyond traditional, single-threaded applications running on x86 architectures, the HEP community is writing software to take advantage of emerging architectures. These include vectorized versions of x86 architectures (including Xeon, KNL and AMD) and various types of GPU-based accelerator computing. The types of resources being requested are becoming more varied in other ways. Deep learning is currently most efficient on specialized GPUs and similar

architectures. Containers are being used to run software reliably and reproducibly moving from one computing environment to another. Providing the software and operations infrastructure to access scalable, elastic, and heterogeneous resources is an essential challenge for LHC and HL-LHC computing and the OSG is helping to address that challenge.

The software and computing leaders of the U.S. LHC Operations Program, together with input from the OSG Executive Team, have defined a minimal set of services needed for the next several years. These services and their expected continued FTE levels are listed in Table 2 below. They are orthogonal to the S^2I^2 R&D program for HL-LHC era software, including prototyping. Their focus is on operating the currently needed services. They include R&D and prototyping only to the extent that this is essential to support the software lifecycle of the distributed DHTC infrastructure. The types of operations services supported by the OSG for US-LHC fall into six categories, plus coordination.

Category	ATLAS-only	Shared ATLAS and CMS	CMS only	Total
Infrastructure software maintenance and integration	0.85	2.9	1.7	5.45
CVMFS service operation	0.2	0.1	0.4	0.7
Accounting, registration, monitoring	0.35	0.3	0.2	0.85
Job submission infrastructure operations	1.5	0.0	1.0	2.5
Cybersecurity infrastructure	0.0	0.3	0.0	0.3
Ticketing and front-line support	1.0	1.2	1.0	3.2
Coordination	0.0	0.5	0.0	0.5
Services Total	3.9	5.2	4.2	13.3
Technology evaluation		3.0		3.0

Table 2: OSG LHC Services (in FTEs), organized into six categories that are described in the text. Also shown at the bottom is the FTE level for the OSG technology evaluation area.

Infrastructure software maintenance and integration includes creating, maintaining, and supporting an integrated software stack that is used to deploy production services at compute and storage clusters that support the HL-LHC science program in the U.S. and South America. The entire software lifecycle needs to be supported, from introducing a new product into the stack, to including updated versions in future releases that are fully integrated with all other relevant software to build production services, to retirement of software from the stack. The retirement process typically includes a multi-year “orphanage” during which OSG has to assume responsibility for a software package between the time the original developer abandons support for it, and the time it can be retired from the integrated stack. This is because the software has been replaced with a different product or is otherwise no longer needed.

CVMFS service operations includes operating three types of software library infrastructures. Those that are specific to the two experiments, and the one that both experiments share. As the bulk of the application level software presently is not shared between the experiments, the effort for the shared instance is smallest in Table 2. The shared service instance is also shared with most, but not all other user communities on OSG.

Accounting, registration, and monitoring includes any and all production services that allow

U.S. institutions to contribute resources to WLCG.

Job submission infrastructure operations is presently not shared between ATLAS and CMS because both have chosen radically different solutions. CMS shares its job submission infrastructure with all other communities on OSG, while ATLAS uses its own set of dedicated services. Both types of services need to be operated.

Cybersecurity infrastructure US-ATLAS and US-CMS depend on a shared cybersecurity infrastructure that includes software and processes, as well as a shared *coordination with the Worldwide LHC Computing Grid (WLCG)*. Both of these are also shared with all other communities on OSG.

Ticketing and front-line support The OSG operates a ticketing system to provide support for users and individual sites, including feature requests and handling issues related to security, wide-area networking, and installation and configuration of the software. The OSG also actively tracks and pushes to resolution issues reported by the WLCG community by synchronizing their respective problem ticket systems.

Technology evaluation In addition to these production services, the OSG presently includes a *technology evaluation* area that comprises 3 FTE. This area provides OSG with a mechanism for medium- to long-term technology evaluation, planning and evolution of the OSG software stack. It includes a blueprint activity that OSG uses to engage with computer scientists on longer term architectural discussions that sometimes lead to new projects that address functionality or performance gaps in the software stack. Given the planned role of the S^2I^2 as an intellectual hub for software and computing (see Section 6), it could be natural for this part of the current OSG activities to reside within a new Institute. Given the operational nature of the remainder of current OSG activities, and their focus on the present and the near future, it may be more appropriate for the remaining 13.3 FTE to be housed in an independent *but* collaborating project.

The full scope of whatever project houses OSG-like services in support of the LHC experiments moving forward, in terms of domain sciences, remains ill-defined. The OSG project has demonstrated that a single organization with users that span many different domains and experiments provides a valuable set of synergies and cross-fertilization of tools, technologies and ideas. The DHTC paradigm serves science communities beyond the LHC experiments, communities even more diverse than those of HEP. As clearly identified in the NAS NSF Advanced Computing Infrastructure report [88], *many fields today rely on high-throughput computing for discovery*. We encourage the NSF to develop a funding mechanism to deploy and maintain a common DHTC infrastructure for HL-LHC as well as LIGO, DES, IceCube, and other current and future science programs.

7.7 Backbone for Sustainable Software

In addition to enabling technical advances, the Institute must also focus on how these software advances are communicated to and taken up by students, researchers developing software (both within the HEP experiments and outside), and members of the general public with scientific interests in HEP and big data and software. The Institute will play a central role in elevating the recognition of software as a critical research cyberinfrastructure within the HEP community and beyond. To enable this elevation, we envision a “backbone” activity of the Institute that focuses on finding, improving, and disseminating best practices; determining and applying incentives around software; developing, coordinating and providing training; and making data and tools accessible by and useful to the public.

The experimental HEP community is unique in that the organization of its researchers into very large experiments results in significant community structure on a global scale. It is possible within this structure to explore the impact of changes to the software development processes with concrete metrics, as much of the software development is an open part of the collaborative process. This

1712 makes it a fertile ground both for study and for concretely exploring the nature and impact of best
 1713 practices. This large community also provides the infrastructure to conduct surveys and interviews
 1714 of project personnel to supplement the metrics with subjective and qualitative evaluations of the
 1715 need for and the changes to software process. An Institute Backbone for Sustainable Software,
 1716 with a mandate to pursue these activities broadly within and beyond the HEP community, would
 1717 be well placed to leverage this community structure.

1718 **Best Practices:** The Institute should document, disseminate, and work towards community adop-
 1719 tion of the best practices (from HEP and beyond) in the areas of software sustainability, including
 1720 topics in software engineering, data/software preservation, and reproducibility. Of particular im-
 1721 portance are best practices surrounding the modernization of the software development process
 1722 for scientists. Individual experts can improve the technical performance of software significantly
 1723 (sometimes by more than an order of magnitude) by understanding the algorithms and intended
 1724 optimizations and providing advice on how to achieve the best performance. Surveys and interviews
 1725 of HEP scientists can provide the information need to elicit and document best practices as well as
 1726 to identify the area still in need of improvement. The Institute can improve the overall process so
 1727 that the quality of software written by the original scientist author is already optimized. In some
 1728 cases tool support, including packaging and distribution, may be an integral part of the best
 1729 practices. Best practices should also include the use of testbeds for validation and scaling. This is
 1730 a natural area for collaboration between the Institute and the LHC Ops programs: the Institute
 1731 can provide the effort for R&D and capabilities while the Ops programs can provide the actual
 1732 hardware testbeds. The practices can be disseminated in general outreach to the HEP software
 1733 development community and integrated into training activities. The Backbone can also engage in
 1734 planning exercises and modest, collaborative efforts with the experiments to lower the barrier to
 1735 adoption of these practices.

1736 The Institute should also leverage the experience of the wider research community interested
 1737 in sustainable software issues, including the NSF SI2 community and other S^2I^2 institutes, the
 1738 Software Sustainability Institute in the UK [89], the HPC centers, industry, open source software
 1739 communities, and other organizations and adopt this experience for the HEP community. It should
 1740 also collaborate with empirical software engineers and external experts to (a) study HEP processes
 1741 and suggest changes and improvements and (b) develop activities to deploy and study the imple-
 1742 mentation of these best practices in the HEP community. These external collaborations may involve
 1743 a combination of unfunded collaborations, official partnerships, (funded) Institute activities, and
 1744 potentially even the pursuit of dedicated proposals and projects. The Institute should provide the
 1745 fertile ground in which all of these possibilities can grow.

1746 **Incentives:** The Institute should also play a role in developing incentives within the HEP commu-
 1747 nity for (a) sharing software and for having your software used (in discoveries, by others building
 1748 off it), (b) implementing best practices (as above) and (c) valuing research software development
 1749 as a career path. This may include defining metrics regarding HEP research software (including
 1750 metrics related to software productivity and scientific productivity) and publicizing them within
 1751 the HEP community. It could involve the use of blogs, webinars, talks at conferences, or dedicated
 1752 workshops to raise awareness and to publicize useful software development practices used within
 1753 the institute. Most importantly, the Institute can advocate for use of these metrics in hiring, pro-
 1754 motion, and tenure decisions at Universities and laboratories. To support this, the Institute should
 1755 create sample language and circulate these to departments and to relevant individuals.

1756 8 Institute Organizational Structure and Evolutionary Process

1757 During the S^2I^2 conceptualization process, the U.S. community had a number of discussions re-
 1758 garding possible management and governance structures. In order to organize these discussions,

it was agreed that the management and governance structures chosen for the Institute should be guided by answers the following questions:

1. **Goals:** What are the goals of the Institute?
2. **Interactions:** Who are the primary clients/beneficiaries of the Institute? How are their interests represented? How can the Institute align its priorities with those of the LHC experiments?
3. **Operations:** How does the Institute execute its plan with the resources it directly controls? How does the Institute leverage and collaborate with other organizations? How does the Institute maintain transparency?
4. **Metrics:** How is the impact of the Institute evaluated? And by whom?
5. **Evolution:** What are the processes by which the Institutes areas of focus and activities evolve?

The S^2I^2 discussions converged on the strawman model described show in Figure 8 as a baseline. The specific choices may evolve in an eventual implementation phase depending on funding levels, specific project participants, etc., but the basic functions here are expected to be relevant and important.

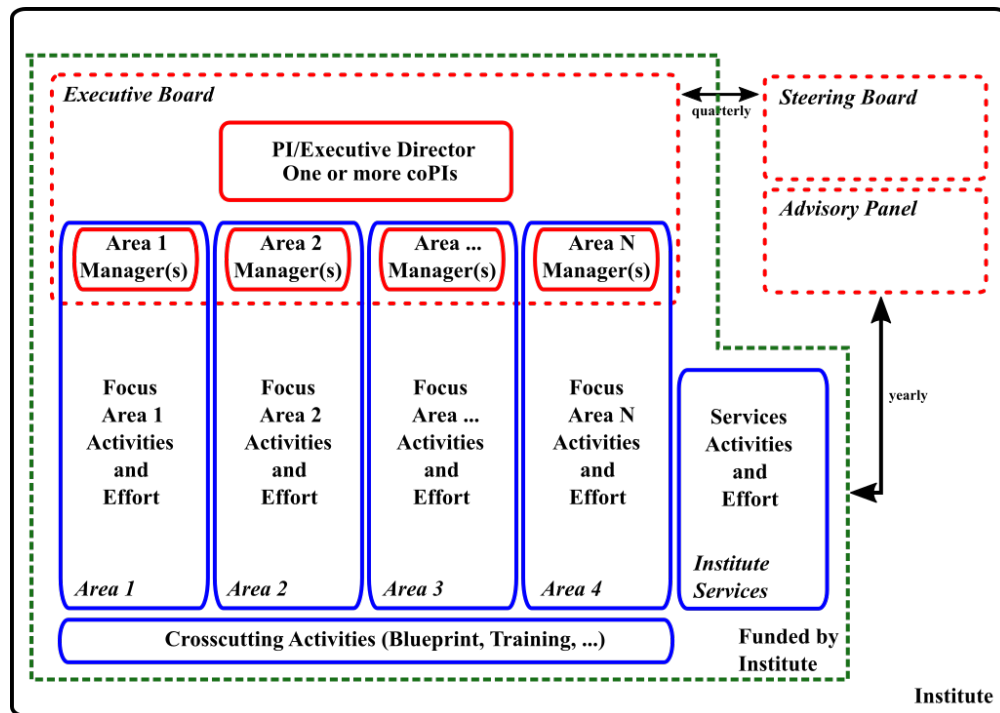


Figure 8: Strawman Model for Institute Management and Governance. (Figure to be remade!)

The main elements in this organizational structure and their roles within the Institute are:

PI/co-PIs: The PI/co-PIs on an eventual Institute implementation proposal will have project responsibilities as defined by NSF.

Focus Areas: A number of Focus Areas will be defined for the institute at any given point in time. These areas will represent the main priorities of the institute in terms of activities aimed at developing the software infrastructure to achieve the mission of the Institute. The S^2I^2 -HEP conceptualization process has identified a initial set of high impact focus areas. These are described in Section 7 of this document. The number and size of focus areas which will be included in an Institute implementation will depend on funding available and resources needed to achieve the goals. The areas could also evolve over the course of the institute, but it is expected to be typically between three and five. Each focus area within an Institute will have a written set of goals for the year and corresponding institute resources. The active focus areas will be reviewed together with the Advisory Panel once/year and decisions will be taken on updating the list of areas and their yearly goals, with input from the Steering Board.

Area Manager(s): Each Area Manager will manage the day to day activities within a focus area. It is for the moment undefined whether there will be an Area Manager plus a deputy, co-managers or a single manager. An appropriate mix of HEP, Computer Science and representation from different experiments will be a goal.

Executive Board: The Executive Board will manage the day to day activities of the Institute. It will consist of the PI, co-PIs, and the managers of the focus areas. A weekly meeting will be used to manage the general activities of the institute and make shorter term plans. In many cases, a liaison from other organizations (e.g. the US LHC Ops programs) would be invited as an “observer” to weekly Executive Board meetings in order to facilitate transparency and collaboration (e.g. on shared services or resources).

Steering Board: A Steering Board will be defined to meet with the executive board approximately quarterly to review the large scale priorities and strategy of the institute. (Areas of focus will also be reviewed, but less frequently.) The steering board will consist of two representatives for each participating experiment, representatives of the US-LHC Operations programs, plus representatives of CERN, FNAL, etc. Members of the Steering Board will be proposed by their respective organizations and accepted by the Executive Director in consultation with the Executive Board.

Executive Director: An Executive Director will manage the overall activities of the institute and its interactions with external entities. In general day-to-day decisions will be taken by achieving consensus in the Executive Board and strategy and priority decisions based on advice and recommendations by the Steering and Executive Boards. In cases where consensus cannot be reached, the Executive Director will take a final decision. It would also be prudent for the Institute to have a Deputy Director who is able to assume the duties during periods of unavailability of the Executive Director.

Advisory Panel: An Advisory Panel will be convened to conduct an internal review of the project once per year. The members of the panel will be selected by the PI/co-PIs with input from the Steering Board. The panel will include experts not otherwise involved with the institute in the areas of physics, computational physics, sustainable software development and computer science.

9 Building Partnerships

9.1 Partners

The roles envisioned for the Institute in Section 6 will require collaborations and partnerships with external entities, as illustrated in Figure 9. These include:

HEP Researchers (University, Lab, International): LHC researchers are the primary repository of expertise related to all of the domain-specific software to be developed and deployed; they




Figure 9: Partners of the Institute.

also define many of the goals for domain-specific implementations of more general types of software such as workflow management. Areas in which collaboration with HEP researchers will be especially close include technical aspects of the detectors and their performance, algorithms for reconstruction and simulation, analysis techniques and, ultimately, the potential physics reach of the experiments. These researchers will define the detailed and evolving physics goals of the experiments. They will participate in many of the roles described in Section 8, and some will be co-funded by the Institute. In addition, the Institute should identify, engage and build collaborations with other non-LHC HEP researchers whose interests and expertise align with the Institute's focus areas.

LHC Experiments: The LHC experiments, and especially the US-ATLAS and US-CMS collaborations, are key partners. In large measure, the success of an Institute will be judged in terms of its impact on their Computing Technical Design Reports (CTDRs), to be submitted in 2020, and its impact on software deployed (at least as a test for HL-LHC) during LHC's Run 3. The experiments will play leading roles in understanding and defining software requirements and how the pieces fit together into coherent computing models with maximum impact on cost/resources, physics reach, and sustainability. As described in Section 8, representatives of the experiments will participate explicitly in the Institute Steering Board to help provide big-picture guidance and oversight. In terms of daily work, the engagement will be deeper. Many people directly supported by the Institute will be collaborators on the LHC experiments, and some will have complementary roles in the physics or software & computing organizations of their experiments. Building on these natural connections will provide visibility for Institute activities within the LHC experiments, foster collaboration across experiments, and provide a feedback mechanism from the experiments to the Institute at the level of individual researchers. The experiments will also be integral to developing sustainability paths for software products they deploy that emerge from R&D performed by the Institute; therefore, they must be partners starting early in the software lifecycle.

US LHC Operations Programs: As described in Section 6, the Institute will be an R&D engine in the earlier phases of the software life cycle. The Operations Programs will be one of the primary

partners within the U.S. for integration activities, testing “at-scale” on real facilities, and eventual deployment. In addition they will provide a long run sustainability path for some elements of the software products. Ultimately, much of the software emerging from Institute efforts will be essential for the LHC Operations Programs or run in facilities they operate. The Institute will address many of the issues that the Operations Programs expect to encounter in the HL-LHC era. Thus, the Institute must have, within the U.S., a close relationship to the Operations Programs. Their representatives will serve as members of the Steering Board, and they will be invited to participate in Executive Board meetings as observers.

Computer Science (CS) Community: During the S^2I^2 -HEP conceptualization process we ran two workshops focused on how the HEP and CS communities could work to their mutual benefit in the context of an Institute, and, also, more generally. We identified some specific areas for collaboration, and others where the work in one field can inform the other. Several joint efforts have started as results of these conversations. More importantly, we discussed the challenges, as well as the opportunities, in such collaborations, and established a framework for continued exchanges. For example, we discussed the fact that the computer science research interest in a problem is often to map specific concrete problems to more abstract problems whose solutions are of research interest, as opposed to simply providing software engineering solutions to the concrete problems. This can nonetheless bring intellectual rigor and new points of view to the resolution of the specific HEP problem, and the HEP domain can provide realistic environments for exploring CS solutions at-scale, but it is very important to keep in mind the differing incentives of the two communities for collaboration. We anticipate direct CS participation in preparing a proposal if there is a solicitation, and collaboration in Institute R&D projects if it comes to fruition. Continued dialogue and engagement with the CS community will help assure the success of the Institute. This may take the form of targeted workshops focused on specific software and computing issues in HEP and their relevance for CS, or involvement of CS researchers in blueprint activities (see below). It may also take the form of joint exploratory projects. Identified topics of common interest include: science practices & policies, sociology and community issues; machine learning; software life cycle; software engineering; parallelism and performance on modern processor architectures, software/data/workflow preservation & reproducibility, scalable platforms; data organization, management and access; data storage; data intensive analysis tools and techniques; visualization; data streaming; training and education; and professional development and advancement. We also expect that one or two members of the CS and Cyberinfrastructure communities will serve on the Institute Advisory Panel, as described in Section 8, to provide a broad view of CS research.

External Software Providers: The LHC experiments depend on numerous software packages developed by external providers, both within HEP and from the wider open source software community. For the non-HEP software packages, the HEP community interactions are often a bit diffuse and unorganized. The Institute could play a role in developing the collaborations with these software providers, as needed, including engaging them for relevant planning, discussions regarding interoperability with other HEP packages, and software packaging and performance issues. For non-HEP packages the Institute can also play a role in introducing key developers of these external software packages to the HEP community. This can be done through invited seminars or sponsored visits to work at HEP institutions or by raising the visibility of HEP use cases in the external development communities. Examples of these types of activity can be found in the “topical meetings” being organized by the DIANA-HEP project [90,91].

Open Science Grid (OSG): The strength of the OSG project is its fabric of services that allows the integration of an at-scale globally distributed computing infrastructure for HTC that is fundamentally elastic in nature, and can scale out across many different types of hardware, software, and business models. It is the natural partner for the Institute on all aspects of “productizing”

1896 prototypes, or testing prototypes at scale. For example, the OSG already supports machine learning
1897 environments across a range of hardware and software environments. New environments could be
1898 added in support of the ML focus area. It is also a natural partner to facilitate discussions with
1899 IT infrastructure providers, and deployment experts, e.g. in the context of the DOMA and Data
1900 Analysis Systems focus areas. Because of its strong connections to the computer science community,
1901 the OSG also may also provide opportunities for engaging computer scientists (as described above)
1902 in other areas of interest to the Institute.

1903 **DOE and the National Labs:** The R&D roadmap outlined in the Community White Paper [13] is
1904 much broader than what will be possible within an Institute. The DOE labs will necessarily engage
1905 in related R&D activities both for the HL-LHC and for the broader U.S. HEP program in the 2020s.
1906 Many DOE lab personnel participated in both the CWP and S^2I^2 -HEP processes. In addition,
1907 a dedicated workshop was held in November 2017 to discuss how S^2I^2 - and DOE-funded efforts
1908 related to HL-LHC upgrade software R&D might be aligned to provide for maximum coherence
1909 (see Appendix B). Collaborations between university personnel and national laboratory personnel
1910 will be critical, as will be collaborations with foreign partners. In particular, the HEP Center for
1911 Computational Excellence (HEP-CCE) [92], a DOE cross-cutting initiative focused on preparations
1912 for effectively utilizing DOE’s future high performance computing (HPC) facilities, and the R&D
1913 projects funded as part of DOE’s SciDAC program are critical elements of the HL-LHC software
1914 upgrade effort. While S^2I^2 R&D efforts will tend to be complementary, the Institute will establish
1915 contacts with all of these projects and will use the blueprint process (described below) to establish
1916 a common vision of how the various efforts align into a coherent set of U.S. activities.

1917 **CERN:** As the host lab for the LHC experiments, CERN must be an important collaborator for
1918 the Institute. Two entities within CERN are involved with software and computing activities. The
1919 IT department is focused on computing infrastructure and hosts CERN openlab (for partnerships
1920 with industry, see below). The Software (SFT) group in the CERN Physics Department develops
1921 and supports critical software application libraries relevant for both the LHC experiments and the
1922 HEP community at large, most notably the ROOT analysis framework and the Geant4 Monte
1923 Carlo detector simulation package. There are currently many ongoing collaborations between the
1924 experiments and U.S. projects and institutions with the CERN software efforts. CERN staff from
1925 these organizations were heavily involved the CWP process. The Institute will naturally build on
1926 these existing relationships with CERN. A representative of CERN will be invited to serve as a
1927 member of the Institute Steering Board, as described in Section 8.

1928 **The HEP Software Foundation (HSF):** The HSF was established in 2015 to facilitate inter-
1929 national coordination and common efforts in high energy physics (HEP) software and computing.
1930 Although a relatively new entity, it has already demonstrated its value. Especially relevant for the
1931 S^2I^2 conceptualization project, it organized the broader roadmap process leading to the parallel
1932 preparation of the Community White Paper. This was a collaboration with our conceptualization
1933 project, and we expect that the Institute will naturally partner with the HSF in future roadmap
1934 activities. Similarly, it will work under the HSF umbrella to sponsor relevant workshops and coor-
1935 dinate community efforts to share information and code.

1936 **Industry:** Partnerships with industry are particularly important. They allow R&D activities to be
1937 informed by technology developments in the wider world and, through dedicated projects, to inform
1938 and provide feedback to industry on their products. HEP has a long history of such collaborations
1939 in many technological areas including software and computing. Prior experience indicates that
1940 involving industry partners in actual collaborative projects is far more effective than simply inviting
1941 them for occasional one-way presentations or training sessions. There are a number of projects
1942 underway today with industry partners. Examples include collaboration with Intel like the Big
1943 Data Reduction Facility [93], through an Intel Parallel Computing Center [94], with Google [95,96]

and AWS [95–97] for cloud computing, etc. A variety of areas will be of interest going forward, including processor, storage and networking technologies, tools for data management at the Exabyte scale, machine learning and data analytics, computing facilities infrastructure and management, cloud computing and software development tools and support for software performance.

In 2001 CERN created a framework for such public-private partnerships with industry called CERN openlab [98]. Initially this was used to build projects between CERN staff and industry on HEP projects, however in recent years the framework has been broadened to include other research institutions and scientific disciplines. Fermilab has recently joined CERN openlab collaboration and Princeton University is currently finishing the process to join. Others may follow. CERN openlab can also be leveraged by the Institute to build partnerships with industry and to make them maximally effective. This can be done in addition to direct partnerships with industry.

9.2 The Blueprint Process

To facilitate the development of effective collaborations with the various partners described above, the Institute should proactively engage and bring together key personnel for small “blueprint” workshops on specific aspects of the full R&D effort. During these blueprint workshops the various partners will not only inform each other about the status and goals of various projects, but actively *articulate and document* a common vision for how the various activities fit together into a coherent R&D picture. The scope of each blueprint workshop should be sized in a pragmatic fashion to allow for convergence on the common vision, and some of the key personnel involved should have the means of realigning efforts within the individual projects if necessary. The ensemble of these small blueprint workshops will be the process by which the Institute can establish its role within the full HL-LHC R&D effort. The blueprint process will also be the mechanism by which the Institute and its various partners can drive the evolution of the R&D efforts over time, as shown in Figure 10.

Following the discussions at the November 2017 S^2I^2 -DOE workshop on HL-LHC R&D, we expect that jointly sponsored blueprint activities between NSF and DOE activities relevant for HL-LHC, the US LHC Operations Programs and resource providers like OSG will likely be possible. All parties felt strongly that an active blueprint process would contribute significantly to the coherence of the combined U.S. efforts. The Institute could also play a leading role to bring other parties into specific blueprint activities, where a formal joint sponsorship is less likely to be possible. This may include specific HEP and CS researchers, other relevant national R&D efforts (non-HEP, non-DOE, other NSF), international efforts and other external software providers, as required for the specific blueprint topic.

Blueprint activities will likely happen 3-4 times per year, typically with a focus on a different specific topics each time. The topics will be chosen based on recognizing areas where a common vision is required for the coordination between partners. Input from the Institute management, the Institute Steering Board and the management of various partner projects and key personnel will be explicitly solicited to identify potential blueprint activities. The Institute will take an active role in organizing blueprint activities by itself and jointly with its partners based on this input. From year to year specific topics may be revisited.

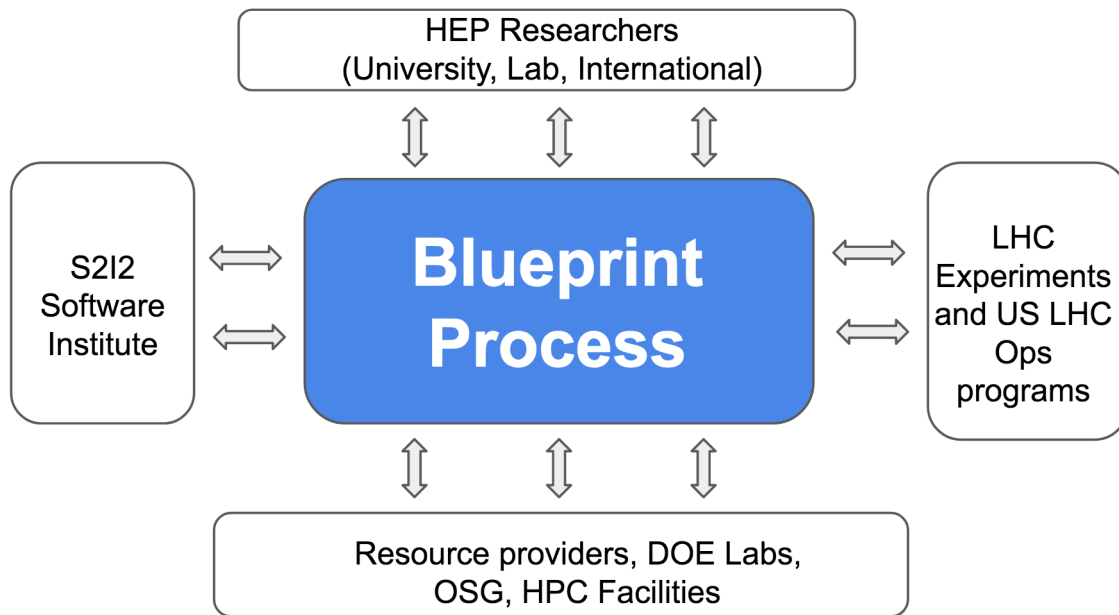


Figure 10: The Blueprint Process will be a primary means of developing a common vision with the major partners.

10 Metrics for Success (Physics, Software, Community Engagement)

11 Training, Education and Outreach

11.1 The HEP Workforce

People are the key to successful software. Computing hardware becomes obsolete after 3 – 5 years. Specific software implementations of algorithms can have somewhat longer lifetimes (or shorter). Developing, maintaining, and evolving algorithms and implementations for HEP experiments can continue for many decades. Using the LEP tunnel at CERN for a hadron collider like the LHC was first considered at a workshop in 1984; the ATLAS and CMS collaborations submitted letters of intent in 1992; the CERN Council approved construction of the LHC in late 1994, and it first delivered beams in 2009. A decade later, the accelerator and the detectors are exceeding their design specifications, producing transformative science. The community is building hardware upgrades and planning for an HL-LHC era which will *start* collecting data circa 10 years from now, and then acquire data for at least another decade. People, working together, across disciplines and experiments, over several generations, are the real cyberinfrastructure underlying sustainable software. Investing in people through training over the course of their careers is vital part of supporting this human facet of scientific research. Training should include scientists and engineers at all stages of their careers, but should take particular care to invest in the young students and postdocs who will be faculty leaders driving the research agenda in the HL-LHC era.

HEP algorithms and their implementations are designed and written by individuals with a broad spectrum of expertise in the underlying technologies, be it physics, or data science, or principles or

computing, or software engineering. Almost all Ph.D. students write analysis software, as do most post-docs. Many students and post-docs write software to acquire data, calibrate and reconstruct it, and reduce data sets to sizes manageable for analysis by teams and individuals. Some of these people have very high levels of domain and software engineering expertise, and some are raw recruits. For example, most experiments have dedicated teams for developing and maintaining code for tracking charged particles. The most senior members of these teams generally have many years of experience and have developed deep understandings of the current algorithms and their performances, both in terms of physics performance and resource usage. This wisdom is passed along in a somewhat unorganized way through what amounts to an unofficial apprenticeship program.

In addition, teams of “core” developers are responsible for designing and implementing software for workflow and workload management. These individuals are often responsible for managing use of these tools to run what are often commonly “central productions” of reconstruction, stripping, and simulation campaigns. Members of these teams are considered software professionals, although many have been formally trained in HEP rather than computer science or software engineering. Matching the educational and training opportunities to the needs of the various levels of software developers across the full spectrum of the community will require carefully assessing what skills and expertise will have the biggest impact on physics. In addition, as most people earning Ph.D.’s in experimental particle physics eventually leave the field, providing educational and training opportunities that prepare them for other career trajectories must be a consideration in setting priorities.

11.2 Current Practices

Training support for software-related activities in HEP is uneven and made up of a patchwork of training activities with some significant holes. Although most universities do provide some relevant computer science and software engineering courses, and many are starting to provide introductory “data science” courses, many HEP graduate students and postdocs are not required to take these classes as part of the curriculum. As students enter the research phase of the graduate student training, many recognize the value of such classes, but are no longer in a position to easily take the classes. No “standard” recommendations exist for incoming students, either for HEP experiments or the HEP field as a whole. Some universities are developing curriculums for STEM training in general and/or “certificate” programs for basic data science and/or software training, but these are by no means yet universal. The result is that the graduate student and postdoc population has a very diverse spectrum of relevant skills.

HEP collaborations do typically provide opportunities for members to learn the software tools developed by and/or used within the experiments. For example, the week-long CMS Data Analysis School (CMSDAS) [99] pairs software experts with new collaborators to build and run end-to-end examples of real analysis applications. Similarly, LHCb has a training program and workshops called the “Starter Kit” [100]. Other collaborations have similar programs. The goals of these programs are primarily to make new collaborators effective *users* of the complex experiment software ecosystems, rather than effective developers of that ecosystem, even if the latter will be often an important part of their eventual research contribution. In addition these programs need to train collaborators with very uneven backgrounds in basic ideas of computer science and software engineering, as described above.

A number of summer schools focused on more advanced software and computing topics also exist in the global HEP community. These include, among others, the CERN School of Computing [101], the GridKa school [102] in Germany organized by the Karlsruhe Institute of Technology, the “Developing Efficient Large Scale Scientific Applications (ESC)” [103], school organized by the Istituto Nazionale di Fisica Nucleare (INFN) in Italy and (more recently) the “Computational and Data Science for High Energy Physics (CoDaS-HEP)” school [104] in the U.S.

11.3 Workforce Development in HEP

Large HEP experiments organize themselves for global efforts in many areas, including large detector construction projects, globally distributed computing systems, software development involving hundreds of researchers and, of course, extremely complex multi-faceted data analysis activities. Despite the universal need for computational skills for nearly all HEP researchers, little exists to bring together all of the pieces of an end-to-end training program accessible to all HEP students and postdocs, as well as more advanced training for more senior HEP researchers or specialists. In addition many of the individual HEP training efforts suffer from sustainability issues, even when they are successful.

This is a place where the Institute could play a leadership role, by working towards a common vision for what HEP workforce development for software and computing would look like. In practice the Institute should not aim to create such a end-to-end training program by itself, but rather to focus resources in two areas. First, it can build alliances between the existing successful training HEP and non-HEP programs and schools as well as the HEP experiment-specific training efforts. Within the U.S. it can also augment this by documenting and promoting a vision for how university courses, certifications and programs can build the necessary base of skills. Second, it can fill in some of the gaps between those efforts.

There are a lot of experiment-specific training efforts and within a single experiment, different skill sets are needed. In addition to a base skill set that contains basic programming language knowledge, testing and code management tools and experiment-specific framework knowledge, there are more specialized skills that only a subset of the community needs to know, such as software optimization, or low-level hardware interfaces. We should probably strive to extract common knowledge and build common training from these efforts, because it enables us to duplicate less effort on experiment-specific training, and to do the shared training better by accumulating more expertise into it.

Many people in the field believe that core elements of computer science, computer programming, and software engineering should be required of *all* students embarking on a Ph.D. in experimental HEP. Some undergraduate programs provide good opportunities in this regard, but there is no universal expectation that this is prerequisite to beginning graduate level study in a U.S. university. Nor do most Ph.D. programs offer formal coursework like this. As a result, the HEP community needs to decide what it expects all of its students to know, and to prepare appropriate pedagogic material that can be used, either in the formal classroom or for independent study. Elements of this material have been assembled by individual instructors, or is taught piecemeal by experiments, but a coherent approach should be developed.

HEP has a set of concepts and a software infrastructure for analyzing data that is approximately domain-specific and transcends individual experiments. The most common analysis framework is the ROOT library developed principally at CERN. It encodes methods for selecting datasets, visualizing data, extracting parameters that describe data, etc. The community is rapidly adopting similar tools from the larger scientific Python community. Some students are introduced to these very informally by mentors who give them tutorials and/or working examples to get started. Some are provided experiment-specific tutorials (in-person or online) to get started. A software institute can take a leading role in collecting, developing, and maintaining a curated set of educational materials that addresses the common software needs of all students starting to do analysis. It can also organize video-based classes or in-person “summer schools” to teach this material.

In addition to writing analysis code, many members of the HEP community write software which becomes part of the experimental infrastructure. Examples of this are reconstruction software, event selection software (at either the trigger level or the offline “stripping” level), simulation software, and data visualization software. Each of these requires both domain expertise and algorithmic design plus software engineering expertise. Providing the training to build high-quality, performant,

sustainable software for these types of applications is qualitatively different – it requires a much higher level of instructor expertise, and the target audience is generally smaller. As such a large fraction of the processing power is deployed for reconstruction, training the lead developers how to use performance tools to study hot spots and memory access patterns, how to design data structures and algorithms to take advantage of vector processors in modern architectures, and how to write thread-safe algorithms is absolutely critical to using computing resources efficiently. Similarly, if we want event selection software to use algorithms built on top of ML learning tools, we must train the developers of that software the underlying principles of ML, what tools exist, how to use those tools to train neural networks or BDTs efficiently, and how to deploy inference engines that execute quickly. In many cases, the state-of-the-art is evolving very rapidly. This means that developers will need continuing education, and much of it should be hands-on and interactive. An Institute will be a natural home for this type of training.

Where appropriate, training programs should take advantage of developments in pedagogy, such as active learning¹ or peer learning². In some cases, it may be advantageous to have code samples that are purposely broken or flawed, and ask students to fix or improve them. Learning material so that it sticks with the students often takes more effort by both the students and the instructors; it often takes more time than we would prefer. However, it is the best way to ensure an educated community that can fully contribute to the physics programs at large, which is really the ultimate goal training programs.

A difficulty that has emerged in the past with respect to implementation of training courses is the lack of funding along with the lack of available time by experts in the field. People with enough expertise or insight in the field have usually no time to devote to prolonged periods of student’s training, and, even when they can find some, the cost of setting up a training course in an effective way is often beyond what’s made available by funding agencies (funds for travel, hosting, setting up a room with a computing infrastructure to allow interactive hands-on session, etc.). A possible way out is a completely different approach to training (but complementary to the already existing and successful classical efforts such as the CERN School of Computing’s Bertinoro and KIT ones): instead of directly teaching to students, trainees could make use of a web-based platform to provide training materials to students. This complementary approach has several advantages over traditional ones:

11.4 Knowledge that needs to be transferred

At all stages of software & computing training, we should take care to encourage Good Practices Across the Community (GPAC), such as error checking, modularity of code design, writing tests, etc. All the key concepts addressed in the training should not be specific to a particular experiment or field of application, but general enough to be useful for the whole HEP community and possibly beyond. In this section, we present a list of specific concepts that need to be taught to members of the community, in order to guarantee the base level of competence needed to write efficient code for the different tasks performed in HEP experiments.

Base knowledge to be transferred includes basic programming concepts, data structures, basics of code design, error checking, code management tools, validation and debugging tools. More advanced topics include modularity of code design, advanced data structures, evaluation metrics, writing tests and working with different types of hardware accelerators. Special emphasis should be made on reporting results and documenting them.

- Basic Programming Concepts
 - Object oriented paradigm

¹<http://www.crlt.umich.edu/tstrategies/tsal>

²https://en.wikipedia.org/wiki/Peer_learning

- 2145 – Compiled languages (C++)
- 2146 – Scripting languages (Python, Javascript,...)
- 2147 • Algorithms
- 2148 – Boost library
- 2149 – STL algorithms for containers
- 2150 – R and/or ROOT
- 2151 • Existing frameworks (development or application level)
- 2152 – Qt
- 2153 – ROOT
- 2154 – experiment specific framework (possibly if of potential interest outside the native exper-
- 2155 – iment)
- 2156 • Code design (design patterns)
- 2157 • Development tools
- 2158 – IDEs (Integrated Development Environment)
- 2159 – Debuggers
- 2160 – Profilers
- 2161 • Evaluation metrics
- 2162 • “Trust” metrics such as data driven tests
- 2163 • Specific software implementation training
- 2164 • Good practices
- 2165 • Code style and clarity
- 2166 • Scripting and data cleaning
- 2167 • Reporting results reproducibly
- 2168 • Writing Documentation

2169 11.5 Roadmap

- 2170 • Work with the Carpentries (software & data) to customize (focusing on what is needed in
- 2171 HEP, making examples HEP-specific) general/basic software training for new students
- 2172 • Work with HPC centers when training needs/goals overlap, e.g. DOE Lab and university
- 2173 computing centers that provide live, virtual, and recorded training
- 2174 • Summer schools
- 2175 • Focused webinars on specific topics (both beginner and advanced)
- 2176 • Focused webinars on specific topics (both beginner and advanced), this could be collaborative
- 2177 with software, HPC, data science communities
- 2178 • Provide advanced/focused hands-on in-person and virtual training on a variety of HEP-
- 2179 specific topics (following CMS-HATS model)
- 2180 – Coordinate with experiments & LHC physics centers, for content, instructors, and train-
- 2181 – ing venues
- 2182 – Initial topics: Analysis in python, analysis in R, histogramming, PyROOT and rootpy,
- 2183 ML to improve Physics Objects, tracking tagging, Modern Tools for Physics Analysis-
- 2184 – Roofit, MVA
- 2185 – Method for bringing in new topics:
- 2186 * Suggestions from users and developers, user survey

- 2187 * Find willing instructors (from LHC Experiments etc)
- 2188 * Institute's role is coordinator, not funder, not instructor (though maybe will fund/help
- 2189 students?, pay for instructor travel for in-person training?)

2190 **11.6 Outreach**

2191 Outreach and use of HEP data by researchers in other fields and members of the public with
2192 scientific interests (linked to software/data preservation and reproducibility within Analysis focus
2193 area)

- 2194 • Provide data and tools to the non-HEP researchers, e.g. computer scientists who want to
- 2195 work on big data problems
- 2196 • Provide data and tools to the interested public
- 2197 • Document data and tools and provide examples of usage
- 2198 • How do members of the public get access to enough computing to work with this data? (HEP
- 2199 data analysis science gateway/portal?)
- 2200 • Bringing together Inreach and Outreach community

12 Broadening Participation

The participation of women and ethnic minorities is generally low in the HEP world, and fractionally, it is even lower in the HEP Software and Computing (S&C) world. We do not understand why this is true, but it is what it is. We estimate that fewer than 10% of people in HEP S&C are women while (from LHC experiment statistics) between 18% and 20% of the LHC’s experiments’ collaborators are women. Nationally, 7.4% of high-tech employment is black [105], while in HEP S&C the fraction is negligibly small. Looking forward, increasing the diversity of the HEP S&C workforce promises two types of benefits. From first principles, the top 5% of a larger pool should always be better than the top 10% of a pool half as large; in addition, teams of people from diverse backgrounds are generally believed to be more creative when crafting solutions to complex problems.

An S^2I^2 will not significantly increase the fraction of under-represented populations in HEP or in S&C; it will be too small a player. However, the Institute must be sensitive to diversity in building its own teams, and it can help build the pipeline by partnering with institutions actively working to do this. At the high school level, programs like QuarkNet engage diverse groups of students. At the undergraduate level, “Women in Science and Engineering” programs like those at the University of Michigan, the University of Arizona, and the University of Cincinnati (as examples) provide REU opportunities targeting women. At the transition from undergraduate to Ph.D. student level, the American Physical Society’s Bridge Program targets under-represented populations (self-identified, so including first generation college, as an interesting example). The Institute’s outreach and education program can include supporting programs like these by providing both financial support and opportunities to work with S^2I^2 teams and their collaborators. One model would be sponsoring undergraduate and graduate student Fellowships, similar to those offered by the DIANA project. In this case, one of the three Fellows who has already completed their projects was a woman, and two of the two lined up for early 2018 Fellowships is a woman. In addition, the S^2I^2 can work with groups like **Data Carpentry** to organize workshops using HEP data to introduce high school and college students to data science.

If there is an S^2I^2 solicitation, the proposal should identify specific models and partners for encouraging participation of under-represented populations in its outreach and education program. The proponents will need to reach out to institutions with programs with track records of increasing diversity to find out what works. They exist. The NYU Material Research Science and Engineering Center runs an REU program with 40% minority student participation and 50% women. The University of Maryland, Baltimore Campus has become a center for cultivating underrepresented minority scholarship and awareness in the math, science, and engineering disciplines. Florida International University, which has a CMS group, serves a student population which is predominantly Hispanic and 85% minority. While increasing diversity will not be the primary goal of the envisioned S^2I^2 , devoting 2% - 3% of its resources to outreach and education efforts targeting the pipeline can have a beneficial impact on diversity in the HL-LHC era.

13 Sustainability

14 Risks and Mitigation

The primary goal of the envisaged S^2I^2 is to enable the science goals of the HL-LHC through software R&D leading to deployment of the requisite software by the experiments. The risks are *social*, *technical*, and *contextual*. Those in the *social* category include risks related to: (i) building and maintaining the S^2I^2 team, (ii) fully engaging in a coherent fashion with the larger HL-LHC software community, and (iii) executing the R&D plan successfully. Those in the *technical* category include: (i) slower improvement of hardware performance than anticipated, (ii) less benefit from

new features like parallelization and SIMD vectorization than anticipated, and (iii) less benefit from Machine Learning than anticipated. Those in the *contextual* category include: (i) substantial changes to the hardware upgrade plans for the accelerator and detector, (ii) substantial changes to the upgrade software R&D funding profiles by other agencies, and (iii) major scientific discoveries at the LHC, before the HL-LHC era begins, that significantly change the physics priorities of the experiments. Each of these requires different specific mitigations, but all require regular review of progress by S^2I^2 management, the outside stakeholders, and “disinterested” external advisors coupled with the agility to redirect resources.

Building the Institute team will be the first major challenge. Subsequently maintaining an effective team will be a continuing challenge that requires careful thought in advance, as well as continuing attention. An Executive Director, and probably a Deputy Executive Director, will lead the Executive Board. The initial choices for these positions will be the responsibility of the lead PI and co-PIs, probably taken in consultation with the NSF while negotiating a Cooperative Agreement before an award is made. The individuals selected for these roles will need to devote substantial fractions of their professional effort to the Institute. The Executive Director will almost certainly need to devote at least 50% of his/her effort to the position. In general, the Deputy Executive Director must be willing to serve as Interim Executive Director should the occasion arise, and must be willing to devote enough effort to the Institute to be ready assume this role on short notice. An initial team of Area Managers should be identified while negotiating a Cooperative Agreement. The specific individuals can be identified only when the number of Focus Areas to be supported is known. In addition to their domain expertise, members of the Executive Board should broadly represent the interests of the LHC experiments. They should all have track records of collaborating effectively. The activities of the core team and in each Focus Area will be formally reviewed each year to prepare annual Statements of Work (SOWs) to be done. At this time, it will be appropriate to consider whether new Area Managers (plus Co-Managers or Deputy-Managers) should be appointed. These decisions will be taken by the Executive Board, in consultation with the Steering Board. Should the Executive Director or Deputy Executive Director step down before the five-year term of the award ends, the lead PI and co-PIs will select a replacement, in consultation with the Executive Board and the Steering Board.

Building a team of approximately 20 FTE physicists and software professionals to undertake the support and R&D responsibilities of the Institute will take time. The number of highly qualified personnel with the requisite domain and software expertise is limited. The Institute will initially build on the existing software development infrastructure by co-funding individuals whose other activities complement those being undertaken by the Institute. For example, it would be appropriate for someone already working in DOMA as part of the ATLAS or CMS operations team to continue that work half time and begin to work on the software upgrade R&D as a member of the Institute. Similarly, someone who already provides continuous integration services, packaging, etc., to an ongoing project could provide similar services in support of Institute projects on a part-time, co-funded basis. While a certain level of finesse will be required to ensure that individuals funded for different projects are splitting their efforts appropriately, co-funding these people will provide opportunities to build a sense of community across experiments and help keep the Institute focussed on efforts of interest to the experiments. The highest priority for the first year will be hiring members of the core team who are anticipated to continue as members of the Institute for the full term of the award. In parallel, the Area Managers for the R&D Focus Areas will identify a mix of post-docs, more senior physicists, and software professionals with the expertise and interest to advance their research programs. As appropriate, these individuals can be co-funded or hired directly by a University group, in conjunction with an SOW.

Building a team is more complicated than hiring individuals to do well-defined jobs – in this case we want people to collaborate with each other effectively, and also with the larger community. This will require defining expectations for collaborative work, and rewarding it meaningfully. Code and

documentation will be reviewed as part of the engineering process; we have observed that this both improves software products *and* tends to build a sense of community, perhaps because it creates joint responsibility and ownership. Similarly, developers will be expected to present their work within the Institute and also to the larger community. Team progress and individual performance will be formally reviewed on an annual basis. Individuals will also be asked to prepare written 3-, 6-, and 12-month goals and plans on a rolling basis, for less formal discussions with their immediate supervisors. This process will provide opportunities to laud excellent work (which is generally expected) and identify the need for remediation, when indicated. Where appropriate, individuals will be members of LHC experiments as well as members of the Institute. In these cases, S^2I^2 management will work with the experiments' managements to make sure that S^2I^2 efforts are explicitly recognized as service work to the experiments. This will be especially important for students, post-docs, and other more junior members of experiments who are expected to engage in a mix of service work and physics analysis as part of their professional development.

Some of the approaches to building and maintaining the team, discussed above, also address a second key issue: *fully engaging with the larger HL-LHC software community*. The Steering Board (discussed in Section 8) will explicitly have representatives of all of the LHC experiments, the US-LHC Operations programs as well as Fermilab and CERN. They will review the large scale priorities and strategy of the Institute quarterly, and provide advice on any changes of direction that should be considered. Just as this Strategic Plan has emerged from a community process, executed in parallel with the broader CWP process, the Institute will sponsor continuing blueprint activities, in conjunction with the HEP Software Foundation, to update the roadmap on a rolling basis and identify any changes in priorities. Additionally, members of the Institute will make presentations to the individual experiments at different levels of technical detail. At the finest level of granularity, presentations of specific algorithms and implementations will be presented to tools Working Groups. At a coarser level of granularity, projects will be presented and discussed during Software and Computing Weeks. When appropriate, overview presentations will be made at general Collaboration Meetings. In all cases, the goal will be two-way communication.

Executing the R&D plan successfully will require that developers be technically strong, that they work together collaboratively, and that they adhere to good software engineering practices. It will also require careful attention to short-term (and longer-term) goals by members of the Executive Board, some of whom may be developers themselves. Many of the software engineering practices described in Section ?? are meant to help keep projects on track and assure high quality. As an example, requiring that all code be reviewed by a second developer before a merge request is accepted will help assure good documentation and correct algorithmic implementation. As discussed in Section 8, the EB will meet weekly to make short term plans to keep efforts properly focussed.

The magnitudes of the computing challenges described in Section 3 *assume* that CPU and mass storage performance per unit price will continue to grow at a rate equivalent to 20% per year, about a factor of 6 over a decade, and that the funding profile will remain flat. The experiments' needs are *another* factor of 5 greater, given current algorithms' use of resources. The purpose of the S^2I^2 is to undertake a software upgrade to provide the enhanced performance required use the anticipated resources that much more effectively. Should effective hardware costs drop more slowly than estimated, or the hardware acquisition budgets drop, the software goals may need to be revised. At the moment, we want to find better algorithms to reconstruct, process, and analyze data with essentially the same fidelity as is done today. If this is not possible, the experiments will need to process as much data as possible with lower fidelity or less data with greater fidelity. The Institute, with advice from the Steering Committee will need to adjust its goals and priorities accordingly. Somewhat similarly, if new algorithms taking advantage of vectorization and machine learning do not deliver the anticipated improvements in performance in the next five years, the HL-LHC experiments will need adapt their plans to live with what is possible.

The American philosopher Yogi Berra warned that *It's tough to make predictions, especially about the future*, and the Scottish poet Robert Burns observed that *The best laid schemes of mice and men Go often askew* (as translated into modern English). We recognize that both these insights apply to preparing a software upgrade for deployment almost 10 years from now. Nonetheless, the plan for R&D over the next 5 years should be relatively robust. Changes in the accelerator and detector upgrade plans are most likely to produce quantitative, not qualitative, changes in the computing and software models for the HL-LHC. Similarly, changes in physics priorities resulting from discoveries made before the HL-LHC turns on may require re-balancing the computing resources – perhaps more reconstruction and less simulation, or vice versa. But the key problems will remain the same. In general, the time scales for these contextual changes should be long enough that the Institute's regular reviews will permit it to adapt its efforts without significant disruption.

15 Funding Scenarios

The costs of an S^2I^2 will depend on its scope and its relationships to other entities. Most are estimated in terms of nominal full-time-equivalent (FTE) professionals. Approximately a third of the funding will support core personnel and other backbone activities. The remaining funding will primarily support personnel, affiliated with other university groups, to lead and contribute to software R&D in the identified focus areas.

Some of the Institute personnel may be working only on S^2I^2 projects. However, most effort will be done by a mixture of software professionals working part-time on S^2I^2 projects and part-time on complementary projects, funded through other mechanisms, plus post-docs and graduate students supported partly by the S^2I^2 for their work on its projects and supported partly by other funds for related and complementary activities. Co-funding individuals with relevant expertise will be a key method of ensuring significant community buy-in and engagement. The Institute may undertake some projects on its own, but *most* should be of sufficient interest to attract support from elements of the community who want to collaborate. For example, one of the topics in the *Reconstruction and Trigger Algorithms* focus area, identified as important by all the experiments, is learning to use vectorization programming techniques effectively. An individual might develop generic toolkits (or algorithms), funded by the Institute, and test them (or deploy them) in experiment-specific software, funded by a partner. In such a case, the Institute is leveraging its resources *and* ensuring that its work is relevant to at least one experiment.

As a first approximation, we estimate that the fully loaded cost of a software professional FTE will average \$200K/year. Typically, this will include salary, fringe benefits, travel, materials and supplies, plus overhead. Based on the experience of the OSG, we estimate that operations personnel will average \$160K/year.

We expect that the core team will include an Executive Director and project/administrative support plus a core set of software professionals who will (i) engage directly in R&D projects related to established focus areas and exploratory studies, (ii) provide software engineering support across the program, (iii) provide the effort for the Institute “backbone” focused on developing, documenting and disseminating best practices and developing incentives, (iv) provide some services (e.g., packaging and infrastructure support across the program), (v) lead the education and outreach effort, (vi) lead the blueprint effort, (vii) coordinate efforts to build bridges beyond the S^2I^2 itself to the larger HEP, Computer Science, Software Engineering, and Data Science communities and to establish the Institute as an intellectual hub for HL-LHC software and computing R&D. Depending on the funding available, and the overall scope of the project, we anticipate that the team will consist of the Executive Director plus 5 – 7 FTEs. As a first approximation, the bottom lines for what be deemed “central” expenses range from \$1200K/year to \$1800/year.

An essential element of building a software R&D will be sponsoring workshops and supporting participation in other relevant workshops. Based on our experience with the S^2I^2 conceptualization

process, a Participant Costs budget of \$200K/year will prove sufficient, in large measure because these funds can be used to supplement those from other sources for many people. Similarly, we estimate that a \$200K/year Participant Costs budget reserved for summer schools and other explicitly pedagogic activities will make a significant impact. In the tighter budget scenarios, these last two items could be reduced stepwise to half in the lowest scenario.

Beyond the core efforts and backbone team, we anticipate funding an average of 4 FTE lines for each of four focus areas in the fully funded scenario, about \$800K/year each. This level of effort would provide *critical mass* to guarantee a significant leading impact on a focus areas, given previous experience in smaller (NSF-funded) projects such as DIANA-HEP [90], DASPOS [106], the Parallel Kalman Filter Tracking Project [107] and the “Any Data, Any Time, Anywhere: Global Data Access for Science” [86] project. Almost none of the personnel funded by these lines would be fully funded by the S^2I^2 – the projects they will work on should be of sufficient interest to the community that collaborators will co-fund individuals whose other projects are closely aligned with their Institute projects. The total expense of these activities in a fully funded project would be \$3200K/year. If sufficient funding is not available, the number of focus areas would be reduced, rather than trying to fund all at insufficient levels. The bare minimum number of focus areas to have a significant impact on HL-LHC software development would be 2, at a cost of \$1600K/year.

Beyond the software R&D scope envisioned for the Institute when the S^2I^2 conceptualization process started, we have considered the possibility that a single institute might serve as an umbrella organization with OSG-like operational responsibilities related to the LHC experiments, as well. As indicated in Table 2, this would require supporting 13.3 FTE operations personnel at an estimated cost of ~\$2100K/year.

scenario	core and backbone	participant costs	focus areas	operations	total
low R&D	1200	200	1600		3000
medium R&D	1400	300	2400		4100
high R&D	1800	400	3200		5400
OSG-HEP				2100	2100

Table 3: Three possible budget scenarios for the R&D efforts, plus the OSG-HEP operations effort. All entries are k\$/year.

Three software R&D scenarios (no OSG-like operations responsibilities) are illustrated in Table 3. The numbers are rough estimates. Funding for OSG-like operations adds another \$2100K to any of these. A proposal responding to a solicitation will need to provide better estimates of the funding required to cover the proposed activities. For the purposes of a strategic plan, we tentatively identify the “Reconstruction and Trigger Algorithms” and “Data Organization, Management and Access” focus areas to be the very highest priority for S^2I^2 funding. The former is closest to the core physics program, and it is where U.S. university groups have the most expertise and interest. The latter covers core technologies tying together processing all the way from data acquisition to final physics analysis. It is inherently cross-disciplinary, and will engage U.S. university HEP, Computer Science, and Software Engineering researchers. Data Analysis Systems R&D is essential to the success of the HL-LHC. If insufficient funding is available through this funding mechanism, efforts in this area might be funded through other mechanisms or might be deferred. However, continuity of effort from the existing NSF-funded DIANA-HEP project [90] and the ability to test run analysis system solutions during LHC Run 3 will be at risk. Applications of Machine Learning garnered the highest level of interest during the CWP and S^2I^2 conceptualization processes, and it is especially well suited to cross-disciplinary research. Deciding not to include this as one of the two highest priority focus areas at this stage was a close call. Depending on the details of a solicitation and the anticipated funding level, it might displace one of the focus areas identified as

higher priority here.

A Appendix - S^2I^2 Strategic Plan Elements

The original S^2I^2 -HEP proposal was written in response to solicitation NSF 15-553 [47]. This solicitation specified that: “The product of a conceptualization award will be a strategic plan for enabling science and education through a sustained software infrastructure that will be freely available to the community, and will address the following elements:”

- the science community and the specific grand challenge research questions that the S^2I^2 will support;
- specific software elements and frameworks that are relevant to the community, the sustainability challenges that need to be addressed, and why addressing these challenges will be transformative;
- appropriate software architectures and lifecycle processes, development, testing and deployment methodologies, validation and verification processes, end usability and interface considerations, and required infrastructure and technologies;
- the required organizational, personnel and management structures and operational processes;
- the requirements and necessary mechanisms for human resource development, including integration of education and training, mentoring of students, postdoctoral fellows as well as software professionals, and proactively addressing diversity and broadening participation;
- potential approaches for long-term sustainability of the software institute as well as the software; and
- potential risks including risks associated with establishment and execution, necessary infrastructure and associated technologies, community engagement, and long-term sustainability.

Moreover the solicitation states that “The strategic plan resulting from the conceptualization phase is expected to serve as the conceptual design upon which a subsequent S^2I^2 Implementation proposal could be based.”. In this “Strategic Plan” document, we have attempted to respond to these criteria.

We note in addition that the same solicitation (NSF 15-553 [47]) also allowed for implementation proposals for “Chemical and Materials Research” and “Science Gateways”. For these implementation proposals, the solicitation requested the following elements in the (20 page) proposals:

- The overall rationale for the envisioned institute, its mission, and its goals.
- A set of software issues and needs and software sustainability challenges faced by a particular, well-defined yet broad community (that is clearly identified in the proposal) that can best be addressed by an institute of the type proposed, a compelling case these are the most important issues faced by the community, and that these issues are truly important.
- A clear and compelling plan of activities that shows how the proposed institute will address these issues and needs by involving (and leveraging) the community, including its software developers, in a way that will benefit the entire community.
- If there are other NSF-funded activities that might appear to overlap the institute’s activities, a discussion clarifying how the funding of each activity will be independent and non-overlapping.
- Metrics of how success will be measured, that include at least impact on the developer and user communities.

- Evidence that the people involved in planning and setting up the institute have the organizational, scientific, technical, and sociocultural skills to undertake such a task, and that they are trusted and respected by the community as a whole.
- Evidence of a high degree of community buy in that a) these are the urgent/critical needs and b) this institute is the way to address them.
- A plan for management of the institute, including 1) the specific roles of the PI, co-PIs, other senior personnel and paid consultants at all institutions involved, 2) how the project will be managed across institutions and disciplines, 3) identification of the specific coordination mechanisms that will enable cross-institution and/or cross-discipline scientific integration, and 4) pointers to the budget line items that support these management and coordination mechanisms.
- A steering committee composed of leading members of the targeted community that will assume key roles in the leadership and/or management of the institute. A brief biography of the members of the steering committee and their role in the conceptualization process should be included.
- A plan for how the institute activities will continue and/or the value of the institute's products will be preserved after the award, particularly if it does not receive additional funds from NSF.

As these criteria are general enough to be relevant also for an S^2I^2 for HEP, we have included also some initial information on these items in this document.

In addition, a National Academy of Science report, *Future Directions for NSF Advanced Computing Infrastructure to Support U.S. Science and Engineering in 2017-2020* [88], appeared shortly before the S^2I^2 -HEP project began. One of its general recommendations is that NSF “collect community requirements and construct and publish roadmaps to allow it to better set priorities and make more strategic decisions about advanced computing” and that these roadmaps should “would reflect the visions of the science communities supported by NSF, including both large users and those (in the “long- tail”) with more modest needs. The goal is to develop brief documents that set forth the overall strategy and approach rather than high-resolution details. They would look roughly 5 years ahead and provide a vision that extends about 10 years ahead.” The S^2I^2 -HEP and CWP community processes should be seen as input regarding the vision of the HEP community for the HL-LHC era.

B Appendix - Workshop List

During the process we have organized a number of workshops and sessions at preexisting meetings. These included (in chronological order):

S^2I^2 HEP/CS Workshop

Date: 7–9 Dec, 2016

Location: University of Illinois at Urbana-Champaign

URL: <https://indico.cern.ch/event/575443/>

Summary report: <http://s2i2-hep.org/downloads/s2i2-hep-cs-workshop-summary.pdf>

Description: This workshop brought together attendees from both the particle physics and computer science (CS) communities to understand how the two communities could work together in the context of a future NSF Software Institute aimed at supporting particle physics research over the long term. While CS experience and expertise has been brought into the HEP community over the years, this was a fresh look at planned HEP and computer science research and brainstorm about engaging specific areas of effort, perspectives, synergies and expertise of mutual benefit to

2520 HEP and CS communities, especially as it relates to a future NSF Software Institute for HEP.
2521

2522 **HEP Software Foundation Workshop**

2523 *Date:* 23–26 Jan, 2017

2524 *Location:* UCSD/SDSC (La Jolla, CA)

2525 *URL:* <http://indico.cern.ch/event/570249/>

2526 *Description:* This HSF workshop at SDSC/UCSD was the first workshop supporting the CWP
2527 process. There were plenary sessions covering topics of general interest as well as parallel sessions
2528 for the many topical working groups in progress for the CWP.

2529

2530 **S^2I^2 -HEP/OSG/US-CMS/US-ATLAS Panel**

2531 *Date:* 8 Mar, 2017

2532 *Location:* UCSD/SDSC (La Jolla, CA)

2533 *URL:* <https://indico.fnal.gov/conferenceTimeTable.py?confId=12973#20170308>

2534 *Description:* This panel took place at Open Science Grid All Hands Meeting (OSG-AHM). Partic-
2535 ipants included Kaushik De (US-ATLAS), Peter Elmer (S^2I^2 -HEP, US-CMS), Oli Gutsche (US-
2536 CMS) and Mark Neubauer (S^2I^2 -HEP, US-ATLAS), with Frank Wuerthwein (OSG, US-CMS) as
2537 moderator. The goal was to inform the OSG community about the CWP and S^2I^2 -HEP processes
2538 and learn from the OSG experience.

2539

2540 **Software Triggers and Event Reconstruction WG meeting**

2541 *Date:* 9 Mar, 2017

2542 *Location:* LAL-Orsay (Orsay, France)

2543 *URL:* <https://indico.cern.ch/event/614111/>

2544 *Description:* This was a meeting of the Software Triggers and Event Reconstruction CWP working
2545 group. It was held as a parallel session at the “Connecting the Dots” workshop, which focuses on
2546 forward-looking pattern recognition and machine learning algorithms for use in HEP.

2547

2548 **IML Topical Machine Learning Workshop**

2549 *Date:* 20–22 Mar, 2017

2550 *Location:* CERN (Geneva, Switzerland)

2551 *URL:* <https://indico.cern.ch/event/595059>

2552 *Description:* This was a meeting of the Machine Learning CWP working group. It was held as
2553 a parallel session at the “Inter-experimental Machine Learning (IML)” workshop, an organization
2554 formed in 2016 to facilitate communication regarding R&D on ML applications in the LHC exper-
2555 iments.

2556

2557 **Community White Paper Follow-up at FNAL**

2558 *Date:* 23 Mar, 2017

2559 *Location:* FNAL (Batavia, IL)

2560 *URL:* <https://indico.fnal.gov/conferenceDisplay.py?confId=14032>

2561 *Description:* This one-day workshop was organized to engage with the experimental HEP commu-
2562 nity involved in computing and software for Intensity Frontier experiments at FNAL. Plans for the
2563 CWP and the S^2I^2 -HEP project were described, with discussion about commonalities between the
2564 HL-LHC challenges and the challenges of the FNAL neutrino and muon experiments.

2565

CWP Visualization Workshop

Date: 28–30 Mar, 2017

Location: CERN (Geneva, Switzerland)

URL: <https://indico.cern.ch/event/617054/>

Description: This workshop was organized by the Visualization CWP working group. It explored the current landscape of HEP visualization tools as well as visions for how these could evolve. There was participation both from HEP developers and industry.

2nd S^2I^2 HEP/CS Workshop

Date: 1–3 May, 2017

Location: Princeton University (Princeton, NJ)

URL: <https://indico.cern.ch/event/622920/>

Description: This 2nd HEP/CS workshop built on the discussions which took place at the the first S^2I^2 HEP/CS workshop to take a fresh look at planned HEP and computer science research and brainstorm about engaging specific areas of effort, perspectives, synergies and expertise of mutual benefit to HEP and CS communities, especially as it relates to a future NSF Software Institute for HEP.

DS@HEP 2017 (Data Science in High Energy Physics)

Date: 8–12 May, 2017

Location: FNAL (Batavia, IL)

URL: <https://indico.fnal.gov/conferenceDisplay.py?confId=13497>

Description: This was a meeting of the Machine Learning CWP working group. It was held as a parallel session at the “Data Science in High Energy Physics (DS@HEP)” workshop, a workshop series begun in 2015 to facilitate communication regarding R&D on ML applications in HEP.

HEP Analysis Ecosystem Retreat

Date: 22–24 May, 2017

Location: Amsterdam, the Netherlands

URL: <http://indico.cern.ch/event/613842/>

Summary report: <http://hepsoftwarefoundation.org/assets/AnalysisEcosystemReport20170804.pdf>

Description: This was a general workshop, organized about the HSF, about the ecosystem of analysis tools used in HEP and the ROOT software framework. The workshop focused both on the current status and the 5-10 year time scale covered by the CWP.

CWP Event Processing Frameworks Workshop

Date: 5-6 Jun, 2017

Location: FNAL (Batavia, IL)

URL: <https://indico.fnal.gov/conferenceDisplay.py?confId=14186>

Description: This was a workshop held by the Event Processing Frameworks CWP working group.

HEP Software Foundation Workshop

Date: 26–30 Jun, 2017

Location: LAPP (Annecy, France)

URL: <https://indico.cern.ch/event/613093/>

Description: This was the final general workshop for the CWP process. The CWP working groups

came together to present their status and plans, and develop consensus on the organization and context for the community roadmap. Plans were also made for the CWP writing phase that followed in the few months following this last workshop.

S^2I^2 -HEP Workshop

Date: 23–26 Aug, 2017

Location: University of Washington, Seattle (Seattle, WA)

URL: <https://indico.cern.ch/event/640290/>

Description: This final S^2I^2 -HEP workshop was held as a satellite workshop of the ACAT 2017 Conference. The workshop built on the emerging consensus from the CWP process and focused on the role an NSF-supported Software Institute could play. Specific discussions focused on establishing which areas would be both high impact and appropriate for leadership role in the U.S. universities. In addition the relative roles of an Institute, the US LHC Ops programs and the international LHC program were discussed, along with possible management structures for an Institute.

Data Organisation, Management and Access (DOMA) in Astronomy, Genomics and High Energy Physics

Date: 16–17 Nov, 2017

Location: Flatiron Institute (New York City, NY)

URL: <https://indico.cern.ch/event/669506/>

Description: This workshop was co-sponsored by the Simons Foundation and the S^2I^2 -HEP project. The workshop focused on the current research practices and future needs for Data Organization, Management and Access (DOMA) across the fields of Astronomy, Genomics and High Energy Physics. Discussions centered on identifying possibilities for integration across the different fields, as well as opportunities for common research and development activities.

S^2I^2 /DOE mini-workshop on HL-LHC Software and Computing R&D

Date: 28–29 Nov, 2017

Location: Catholic University of America (Washington DC)

URL: <https://indico.cern.ch/event/678121/>

Description: The goals of this workshop was to (1) review the vision for the ensemble of possible R&D efforts for the HL-LHC as articulated via the international CWP effort, and (2) articulate how R&D efforts such as an NSF S^2I^2 would interact with the US-LHC Operations programs and DOE efforts (in the context of the full, international efforts). Discuss the broad scope of relevant capabilities and current DOE and NSF funded efforts.

This full list of workshops and meetings (with links) is also available on the <http://s2i2-hep.org> website. In addition there were “internal” sessions regarding the CWP in the LHC experiment collaboration meetings, which are not listed above.

More than 260 people participated in one or more of the workshops which had an explicit registration and participant list. This does not include those who participated in the many “outreach” or panel sessions at pre-existing workshops/meetings such as DS@HEP, the OSG AHM, the IML Workshop or the sessions at LHC experiment collaboration meetings which not listed above, for which no explicit participant list was tracked. The combined list of known registered participants is:

Aaron Elliott (Aegis Research Labs), Aaron Dominguez (Catholic University of America), Aaron Sauers (Fermilab), Aashrita Mangu (California Institute of Technology), Abid Patwa (DOE), Adam

2660 Aurisano (University of Cincinnati), Adam Lyon (FNAL), Ajit Majumder (Wayne State), Alexei
 2661 Klimentov (Brookhaven National Lab), Alexey Svyatkovskiy (Princeton University), Alja Mrak
 2662 Tadel (University of California San Diego), Amber Boehnlein (Jefferson Lab), Amir Farbin (University
 2663 of Texas at Arlington), Amit Kumar (Southern Methodist), Andrea Dotti (SLAC National Accelerator
 2664 Laboratory), Andrea Rizzi (INFN-Pisa), Andrea Valassi (CERN), Andrei Gheata (CERN), Andrew
 2665 Gilbert (KIT), Andrew Hanushevsky (SLAC National Accelerator Laboratory), Anton Burtsev
 2666 (University of California, Irvine), Anton Poluektov (University of Warwick), Antonio Augusto Alves
 2667 Junior (University of Cincinnati), Antonio Limosani (CERN / University of Sydney), Anyes Taffard
 2668 (UC Irvine), Ariel Schwartzman (SLAC), Attila Krasznahorkay (CERN), Avi Yagil (UCSD), Axel
 2669 Naumann (CERN), Ben Hooberman (Illinois), Benedikt Hegner (CERN), Benedikt Riedel (University
 2670 of Chicago), Benjamin Couturier (CERN), Bill Nitzberg (Altair), Bo Jayatilaka (FNAL),
 2671 Bogdan Mihaila (NSF), Brian Bockelman (University of Nebraska - Lincoln), Brian O'Connor
 2672 (University of California at Santa Cruz), Burt Holzman (Fermilab), Carlos Maltzahn (University
 2673 of California - Santa Cruz), Catherine Biscarat (CNRS), Cecile Barbier (LAPP), Charles Leggett
 2674 (LBNL), Charlotte Lee (University of Washington), Chris Green (FNAL), Chris Tunnell (University
 2675 of Chicago, KICP), Christopher Jones (FNAL), Claudio Grandi (INFN), Conor Fitzpatrick
 2676 (EPFL), Daniel S. Katz (University of Illinois at Urbana-Champaign/NCSA), Dan Riley (Cornell
 2677 University), Daniel Whiteson (UC Irvine), Daniele Bonacorsi (University of Bologna), Danko
 2678 Adrovic (DePaul), Dario Berzano (CERN), Dario Menasce (INFN Milano-Bicocca), David Ab-
 2679 durachmanov (University of Nebraska-Lincoln), David Lange (Princeton University), David Lesny
 2680 (Illinois), David Malon (Argonne National Laboratory), David Rousseau (LAL-Orsay), David Smith
 2681 (CERN), Dick Greenwood (Louisiana Tech University), Dirk Duellmann (CERN), Dirk Hufnagel
 2682 (Fermilab), Don Petravick (Illinois/NCSA), Dorian Kcira (California Institute of Technology),
 2683 Doug Benjamin (Duke University), Doug Thain (Notre Dame), Douglas Thain (University of Notre
 2684 Dame), Dustin Anderson (California Institute of Technology), Dustin Tran (Columbia University),
 2685 Eduardo Rodrigues (University of Cincinnati), Elizabeth Sexton-Kennedy (FNAL), Enric Tejedor
 2686 Saavedra (CERN), Eric Lancon (BNL), Eric Vaandering (FNAL), Farah Hariri (CERN), Federico
 2687 Carminati (CERN), Fernanda Psihas (Indiana University), Fons Rademakers (CERN), Frank
 2688 Gaede (DESY), Frank Wuerthwein (University of California at San Diego/SDSC), Frederique Chol-
 2689 let (LAPP), Gabriel Perdue (Fermilab), Gerardo Ganis (CERN), Gerhard Raven (Nikhef), Giacomo
 2690 Govi (FNAL), Giacomo Tenaglia (CERN), Gianluca Cerminara (CERN), Giulio Eulisse (CERN),
 2691 Gloria Corti (CERN), Gordon Watts (University of Washington), Graeme Stewart (University of
 2692 Glasgow), Graham Mackintosh (IBM), H. Biral Runesha (University of Chicago), Hadrien Grasland
 2693 (Universite de Paris-Sud), Harvey Newman (Caltech), Helge Meinhard (CERN), Henry Schreiner III
 2694 (University of Cincinnati), Horst Severini (University of Oklahoma), Ian Bird (CERN), Ian Collier
 2695 (RAL), Ian Cosden (Princeton University), Ian Fisk (Simons Foundation), Ian Stockdale (Altair En-
 2696 gineering), Ilija Vukotic (University of Chicago), Isobel Ojalvo (Princeton University), Ivo Jimenez
 2697 UC (University of California - Santa Cruz), Jakob Blomer (CERN), Jamie Bedard (Siena College),
 2698 Jean Jacquemier (LAPP), Jean-Roch Vlimant (California Institute of Technology), Jeff Carver
 2699 (University of Alabama), Jeff Hammond (Intel), Jeff Lefevre (University of California at Santa
 2700 Cruz), Jeff Porter (LBNL), Jeff Templon (Nikhef), Jeffrey Carver (University of Alabama), Jerome
 2701 Lauret (BNL), Jim Kowalkowski (FNAL), Jim Pivarski (Princeton University), Johannes Albrecht
 2702 (TU Dortmund), John Apostolakis (CERN), John Harvey (CERN), John Towns (Illinois/NCSA),
 2703 Joon Kim (Princeton University), Joseph Boudreau (University of Pittsburgh), Justas Balcas (Cal-
 2704 tech), Justin Wozniak (University of Chicago/ANL), Karan Bhatia (Google Cloud), Karen Tomko
 2705 (Ohio Supercomputer Center), Kathryn Huff (Illinois), Kaushik De (University of Texas at Ar-
 2706 lington), Ken Bloom (University of Nebraska-Lincoln), Kevin Jorissen (Amazon Web Services),
 2707 Kevin Lannon (University of Notre Dame), Konstantin Toms (University of New Mexico), Kurt
 2708 Rinnert (U.Liverpool), Kyle Chard (University of Chicago), Kyle Cranmer (New York University),
 2709 Kyle Knoepfel (FNAL), Lauren Anderson (Flatiron Institute), Lawrence R Frank (UCSD), Lind-

2710 sey Gray (Fermilab), Liz Sexton-Kennedy (FNAL), Lorenzo Moneta (CERN), Lothar Bauerdick
 2711 (FNAL), Louis Capps (NVIDIA), Lukas Heinrich (New York University), Lukasz Kreczko (Bris-
 2712 tol), Madeline Hagen (Siena College), Makoto Asai (SLAC), Manish Parashar (Rutgers University),
 2713 Marc Paterno (FNAL), Marc Verderi (Ecole Polytechnique), Marcin Nowak (CERN), Maria Girone
 2714 (CERN), Maria Spiropulu (Caltech), Mario Lassnig (CERN), Mark Neubauer (University of Illi-
 2715 nois at Urbana-Champaign), Markus Klute (MIT), Markus Schulz (CERN), Martin Ritter (LMU
 2716 Munich), Matevz Tadel (UCSD), Matthew Bellis (Siena College), Matt Zhang (Illinois), Matthew
 2717 Feickert (Southern Methodist University), Matthew Turk (University of Illinois), Matthieu Lefeb-
 2718 vre (Princeton University), Max Baak (KPMG), Meghan Frate (University of California, Irvine),
 2719 Meghan Kane (SoundCloud, MIT), Michael Andrews (Carnegie Mellon University/CERN), Michael
 2720 Kirby (FNAL), Michael Sevilla (University of California, Santa Cruz), Michael Sokoloff (Univer-
 2721 sity of Cincinnati), Michel Jouvin (LAL/Universite de Paris-Sud), Michela Paganini (Yale Univer-
 2722 sity), Michela Taufer (University of Delaware), Mike Hildreth (University of Notre Dame), Mike
 2723 Williams (MIT), Miron Livny (University of Wisconsin-Madison), Mohammad Al-Turany (GSI),
 2724 Nadine Neyroud (LAPP), Nan Niu (University of Cincinnati), Nancy Wilkins-Diehr (University
 2725 of California San Diego), Natalia Volfovsky (Simons Foundation), Nathalie Rauschmayr (CERN),
 2726 Neil Ernst (Software Engineering Institute), Noah Watkins (University of California, Santa Cruz),
 2727 Oliver Gutsche (FNAL), Oliver Keeble (CERN), Panagiotis Spentzouris (FNAL), Paolo Calafiura
 2728 (LBNL), Parag Mhashikar (Fermilab), Patricia Mendez Lorenzo (CERN), Patrick Bos (Nether-
 2729 lands eScience Center), Patrick Skubic (University of Oklahoma), Patrick de Perio (Columbia Uni-
 2730 versity), Paul Laycock (CERN), Paul Mattione (Jefferson Lab), Paul Rossman (Google Inc.), Pere
 2731 Mato (CERN), Peter Elmer (Princeton University), Peter Hristov (CERN), Peter Onyisi (University
 2732 of Texas at Austin), Philippe Canal (FNAL), Pierre Aubert (LAPP), Rajesh Ranganath (Princeton
 2733 University), Riccardo Maria Bianchi (University of Pittsburgh), Richard Hay Jr (Princeton Uni-
 2734 versity), Richard Mount (SLAC), Rick Wagner (Globus), Rob Gardner (University of Chicago),
 2735 Rob Kutschke (FNAL), Rob Quick (Indiana University), Robert Illingworth (Fermilab), Robert
 2736 Kalescky (Southern Methodist), Robert Knight (Princeton University), Robert Kutschke (Fermi-
 2737 lab), Roger Jones (Lancaster), Ruslan Mashinistov (University of Texas at Arlington), Sabine Elles
 2738 (LAPP), Sally Seidel (New Mexico), Sandra Gesing (University of Notre Dame), Sandro Wen-
 2739 zel (CERN), Sascha Caron (Nikhef), Sebastien Binet (IN2P3/LPC), Sergei Gleyzer (University
 2740 of Florida), Shantenu Jha (Rutgers University), Shaun Astarabadi (Western Digital), Shawfeng
 2741 Dong (University of California at Santa Cruz), Shawn McKee (University of Michigan), Shy Genel
 2742 (Flatiron Institute), Simone Campana (CERN), Slava Krutelyov (University of California at San
 2743 Diego), Spencer Smith (McMaster University), Stefan Roiser (CERN), Steven Schramm (Univer-
 2744 site de Geneve), Sudhir Malik (University of Puerto Rico Mayaguez), Sumanth Mannam (DePaul),
 2745 Sumit Saluja (Princeton University), Sunita Chandrasekaran (University of Delaware), Tanu Malik
 2746 (Depaul University), Taylor Childers (Argonne Nat. Lab), Thomas Hacker (Purdue University),
 2747 Thomas Kuhr (LMU), Thomas McCauley (University of Notre Dame), Thomas Vuillaume (LAPP),
 2748 Thorsten Kollegger (GSI), Tom Gibbs (NVIDIA), Tom Lecompte (DOE/ANL), Tommaso Boccali
 2749 (INFN Pisa), Torre Wenaus (BNL), V. Daniel Elvira (Fermilab), Vakho Tsulaia (LBNL), Valentin
 2750 Kuznetsov (Cornell University), Vassil Vassilev (Princeton University), Vincent Croft (Nikhef),
 2751 Vinod Gupta (Princeton University), Vladimir Gligorov (CNRS), Wahid Bhimji (NERSC/LBNL),
 2752 Wenjing Wu (Institute of High Energy Physics, Beijing), Wouter Verkerke (Nikhef)

References

- [1] S. Campana, presentation to the 2016 Aix-les-Bains ECFA HL-LHC workshop, 3 Oct 2016. <https://indico.cern.ch/event/524795/contributions/2236590/attachments/1347419/2032314/ECFA2016.pdf>.
- [2] S2I2-HEP project webpage: <http://s2i2-hep.org>.
- [3] G. Aad et al. Observation of a new particle in the search for the Standard Model Higgs boson with the ATLAS detector at the LHC. *Phys.Lett.*, B716:1–29, 2012.
- [4] Serguei Chatrchyan et al. Observation of a new boson at a mass of 125 GeV with the CMS experiment at the LHC. *Phys.Lett.*, B716:30–61, 2012.
- [5] Gino Isidori, Yosef Nir, and Gilad Perez. Flavor Physics Constraints for Physics Beyond the Standard Model. *Ann.Rev.Nucl.Part.Sci.*, 60:355, 2010.
- [6] Particle Physics Project Prioritization Panel. Building for Discovery: Strategic Plan for U.S. Particle Physics in the Global Context. http://science.energy.gov/~media/hep/hepap/pdf/May%202014/FINAL_DRAFT2_P5Report_WEB_052114.pdf.
- [7] ALICE Collaboration public website. <http://aliceinfo.cern.ch/>.
- [8] D Lucchesi. Computing Resources Scrutiny Group Report. Technical Report CERN-RRB-2016-049, CERN, Geneva, Feb 2016.
- [9] Concezio Bozzi. LHCb Computing Resources: 2019 requests and reassessment of 2018 requests. Technical Report LHCb-PUB-2017-019. CERN-LHCb-PUB-2017-019, CERN, Geneva, Sep 2017.
- [10] S. Foffano. Computing Resources Scrutiny Group Report. Technical Report CERN-RRB-2017-067, CERN, Geneva, April 2017.
- [11] Samuel H. Fuller and Editors; Committee on Sustaining Growth in Computing Performance; National Research Council Lynette I. Millett. *The Future of Computing Performance: Game Over or Next Level?* The National Academies Press, 2011.
- [12] M. Butler, R. Mount, and M. Hildreth. Snowmass 2013 Computing Frontier Storage and Data Management. *ArXiv e-prints*, November 2013.
- [13] HSF Community White Paper webpages. <http://hepsoftwarefoundation.org/activities/cwp.html>.
- [14] Charge for Producing the HSF Community White Paper. <http://hepsoftwarefoundation.org/assets/CWP-Charge-HSF.pdf>.
- [15] Summary of Workshop 1: Fostering Collaboration between HEP and Computer Science Communities. <http://s2i2-hep.org/downloads/s2i2-hep-cs-workshop-summary.pdf>.
- [16] NSF Software Infrastructure for Sustained Innovation (SI2) Program page. https://www.nsf.gov/funding/pgm_summ.jsp?pims_id=503489.
- [17] Ruth Pordes, Don Petravick, Bill Kramer, Doug Olson, Miron Livny, Alain Roy, Paul Avery, Kent Blackburn, Torre Wenaus, Frank Wuerthwein, Ian Foster, Rob Gardner, Mike Wilde, Alan Blatecky, John McGee, and Rob Quick. The open science grid. *Journal of Physics: Conference Series*, 78(1):012057, 2007.

- [18] Open Science Grid webpage: <https://www.opensciencegrid.org>.
- [19] CHEP 2016 conference webpage: <http://chep2016.org>.
- [20] ACAT 2017 conference webpage: <https://indico.cern.ch/event/567550/>.
- [21] ROOT home page. <http://root.cern.ch/drupal/>.
- [22] S. Agostinelli, J. Allison, K. Amako, J. Apostolakis, H. Araujo, P. Arce, M. Asai, D. Axen, S. Banerjee, G. Barrand, F. Behner, L. Bellagamba, J. Boudreau, L. Broglia, A. Brunengo, H. Burkhardt, S. Chauvie, J. Chuma, R. Chytrcek, G. Cooperman, G. Cosmo, P. Degt-yarenko, A. Dell’Acqua, G. Depaola, D. Dietrich, R. Enami, A. Feliciello, C. Ferguson, H. Fesefeldt, G. Folger, F. Foppiano, A. Forti, S. Garelli, S. Giani, R. Giannitrapani, D. Gibin, J.J. Gmez Cadenas, I. Gonzlez, G. Gracia Abril, G. Greeniaus, W. Greiner, V. Grichine, A. Grossheim, S. Guatelli, P. Gumplinger, R. Hamatsu, K. Hashimoto, H. Hasui, A. Heikkinen, A. Howard, V. Ivanchenko, A. Johnson, F.W. Jones, J. Kallenbach, N. Kanaya, M. Kawabata, Y. Kawabata, M. Kawaguti, S. Kelner, P. Kent, A. Kimura, T. Kodama, R. Kokoulin, M. Kossov, H. Kurashige, E. Lamanna, T. Lampn, V. Lara, V. Lefebure, F. Lei, M. Liendl, W. Lockman, F. Longo, S. Magni, M. Maire, E. Medernach, K. Minamimoto, P. Mora de Freitas, Y. Morita, K. Murakami, M. Nagamatu, R. Nartallo, P. Nieminen, T. Nishimura, K. Ohtsubo, M. Okamura, S. O’Neale, Y. Oohata, K. Paech, J. Perl, A. Pfeiffer, M.G. Pia, F. Ran-jard, A. Rybin, S. Sadilov, E. Di Salvo, G. Santin, T. Sasaki, N. Savvas, Y. Sawada, S. Scherer, S. Sei, V. Sirotenko, D. Smith, N. Starkov, H. Stoecker, J. Sulkimo, M. Takahata, S. Tanaka, E. Tcherniaev, E. Safai Tehrani, M. Tropeano, P. Truscott, H. Uno, L. Urban, P. Urban, M. Verderi, A. Walkden, W. Wander, H. Weber, J.P. Wellisch, T. Wenaus, D.C. Williams, D. Wright, T. Yamada, H. Yoshida, and D. Zschesche. Geant4a simulation toolkit. *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, 506(3):250 – 303, 2003.
- [23] J. Allison, K. Amako, J. Apostolakis, H. Araujo, P. Arce Dubois, M. Asai, G. Barrand, R. Capra, S. Chauvie, R. Chytrcek, G. A. P. Cirrone, G. Cooperman, G. Cosmo, G. Cuttone, G. G. Daquino, M. Donszelmann, M. Dressel, G. Folger, F. Foppiano, J. Generowicz, V. Grichine, S. Guatelli, P. Gumplinger, A. Heikkinen, I. Hrivnacova, A. Howard, S. Incerti, V. Ivanchenko, T. Johnson, F. Jones, T. Koi, R. Kokoulin, M. Kossov, H. Kurashige, V. Lara, S. Larsson, F. Lei, O. Link, F. Longo, M. Maire, A. Mantero, B. Mascialino, I. McLaren, P. Mendez Lorenzo, K. Minamimoto, K. Murakami, P. Nieminen, L. Pandola, S. Parlati, L. Peralta, J. Perl, A. Pfeiffer, M. G. Pia, A. Ribon, P. Rodrigues, G. Russo, S. Sadilov, G. Santin, T. Sasaki, D. Smith, N. Starkov, S. Tanaka, E. Tcherniaev, B. Tome, A. Trindade, P. Truscott, L. Urban, M. Verderi, A. Walkden, J. P. Wellisch, D. C. Williams, D. Wright, and H. Yoshida. Geant4 developments and applications. *IEEE Transactions on Nuclear Science*, 53(1):270–278, Feb 2006.
- [24] J. Allison, K. Amako, J. Apostolakis, P. Arce, M. Asai, T. Aso, E. Bagli, A. Bagulya, S. Banerjee, G. Barrand, B.R. Beck, A.G. Bogdanov, D. Brandt, J.M.C. Brown, H. Burkhardt, Ph. Canal, D. Cano-Ott, S. Chauvie, K. Cho, G.A.P. Cirrone, G. Cooperman, M.A. Corts-Giraldo, G. Cosmo, G. Cuttone, G. Depaola, L. Desorgher, X. Dong, A. Dotti, V.D. Elvira, G. Folger, Z. Francis, A. Galoyan, L. Garnier, M. Gayer, K.L. Genser, V.M. Gri-chine, S. Guatelli, P. Guye, P. Gumplinger, A.S. Howard, I. Hivnov, S. Hwang, S. Incerti, A. Ivanchenko, V.N. Ivanchenko, F.W. Jones, S.Y. Jun, P. Kaitaniemi, N. Karakatsanis, M. Karamitros, M. Kelsey, A. Kimura, T. Koi, H. Kurashige, A. Lechner, S.B. Lee, F. Longo, M. Maire, D. Mancusi, A. Mantero, E. Mendoza, B. Morgan, K. Murakami, T. Nikitina, L. Pandola, P. Paprocki, J. Perl, I. Petrovi, M.G. Pia, W. Pokorski, J.M. Quesada, M. Raine,

M.A. Reis, A. Ribon, A. Risti Fira, F. Romano, G. Russo, G. Santin, T. Sasaki, D. Sawkey, J.I. Shin, I.I. Strakovsky, A. Taborda, S. Tanaka, B. Tom, T. Toshito, H.N. Tran, P.R. Truscott, L. Urban, V. Uzhinsky, J.M. Verbeke, M. Verderi, B.L. Wendt, H. Wenzel, D.H. Wright, D.M. Wright, T. Yamashita, J. Yarba, and H. Yoshida. Recent developments in geant4. *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, 835(Supplement C):186 – 225, 2016.

[25] G. Barrand et al. GAUDI - The software architecture and framework for building LHCb data processing applications. In *Proceedings, 11th International Conference on Computing in High-Energy and Nuclear Physics (CHEP 2000)*, pages 92–95, 2000.

[26] Eulisse G. and Tuura L. IgProf profiling tool. In *Proceedings, 14th International Conference on Computing in High-Energy and Nuclear Physics (CHEP 2004)*, 2004.

[27] Wouter Verkerke and David P. Kirkby. The RooFit toolkit for data modeling. *eConf*, C0303241:MOLT007, 2003.

[28] Andreas Hoecker, Peter Speckmayer, Joerg Stelzer, Jan Therhaag, Eckhard von Toerne, and Helge Voss. TMVA: Toolkit for Multivariate Data Analysis. *PoS*, ACAT:040, 2007.

[29] T. Gleisberg, Stefan. Hoeche, F. Krauss, M. Schonherr, S. Schumann, F. Siegert, and J. Winter. Event generation with SHERPA 1.1. *JHEP*, 02:007, 2009.

[30] Michelangelo L. Mangano, Fulvio Piccinini, Antonio D. Polosa, Mauro Moretti, and Roberto Pittau. ALPGEN, a generator for hard multiparton processes in hadronic collisions. *Journal of High Energy Physics*, 2003(07):001, 2003.

[31] Matteo Cacciari, Gavin P. Salam, and Gregory Soyez. FastJet User Manual. *Eur. Phys. J.*, C72:1896, 2012.

[32] Matteo Cacciari and Gavin P. Salam. Dispelling the N^3 myth for the k_t jet-finder. *Phys. Lett.*, B641:57–61, 2006.

[33] Kosyakov S. et al. FRONTIER: HIGH PERFORMANCE DATABASE ACCESS USING STANDARD WEB COMPONENTS IN A SCALABLE MULTI-TIER ARCHITECTURE. In *Proceedings, 14th International Conference on Computing in High-Energy and Nuclear Physics (CHEP 2004)*, 2004.

[34] A Dorigo, P Elmer, F Furano, and A Hanushevsky. XROOTD - A highly scalable architecture for data access. *WSEAS Transactions on Computers*, 4.3, 2005.

[35] Patrick Fuhrmann. dCache: the commodity cache. In *In Twelfth NASA Goddard and Twenty First IEEE Conference on Mass Storage Systems and Technologies*, 2004.

[36] Andreas J Peters and Lukasz Janyst. Exabyte Scale Storage at CERN. *Journal of Physics: Conference Series*, 331(5):052015, 2011.

[37] AJ Peters, EA Sindrilaru, and G Adde. EOS as the present and future solution for data storage at CERN. *Journal of Physics: Conference Series*, 664(4):042042, 2015.

[38] A A Ayllon, M Salichos, M K Simon, and O Keeble. Fts3: New data movement service for wlcg. *Journal of Physics: Conference Series*, 513(3):032081, 2014.

[39] Jakob Blomer, Carlos Aguado-Sanchez, Predrag Buncic, and Artem Harutyunyan. Distributing LHC application software and conditions databases using the CernVM file system. *Journal of Physics: Conference Series*, 331(4):042003, 2011.

2879 [40] I Sfiligoi. glideinWMS - a generic pilot-based workload management system. *Journal of*
2880 *Physics: Conference Series*, 119(6):062044, 2008.

2881 [41] P Nilsson, J Caballero, K De, T Maeno, A Stradling, T Wenaus, and the Atlas Collabo-
2882 ration. The ATLAS PanDA Pilot in Operation. *Journal of Physics: Conference Series*,
2883 331(6):062040, 2011.

2884 [42] T Maeno, K De, T Wenaus, P Nilsson, R Walker, A Stradling, V Fine, M Potekhin, S Pan-
2885 itkin, and G Compostella. Evolution of the ATLAS PanDA Production and Distributed
2886 Analysis System. *Journal of Physics: Conference Series*, 396(3):032071, 2012.

2887 [43] Douglas Thain, Todd Tannenbaum, and Miron Livny. Distributed computing in practice: the
2888 Condor experience. *Concurrency - Practice and Experience*, 17(2-4):323–356, 2005.

2889 [44] Douglas Thain and Miron Livny. Parrot: Transparent user-level middleware for data-intensive
2890 computing. *Scalable Computing: Practice and Experience*, 6(3), 2005.

2891 [45] P Ferreira, T Baron, C Bossy, J B Gonzalez, M Pugh, A Resco, J Trzaskoma, and C Wachter.
2892 Indico: A collaboration hub. *Journal of Physics: Conference Series*, 396(6):062006, 2012.

2893 [46] J B Gonzalez Lopez, A Avils, T Baron, P Ferreira, B Kolobara, M A Pugh, A Resco, and
2894 J P Trzaskoma. Indico 1.0. *Journal of Physics: Conference Series*, 513(6):062020, 2014.

2895 [47] NSF 15-553. <https://www.nsf.gov/pubs/2015/nsf15553/nsf15553.htm>.

2896 [48] Principles of Agile Software Development. [http://agilemanifesto.org/iso/en/](http://agilemanifesto.org/iso/en/principles.html)
2897 [principles.html](http://agilemanifesto.org/iso/en/principles.html).

2898 [49] CMSSW, <https://github.com/cms-sw/cmssw>.

2899 [50] Based on schedule graphic from HL-LHC Project website (Dec 2017):. [https://](https://project-hl-lhc-industry.web.cern.ch/content/project-schedule)
2900 project-hl-lhc-industry.web.cern.ch/content/project-schedule.

2901 [51] Fons Rademakers and Rene Brun. ROOT: an object-oriented data analysis framework. *Linux*
2902 *J.*, page 6.

2903 [52] Forum on the Interpretation of the LHC Results for BSM studies:. [https://twiki.cern.](https://twiki.cern.ch/twiki/bin/view/LHCPhysics/InterpretingLHCresults)
2904 [ch/twiki/bin/view/LHCPhysics/InterpretingLHCresults](https://twiki.cern.ch/twiki/bin/view/LHCPhysics/InterpretingLHCresults).

2905 [53] Kyle Cranmer and Itay Yavin. Recast — extending the impact of existing analyses. *Journal*
2906 *of High Energy Physics*, 2011(4):38, Apr 2011.

2907 [54] INSPIRE HEP website:. <http://inspirehep.net/>.

2908 [55] Eamonn Maguire, Lukas Heinrich, and Graeme Watt. HEPData: a repository for high energy
2909 physics data. *J. Phys. Conf. Ser.*, 898(10):102006, 2017.

2910 [56] HEPData website:. <https://hepdata.net>.

2911 [57] J Cowton, S Dallmeier-Tiessen, P Fokianos, L Rueda, P Herterich, J Kunar, T imko, and
2912 T Smith. Open data and data analysis preservation services for lhc experiments. *Journal of*
2913 *Physics: Conference Series*, 664(3):032030, 2015.

2914 [58] CERN Analysis Preservation website:. <https://analysispreservation.cern.ch/welcome>.

2915 [59] REANA github:. <http://reanahub.io/>.

- 2916 [60] G. Aad *et al.* [ATLAS Collaboration]. Observation of a new particle in the search for the
2917 Standard Model Higgs boson with the ATLAS detector at the LHC. *Phys. Lett. B*, 716:1,
2918 2012.
- 2919 [61] S. Chatrchyan *et al.* [CMS Collaboration]. Observation of a new boson at a mass of 125 GeV
2920 with the CMS experiment at the LHC. *Phys. Lett. B*, 716:30–61, 2012.
- 2921 [62] Florian Goertz. Accessing Masses Beyond Collider Reach - in EFT. 2017.
- 2922 [63] Philip Chang, Sergei Gleyzer, Mark Neubauer, and Dewen Zhong. Sustainable Matrix Ele-
2923 ment Method. Technical Report HSF-CWP-018, Aug 2017. This document can be found at
2924 <http://hepsoftwarefoundation.org/cwp-whitepapers.html>.
- 2925 [64] Kyle Cranmer. NIPS 2016 Keynote: Machine Learning & Likelihood Free Inference in Particle
2926 Physics. 12 2016.
- 2927 [65] Science Gateways Community Institute website:. <https://sciencegateways.org>.
- 2928 [66] Zaven Akopov et al. Status Report of the DPHEP Study Group: Towards a Global Effort
2929 for Sustainable Data Preservation in High Energy Physics. 2012.
- 2930 [67] ATLAS collaboration (2014). ATLAS Data Access Policy. CERN Open Data Portal:. [http:](http://doi.org/10.7483/OPENDATA.ATLAS.T9YR.Y7MZ)
2931 [//doi.org/10.7483/OPENDATA.ATLAS.T9YR.Y7MZ](http://doi.org/10.7483/OPENDATA.ATLAS.T9YR.Y7MZ).
- 2932 [68] ATLAS Phase-II Upgrade Scoping Document. Technical Report CERN-LHCC-2015-020.
2933 LHCC-G-166, CERN, Geneva, Sep 2015.
- 2934 [69] D Contardo, M Klute, J Mans, L Silvestris, and J Butler. Technical Proposal for the Phase-
2935 II Upgrade of the CMS Detector. Technical Report CERN-LHCC-2015-010. LHCC-P-008.
2936 CMS-TDR-15-02, Geneva, Jun 2015.
- 2937 [70] LHCb Trigger and Online Upgrade Technical Design Report. Technical Report CERN-LHCC-
2938 2014-016. LHCB-TDR-016, May 2014.
- 2939 [71] P Buncic, M Krzewicki, and P Vande Vyvre. Technical Design Report for the Upgrade of the
2940 Online-Offline Computing System. Technical Report CERN-LHCC-2015-006. ALICE-TDR-
2941 019, Apr 2015.
- 2942 [72] I Bird, P Buncic, F Carminati, M Cattaneo, P Clarke, I Fisk, M Girone, J Harvey, B Kersevan,
2943 P Mato, R Mount, and B Panzer-Steindel. Update of the Computing Models of the WLCG
2944 and the LHC Experiments. Technical Report CERN-LHCC-2014-014. LCG-TDR-002, Apr
2945 2014.
- 2946 [73] R. Aaij, S. Amato, L. Anderlini, S. Benson, M. Cattaneo, M. Clemencic, B. Couturier,
2947 M. Frank, V.V. Gligorov, T. Head, C. Jones, I. Komarov, O. Lupton, R. Matev, G. Raven,
2948 B. Sciascia, T. Skwarnicki, P. Spradlin, S. Stahl, B. Storaci, and M. Vesterinen. Tesla : an
2949 application for real-time data analysis in High Energy Physics. *Comput. Phys. Commun.*,
2950 208(CERN-LHCB-DP-2016-001. CERN-LHCB-DP-2016-001):35–42. 8 p, Apr 2016. 14 pages,
2951 8 figures.
- 2952 [74] R Abreu. The upgrade of the ATLAS High Level Trigger and Data Acquisition systems and
2953 their integration. Technical Report ATL-DAQ-PROC-2014-002, CERN, Geneva, May 2014.

- [75] CMS Collaboration. Search for narrow resonances in dijet final states at $\sqrt{s} = 8$ TeV with the novel CMS technique of data scouting. *Phys. Rev. Lett.*, 117(CMS-EXO-14-005. CMS-EXO-14-005. CERN-EP-2016-090):031802. 17 p, Apr 2016. Replaced with published version. All the figures and tables can be found at <http://cms-results.web.cern.ch/cms-results/public-results/publications/EXO-14-005/index.html>.
- [76] SciKIT-LEARN webpage: <http://scikit-learn.org/>.
- [77] F. Chollet, KERAS (2017), GitHub, <https://github.com/fchollet/keras>.
- [78] A. Rogozhnikov *et al.*, REP (2017), GitHub <https://github.com/yandex/rep>.
- [79] A. Rogozhnikov *et al.*, HEPML (2017), GitHub https://github.com/arogozhnikov/hep_ml.
- [80] J. Snoek, SPEARMINT (2017), Github: <https://github.com/HIPS/Spearmint>.
- [81] Philip Ilten, Mike Williams, and Yunjie Yang. Event generator tuning using Bayesian optimization. 2016.
- [82] TUNEMC GitHub repository: <https://github.com/yunjie-yang/TuneMC>.
- [83] Douglas Thain, Sander Klous, Justin Wozniak, Paul Brenner, Aaron Striegel, and Jesus Izaguirre. Separating abstractions from resources in a tactical storage system. In *Proceedings of the 2005 ACM/IEEE conference on Supercomputing*, page 55. IEEE Computer Society, 2005.
- [84] I. Bird, presentation at the WLCG workshop 2016 in conjunction with CHEP 2016, 8 Oct 2016. <https://indico.cern.ch/event/555063/contributions/2285838/>.
- [85] Kenneth Bloom and the CMS Collaboration. CMS Use of a Data Federation. *Journal of Physics: Conference Series*, 513(4):042005, 2014.
- [86] Kenneth Bloom et al. Any Data, Any Time, Anywhere: Global Data Access for Science. 2015.
- [87] Named Data Networks executive summary webpage: <https://named-data.net/project/execsummary/>.
- [88] National Academies of Sciences, Engineering, and Medicine. *Future Directions for NSF Advanced Computing Infrastructure to Support U.S. Science and Engineering in 2017-2020*. The National Academies Press, Washington, DC, 2016.
- [89] Software Sustainability Institute. <https://www.software.ac.uk>.
- [90] DIANA/HEP website. <http://diana-hep.org>.
- [91] Indico category for DIANA/HEP topical meetings. <https://indico.cern.ch/category/7192/>.
- [92] HEP Center for Computational Excellence. <http://hepcce.org>.
- [93] <https://cms-big-data.github.io>. <https://cms-big-data.github.io>.
- [94] IPCC ROOT - Princeton/Intel Parallel Computing Center to Modernize the ROOT Math and I/O Libraries. <https://ipcc-root.github.io>.
- [95] HEPCloud: a new paradigm for particle physics computing. <http://hepcloud.fnal.gov>.

2991 [96] HEPCloud: Provisioning 160,000 Compute Cores for Science. [http://hepcloud.fnal.gov/](http://hepcloud.fnal.gov/wp-content/uploads/2016/05/HEPCloud-DPF.pdf)
2992 [wp-content/uploads/2016/05/HEPCloud-DPF.pdf](http://hepcloud.fnal.gov/wp-content/uploads/2016/05/HEPCloud-DPF.pdf).

2993 [97] B. Holzman, L. A. T. Bauerdick, B. Bockelman, D. Dykstra, I. Fisk, S. Fuess, G. Garzoglio,
2994 M. Girone, O. Gutsche, D. Hufnagel, H. Kim, R. Kennedy, N. Magini, D. Mason, P. Spent-
2995 zouris, A. Tiradani, S. Timm, and E. W. Vaandering. HEPCloud, a New Paradigm for HEP
2996 Facilities: CMS Amazon Web Services Investigation. *ArXiv e-prints*, September 2017.

2997 [98] CERN openlab webpage. <http://openlab.cern>.

2998 [99] S. Malik, F. Hoehle, K. Lassila-Perini, A. Hinzmann, R. Wolf, et al. Maintaining and improv-
2999 ing of the training program on the analysis software in CMS. *J.Phys.Conf.Ser.*, 396:062013,
3000 2012.

3001 [100] LHCb Starter Kit Webpage. <https://lhcb.github.io/starterkit/>.

3002 [101] CERN School of Computing. <https://csc.web.cern.ch/>.

3003 [102] GridKa School (KIT). <http://gridka-school.scc.kit.edu/>.

3004 [103] ESC17 school webpage: <https://web.infn.it/esc17/index.php>.

3005 [104] CoDaS-HEP school webpage: <http://codas-hep.org>.

3006 [105] Segregated Valley: the ugly truth about Google and diversity
3007 in tech: [https://www.theguardian.com/technology/2017/aug/07/](https://www.theguardian.com/technology/2017/aug/07/silicon-valley-google-diversity-black-women-workers)
3008 [silicon-valley-google-diversity-black-women-workers](https://www.theguardian.com/technology/2017/aug/07/silicon-valley-google-diversity-black-women-workers).

3009 [106] Data and Software Preservation for Open Science (DASPOS) website. [https://daspos.crc.](https://daspos.crc.nd.edu)
3010 [nd.edu](https://daspos.crc.nd.edu).

3011 [107] Parallel Kalman Filter Tracking website. <http://trackreco.github.io>.